

CS 7650: Natural Language Processing

Alan Ritter

Administrivia

- ▶ Course website:
<https://aritter.github.io/CS-7650-sp22/>
- ▶ Piazza and Gradescope: links on the course website
 - ▶ We will do our best to answer questions within 24 hours (or Monday for questions asked over the weekend).
- ▶ TA Office hours:
 - ▶ See spreadsheet
<https://tinyurl.com/7650-TAs>

Instructor



Alan Ritter

alan.ritter@cc.gatech.edu

Teaching Assistants

Aman Khullar

akhullar8@gatech.edu

Raj Janardhan

rjanardhan3@gatech.edu

Rucha Sathe

ruchasathe@gatech.edu

Shuyan Lin

slin915@gatech.edu

Prerequisites

- ▶ Probability/Statistics
- ▶ Linear Algebra
- ▶ Multivariable Calculus
- ▶ Programming / Python experience
- ▶ A Machine Learning Course (otherwise this class will be a LOT more work)

Prerequisites

- ▶ Probability/Statistics
- ▶ Linear Algebra
- ▶ Multivariable Calculus
- ▶ Programming / Python experience
- ▶ A Machine Learning Course (otherwise this class will be a LOT more work)

There will be a lot of math and programming!

Coursework

- ▶ 3 Programming Projects (fairly substantial implementation effort)
 - ▶ Text classification
 - ▶ Named entity recognition (BiLSTM-CNN-CRF)
 - ▶ Neural chatbot (Seq2Seq with attention)

Coursework

- ▶ 3 Programming Projects (fairly substantial implementation effort)
 - ▶ Text classification
 - ▶ Named entity recognition (BiLSTM-CNN-CRF)
 - ▶ Neural chatbot (Seq2Seq with attention)
- ▶ 3 written assignments + midterm exam
- ▶ Mostly math problems related to ML / NLP

Coursework

- ▶ 3 Programming Projects (fairly substantial implementation effort)
 - ▶ Text classification
 - ▶ Named entity recognition (BiLSTM-CNN-CRF)
 - ▶ Neural chatbot (Seq2Seq with attention)
- ▶ 3 written assignments + midterm exam
 - ▶ Mostly math problems related to ML / NLP
- ▶ Final project

Coursework

- ▶ 3 Programming Projects (fairly substantial implementation effort)
 - ▶ Text classification
 - ▶ Named entity recognition (BiLSTM-CNN-CRF)
 - ▶ Neural chatbot (Seq2Seq with attention)
- ▶ 3 written assignments + midterm exam
 - ▶ Mostly math problems related to ML / NLP
- ▶ Final project
- ▶ Problem Set 0 (background review) is out now and **due Thursday**.

Final Project

- ▶ Final project (20%)
 - ▶ Groups of 3-4 recommended.
 - ▶ 1 is possible, but will require more work. Grading requirements are the same for individual projects.
 - ▶ 4 page report.
 - ▶ We will discuss more and have a “project kickoff” later in the semester.

Problem Set 1 (Background Review)

- ▶ Due this Thursday.
- ▶ Background review on probability, linear algebra, calculus.
- ▶ **Waitlisted students:** please submit PS1 by Friday if you plan to enroll in the course.
 - ▶ We can't predict whether or not you will get in, as this depends on other students dropping the class...
- ▶ Submit on Gradescope

Project 1 is also out (please look!)

The screenshot shows a Jupyter Notebook interface. The title bar reads "TextClassification_release_v2.ipynb". The menu bar includes "File", "Edit", "View", "Insert", "Runtime", "Tools", "Help", and "All changes saved". Below the menu is a toolbar with "+ Code" and "+ Text" buttons. On the left, there are icons for search, file, and cell. A code cell contains the following text:

```
# Licensing Information: You are free to use or extend this project for  
# educational purposes provided that (1) you do not distribute or publish  
# solutions, (2) you retain this notice, and (3) you provide clear  
# attribution to The Georgia Institute of Technology, including a link to https://aritter.github.io/CS-7650/  
  
# Attribution Information: This assignment was developed at The Georgia Institute of Technology  
# by Alan Ritter (alan.ritter@cc.gatech.edu)
```

Project #1: Text Classification

In this assignment, you will implement the perceptron algorithm, and a simple, but competitive neural bag-of-words model, as described in [this paper](#) for text classification. You will train your models on a (provided) dataset of positive and negative movie reviews and report accuracy on a test set.

In this notebook, we provide you with starter code to read in the data and evaluate the performance of your models. After completing the instructions below, please follow the instructions at the end to submit your notebook and other files to Gradescope.

Make sure to make a copy of this notebook, so your changes are saved.

▼ Download the dataset

First you will need to download the IMDB dataset - to do this, simply run the cell below. We have prepared a small version of the ACL IMDB

Free Textbooks!



- ▶ 2 excellent textbooks for NLP
 - ▶ There will be assigned readings from both
 - ▶ Both freely available online

Natural Language Processing

Speech and Language Processing (3rd ed. draft)

[Dan Jurafsky](#) and [James H. Martin](#)

Jacob Eisenstein

Not free: GPUs



- ▶ Modern NLP methods require non-trivial computation
 - ▶ Training neural networks with many parameters can take a long time (it is a very good idea to start working on the assignments early!)
 - ▶ This is a big part of modern NLP methods. It is important to get experience training these networks.
 - ▶ You will need to use GPUs to complete the programming assignments.
 - ▶ Google Colab: has free GPUs, but with some big limitations that will make the assignments very difficult to complete.
 - ▶ The programming projects are designed with Colab in mind
 - ▶ Colab Pro subscription (\$10/month). This is highly recommended once we start working with PyTorch.

Outline of the Course

- Machine Learning Review (Naive Bayes, Log. Reg. SVMs, Neural Nets)
- Sequence Models (HMMs, CRFs)
- Word Embeddings
- Neural Networks in NLP (NBOW, RNNs, CNNs, Transformers)
- Pre-trained models (e.g. BERT, BART, T5, GPT)
- Machine Translation
- Dialogue
- Question Answering

Course Goals

Course Goals

- ▶ Cover fundamental machine learning techniques used in NLP

Course Goals

- ▶ Cover fundamental machine learning techniques used in NLP
- ▶ Know about modern NLP methods: what is the state-of-the-art in 2023?

Course Goals

- ▶ Cover fundamental machine learning techniques used in NLP
- ▶ Know about modern NLP methods: what is the state-of-the-art in 2023?
- ▶ Make you a “producer” rather than a “consumer” of NLP tools

Course Goals

- ▶ Cover fundamental machine learning techniques used in NLP
- ▶ Know about modern NLP methods: what is the state-of-the-art in 2023?
- ▶ Make you a “producer” rather than a “consumer” of NLP tools
 - ▶ The assignments should teach you what you need to know to understand nearly any modern NLP system.

Course Goals

- ▶ Cover fundamental machine learning techniques used in NLP
 - ▶ Deeper understanding of algorithms beyond “how to use ML/NLP libraries”.
- ▶ Know about modern NLP methods: what is the state-of-the-art in 2023?
- ▶ Make you a “producer” rather than a “consumer” of NLP tools
 - ▶ The assignments should teach you what you need to know to understand nearly any modern NLP system.

Assignments

- ▶ 3 Programming Assignments
 - ▶ Implementation-oriented
 - ▶ ~2 weeks per assignment

Assignments

- ▶ 3 Programming Assignments
 - ▶ Implementation-oriented
 - ▶ ~2 weeks per assignment

These projects require understanding of the concepts, ability to write performant code, and ability to think about how to debug complex systems. **They are challenging, so start early!**

What's the goal of NLP?

What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text

What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems

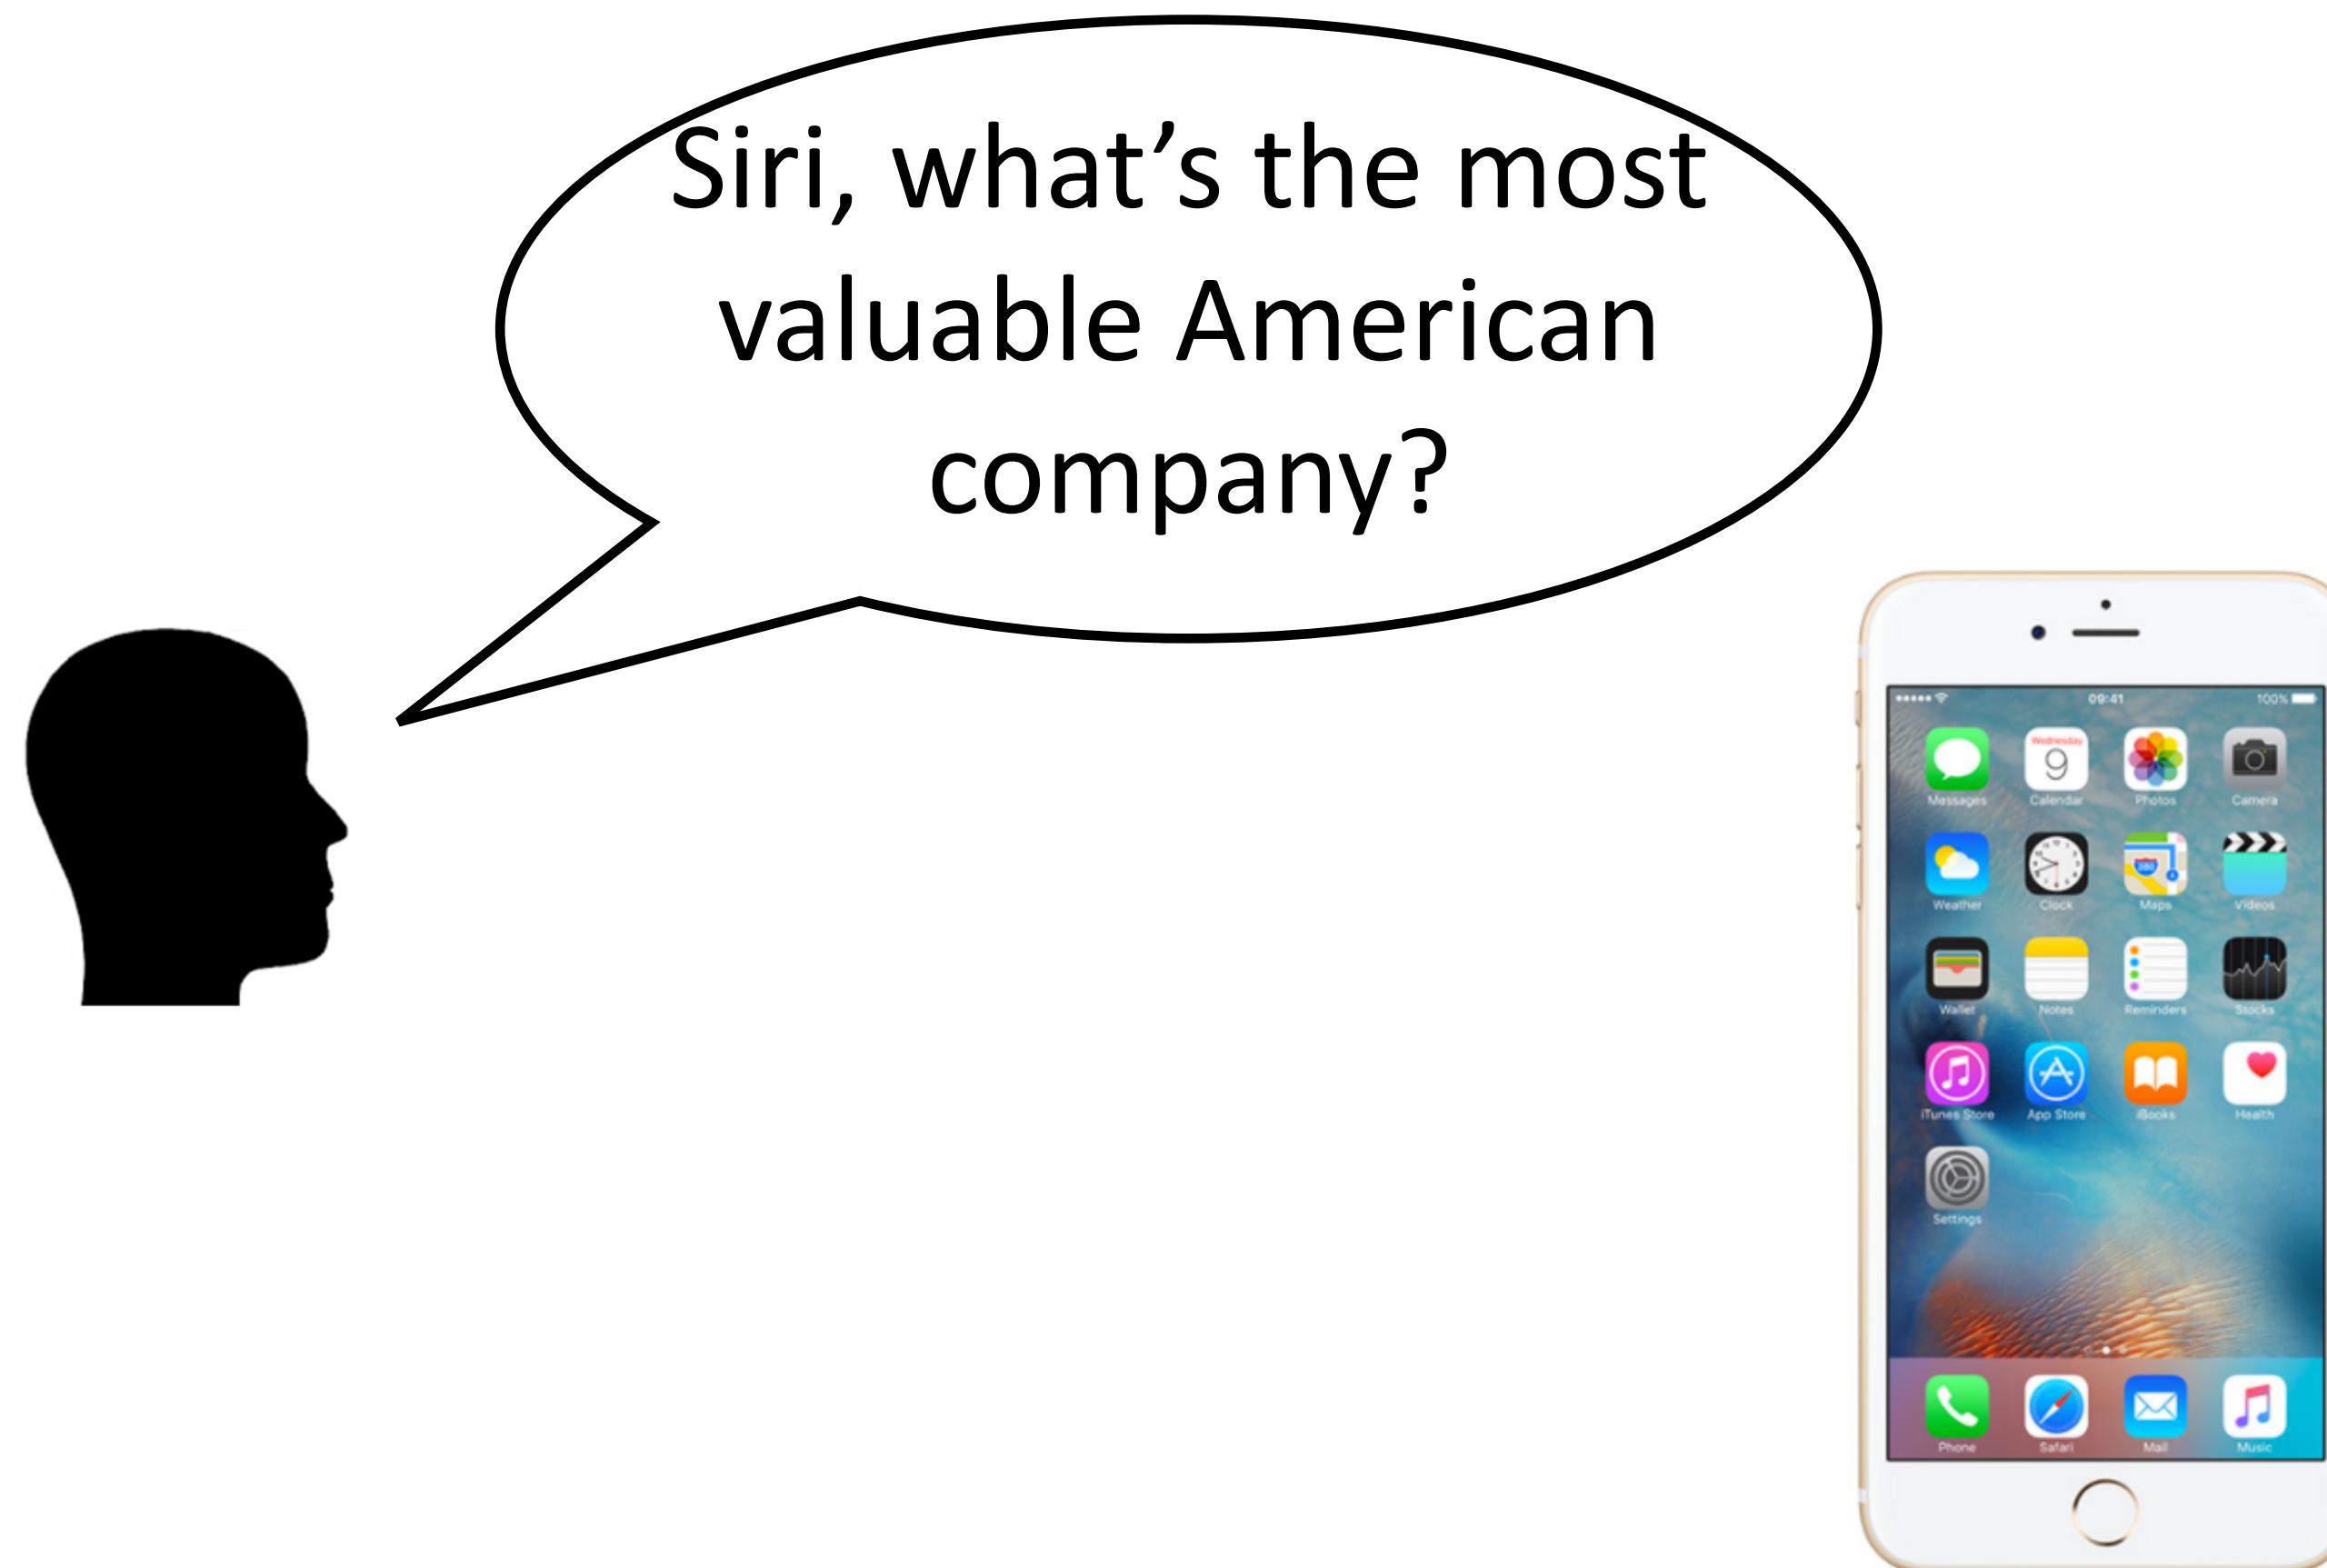
What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



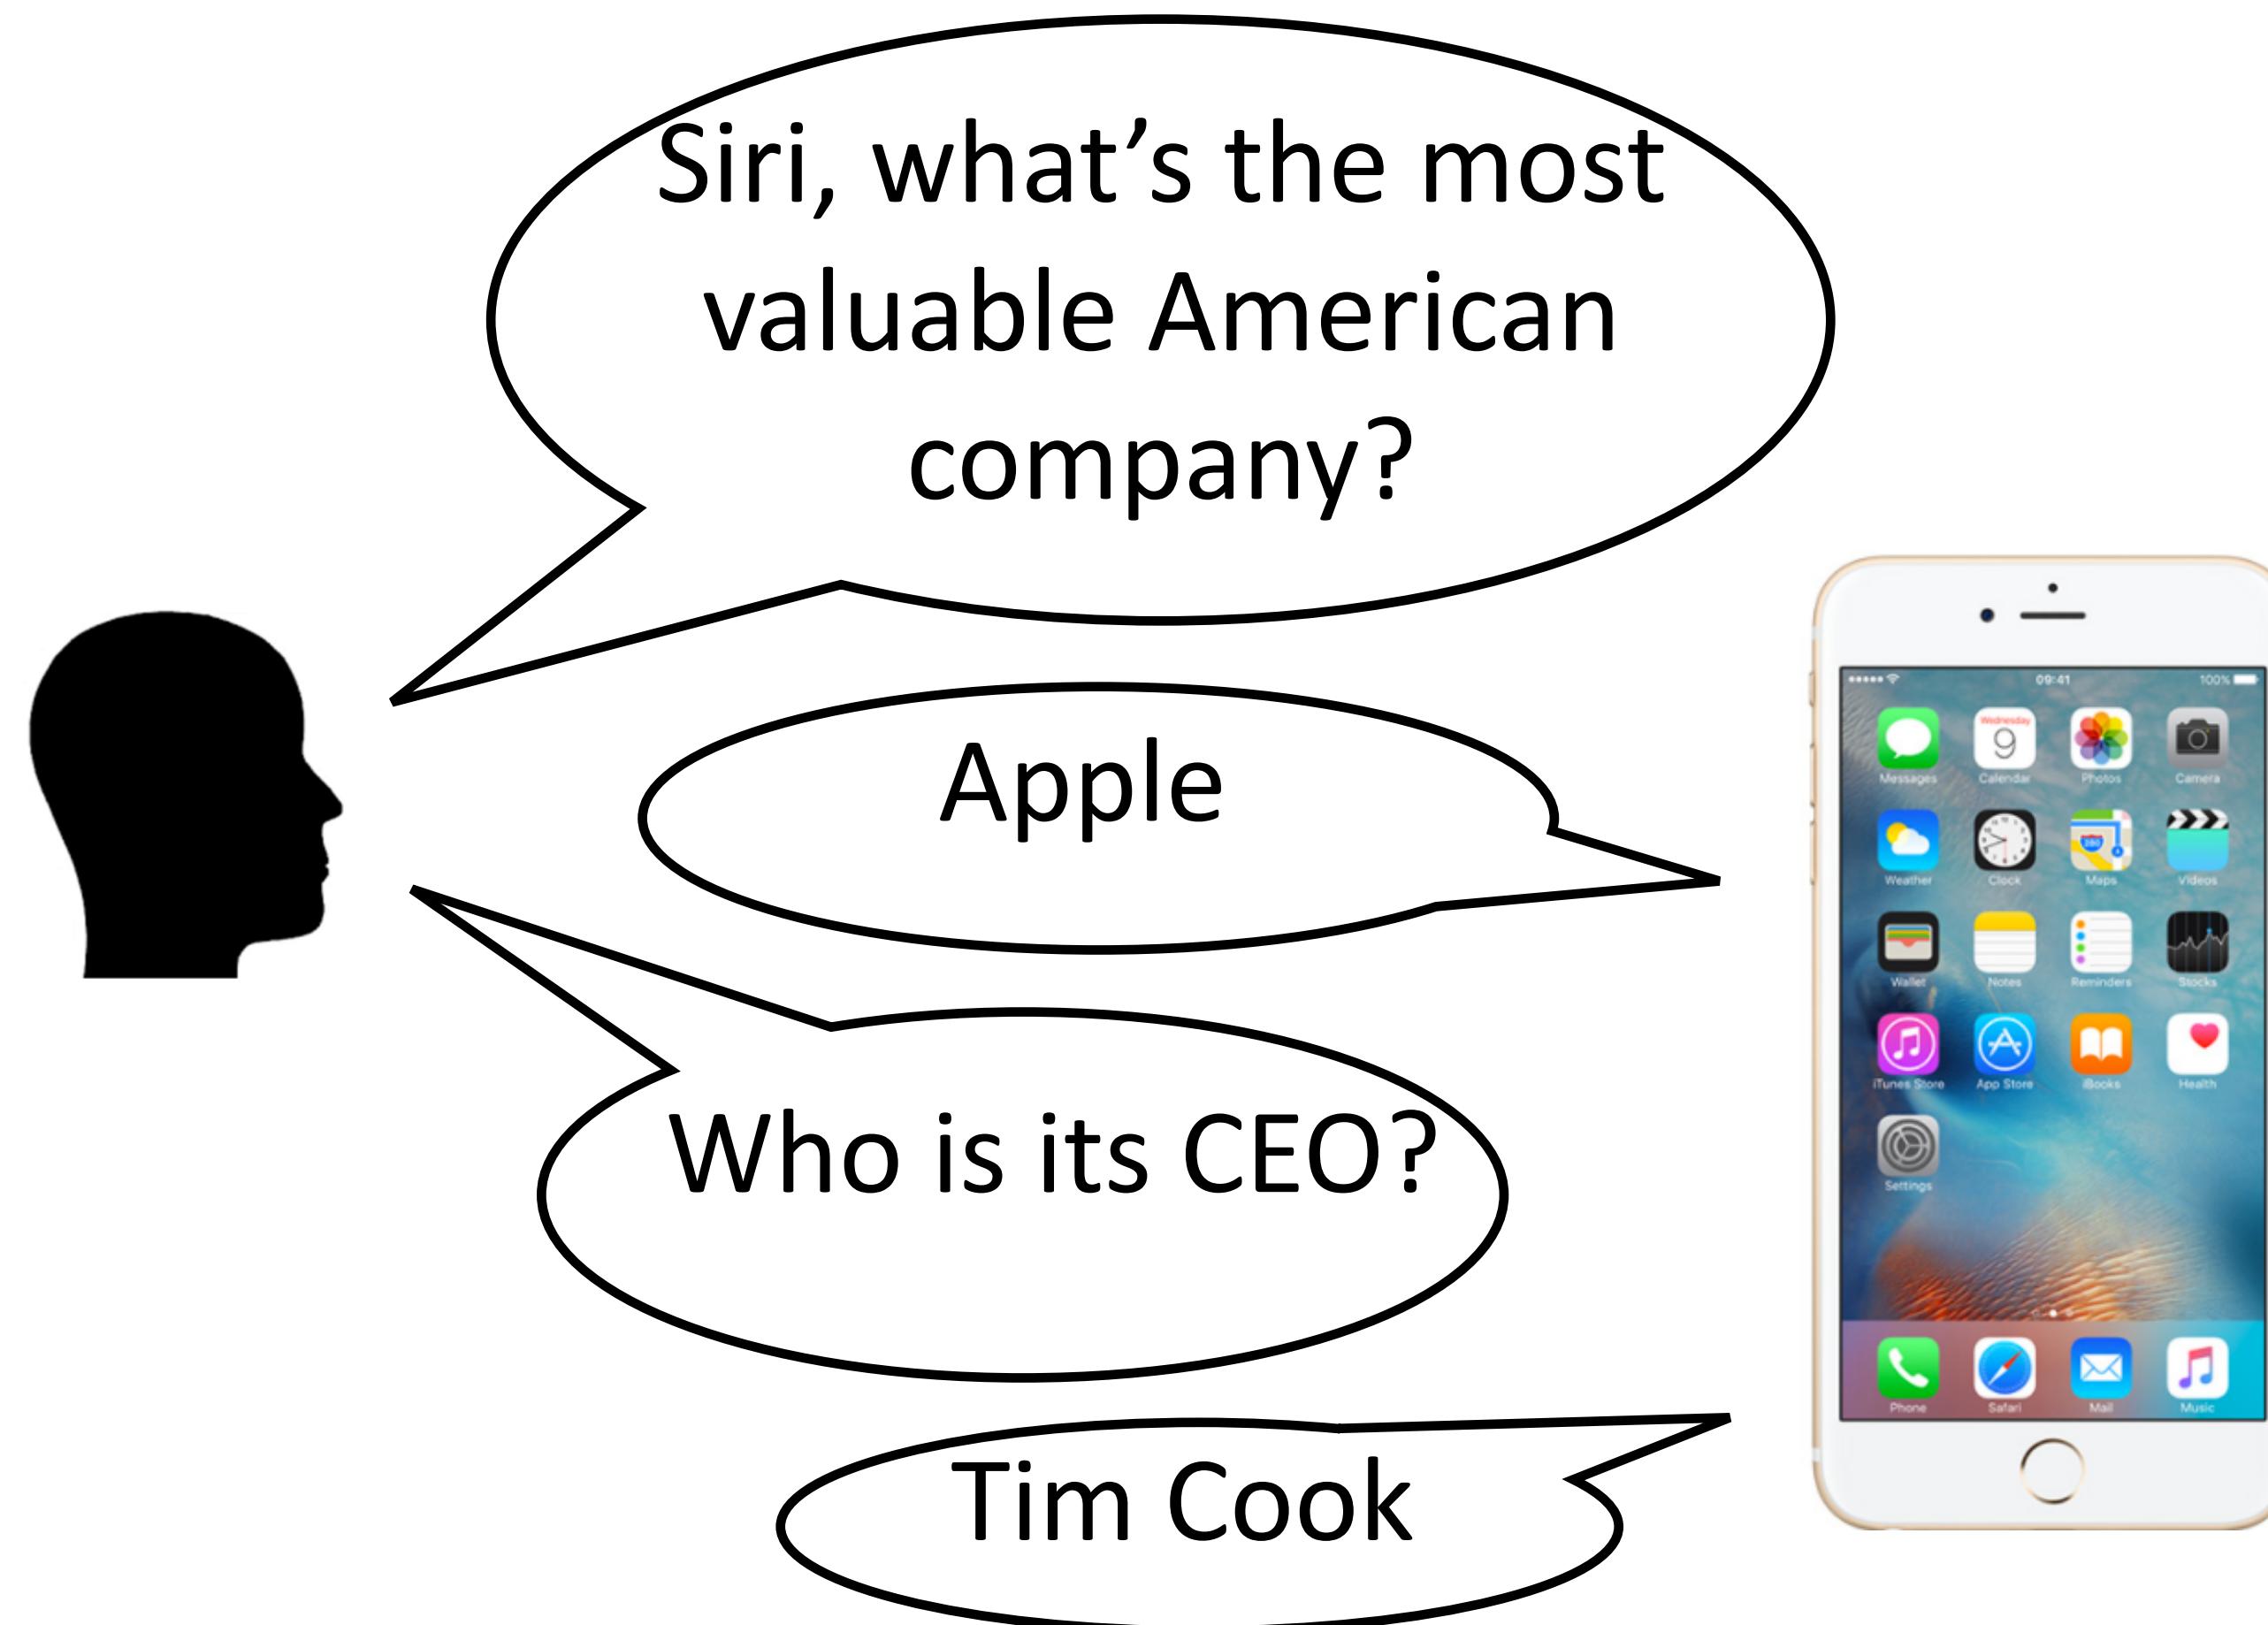
What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



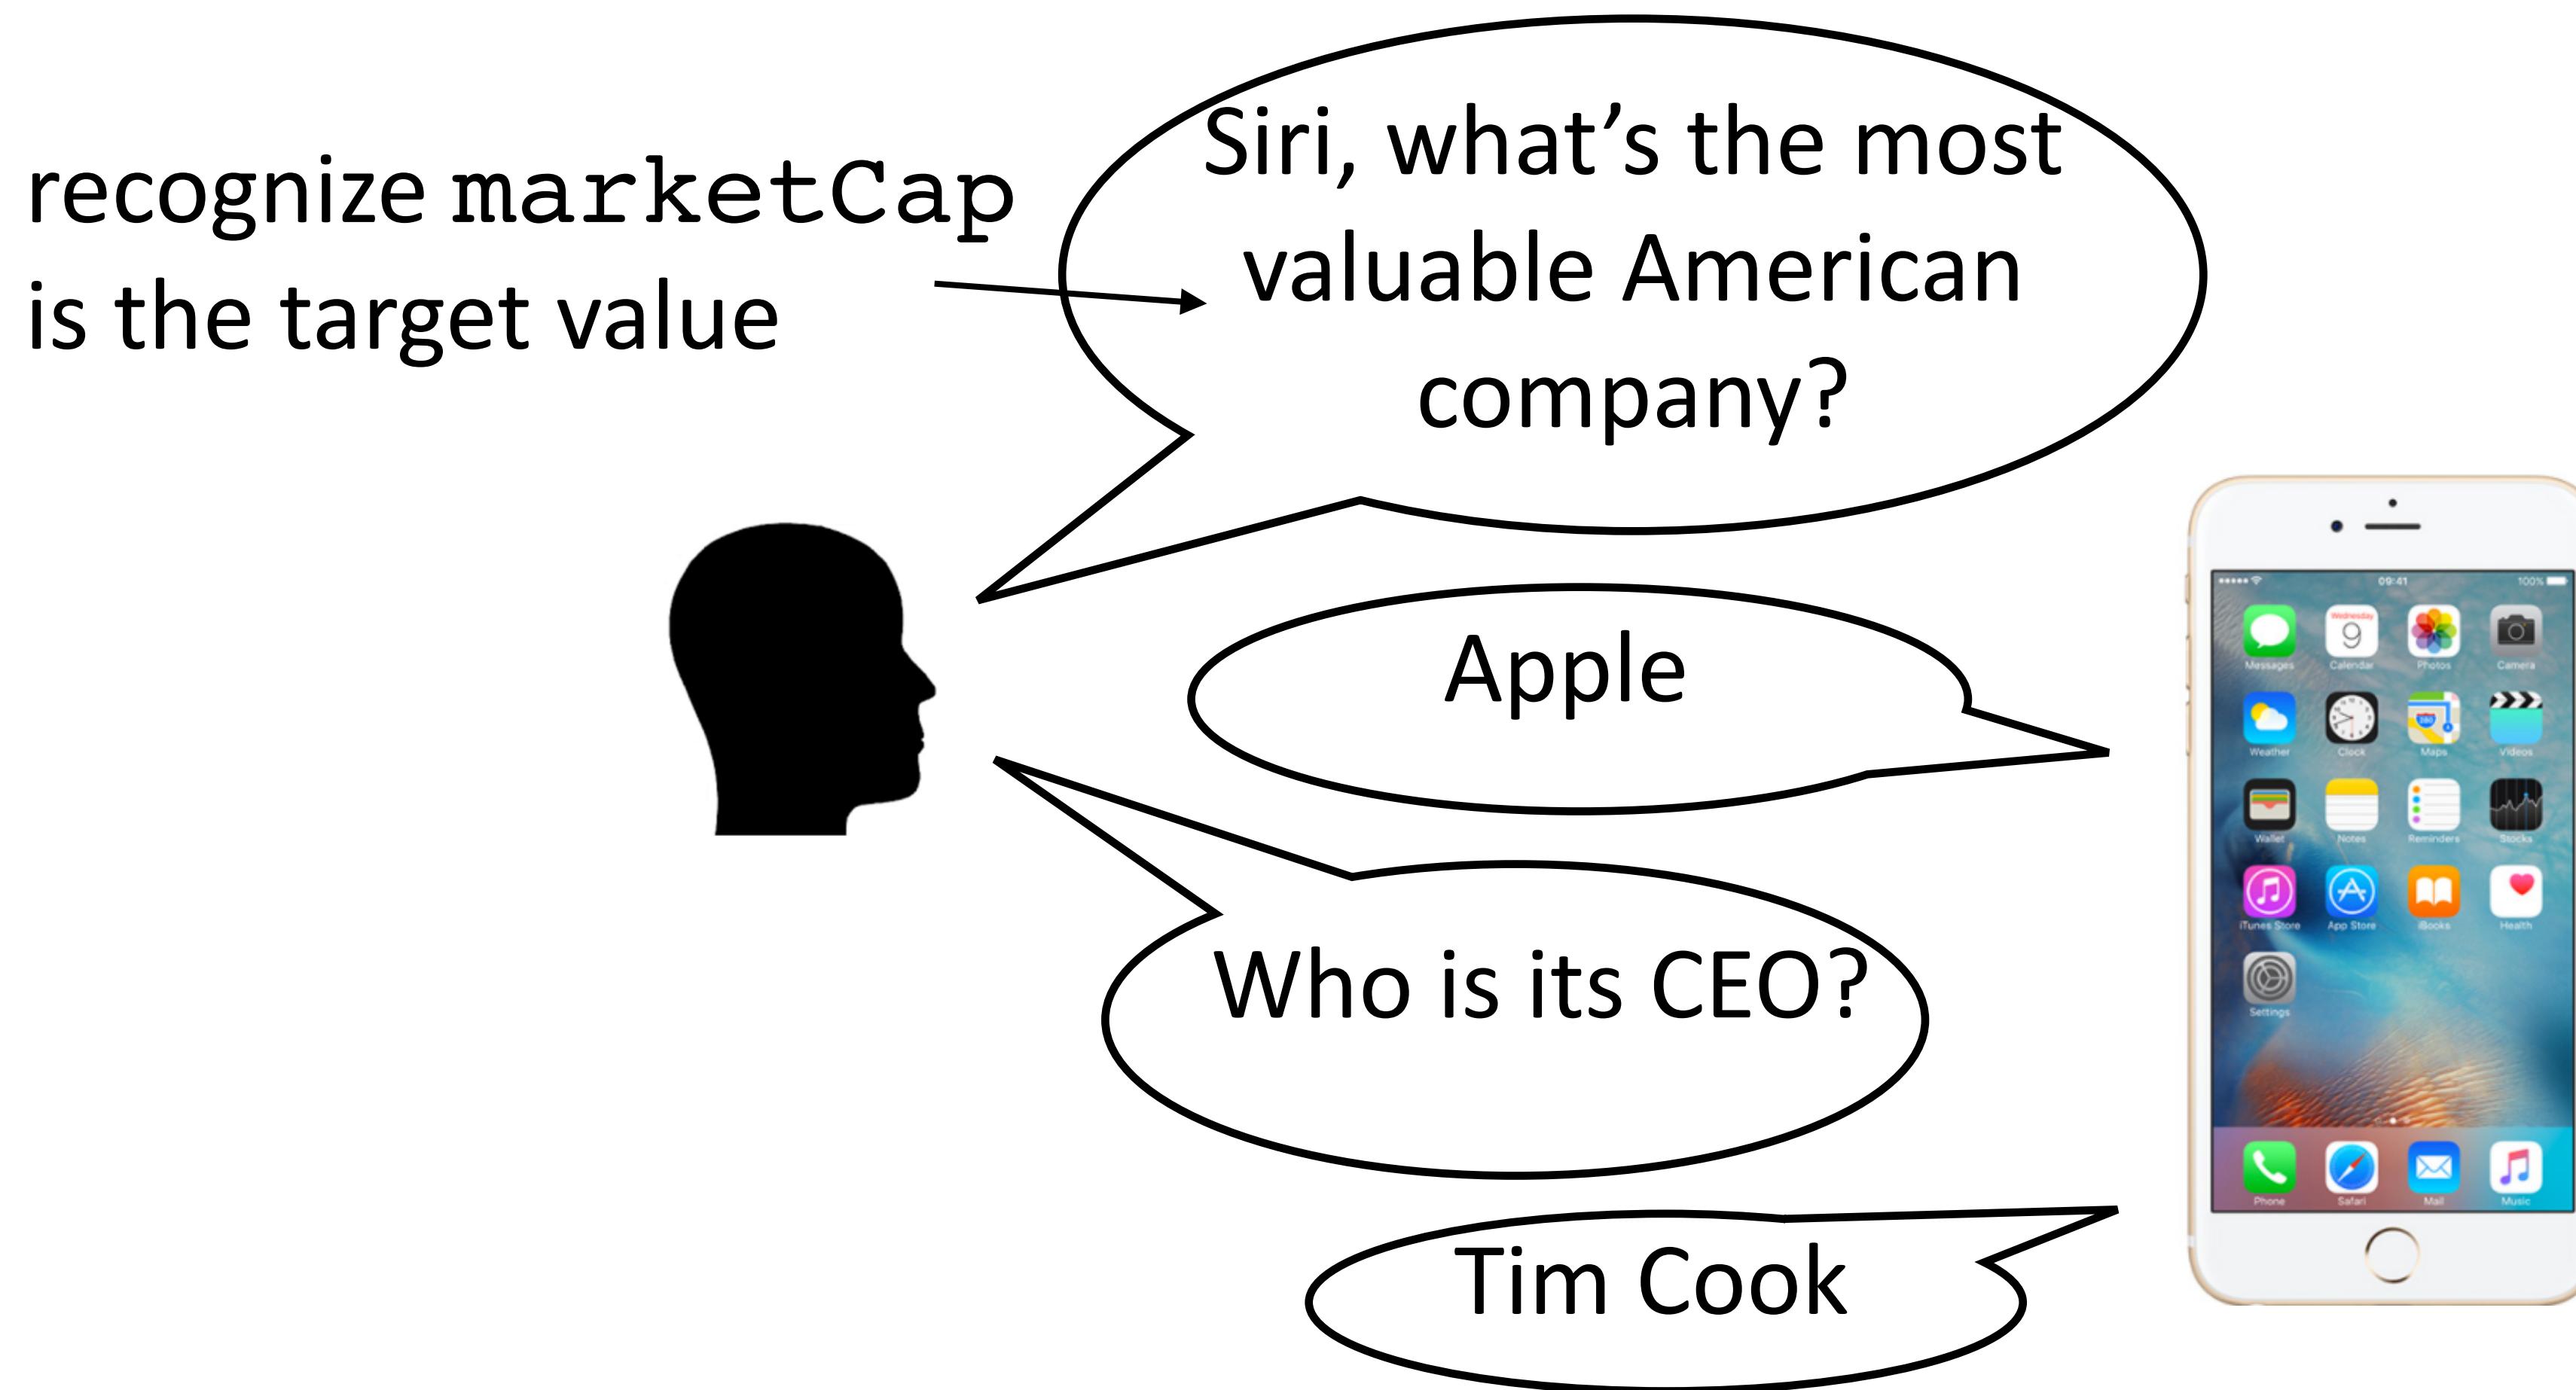
What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



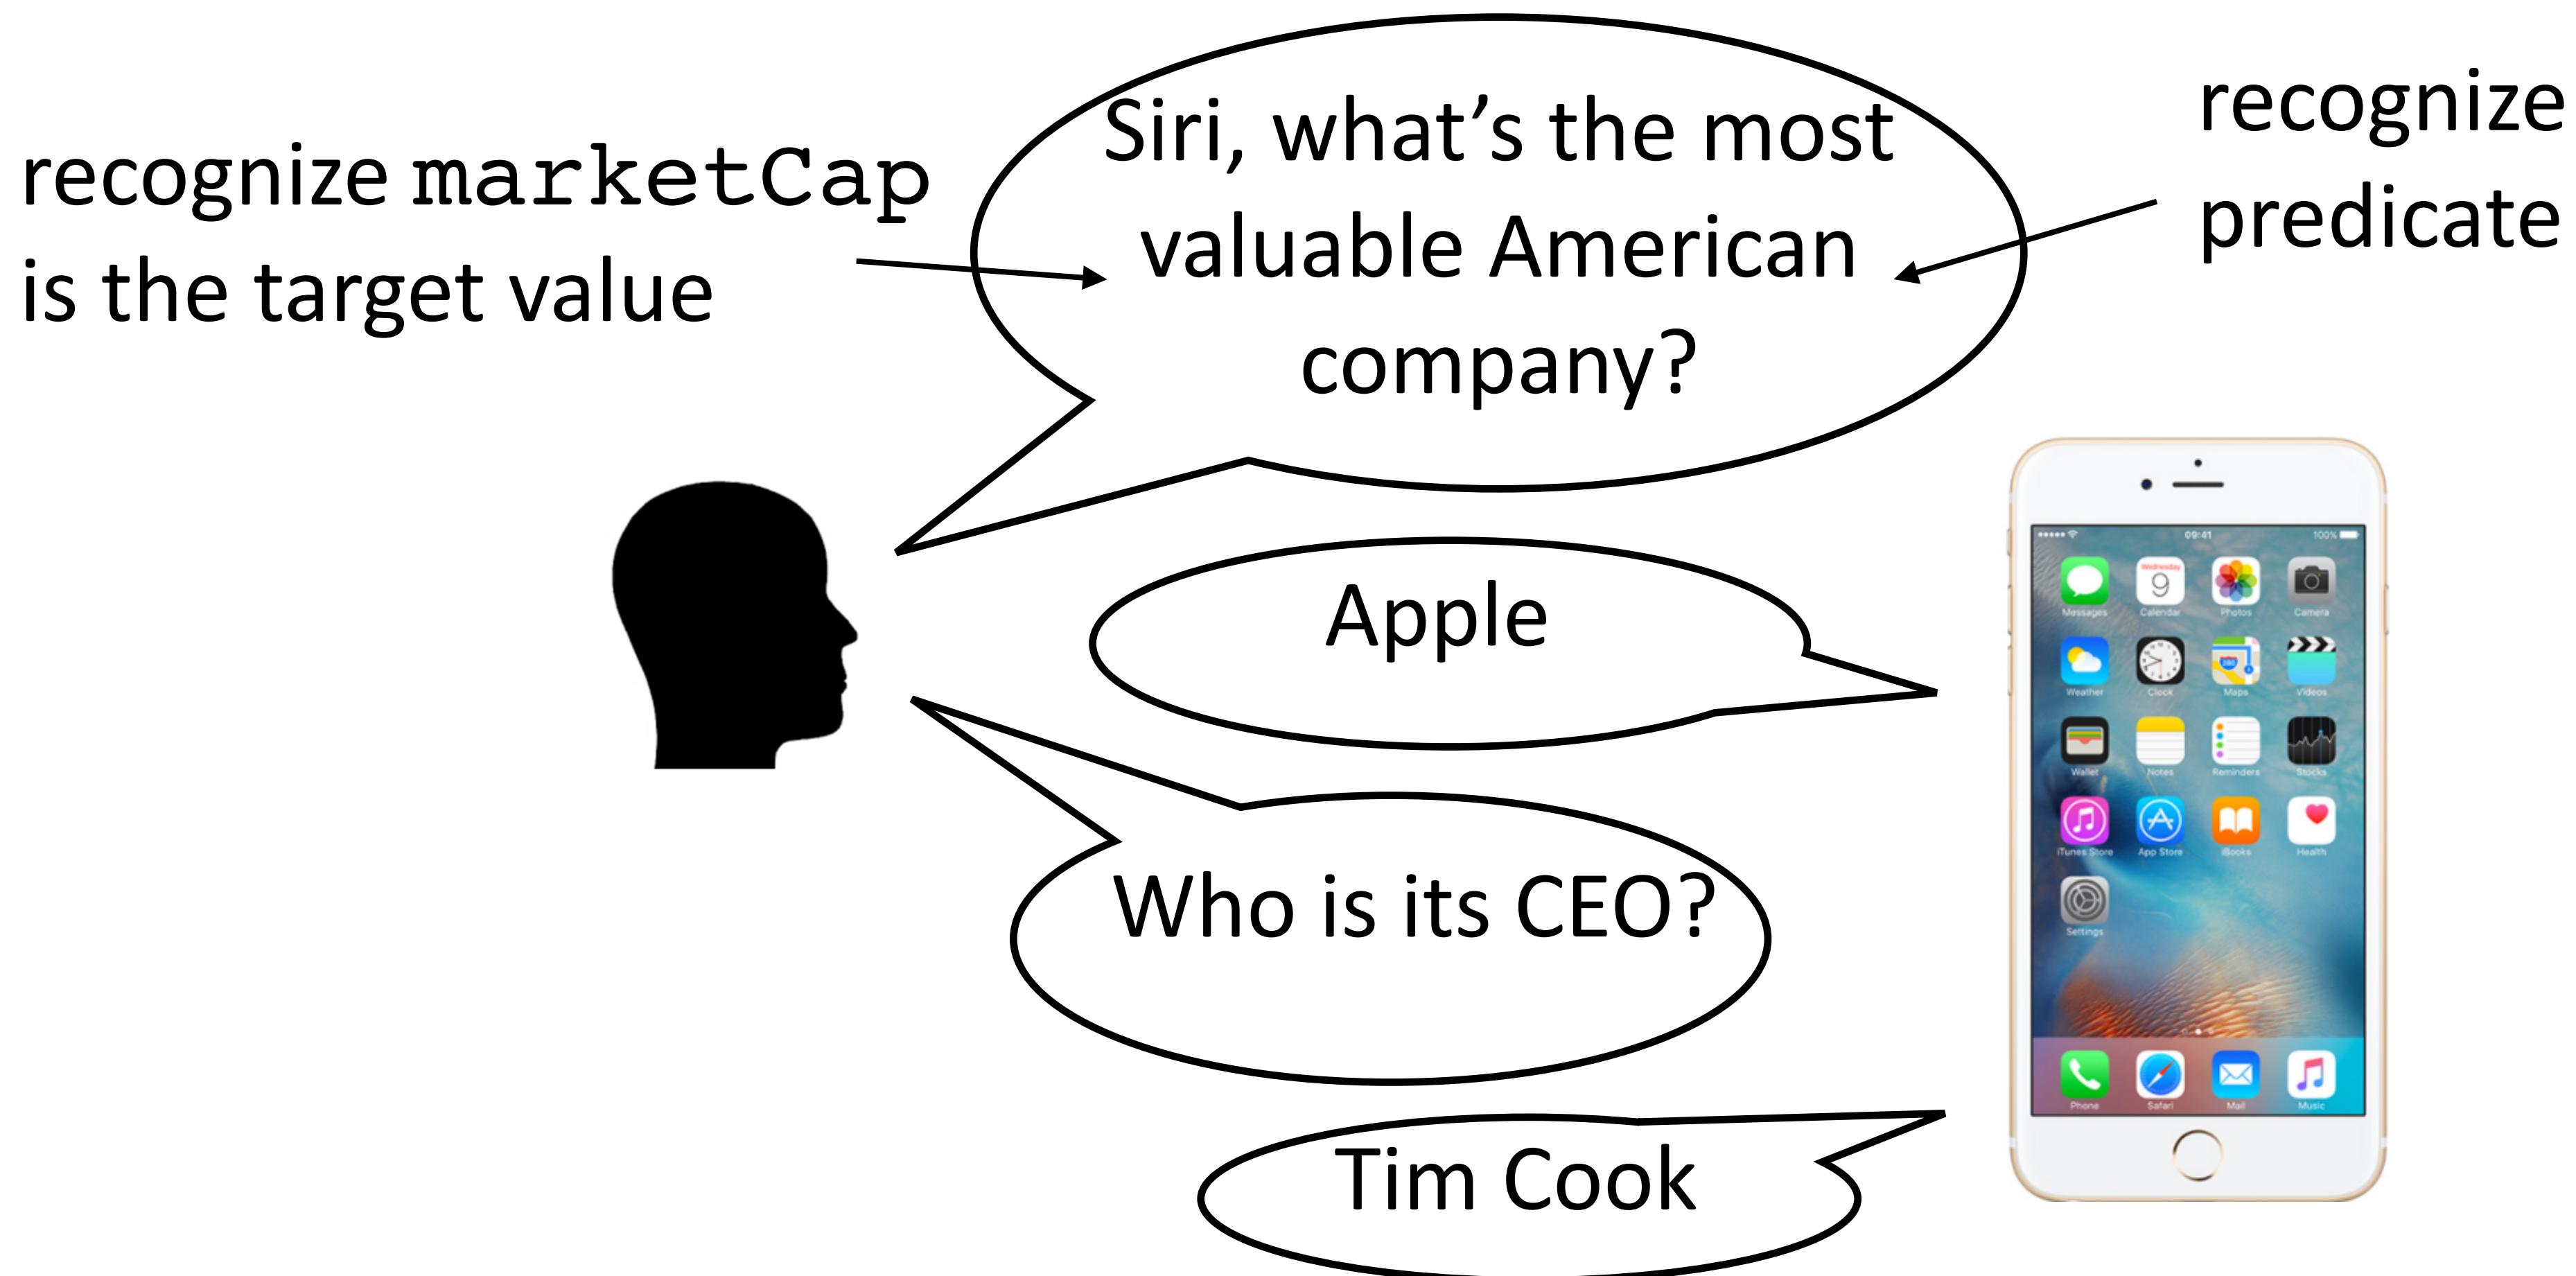
What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



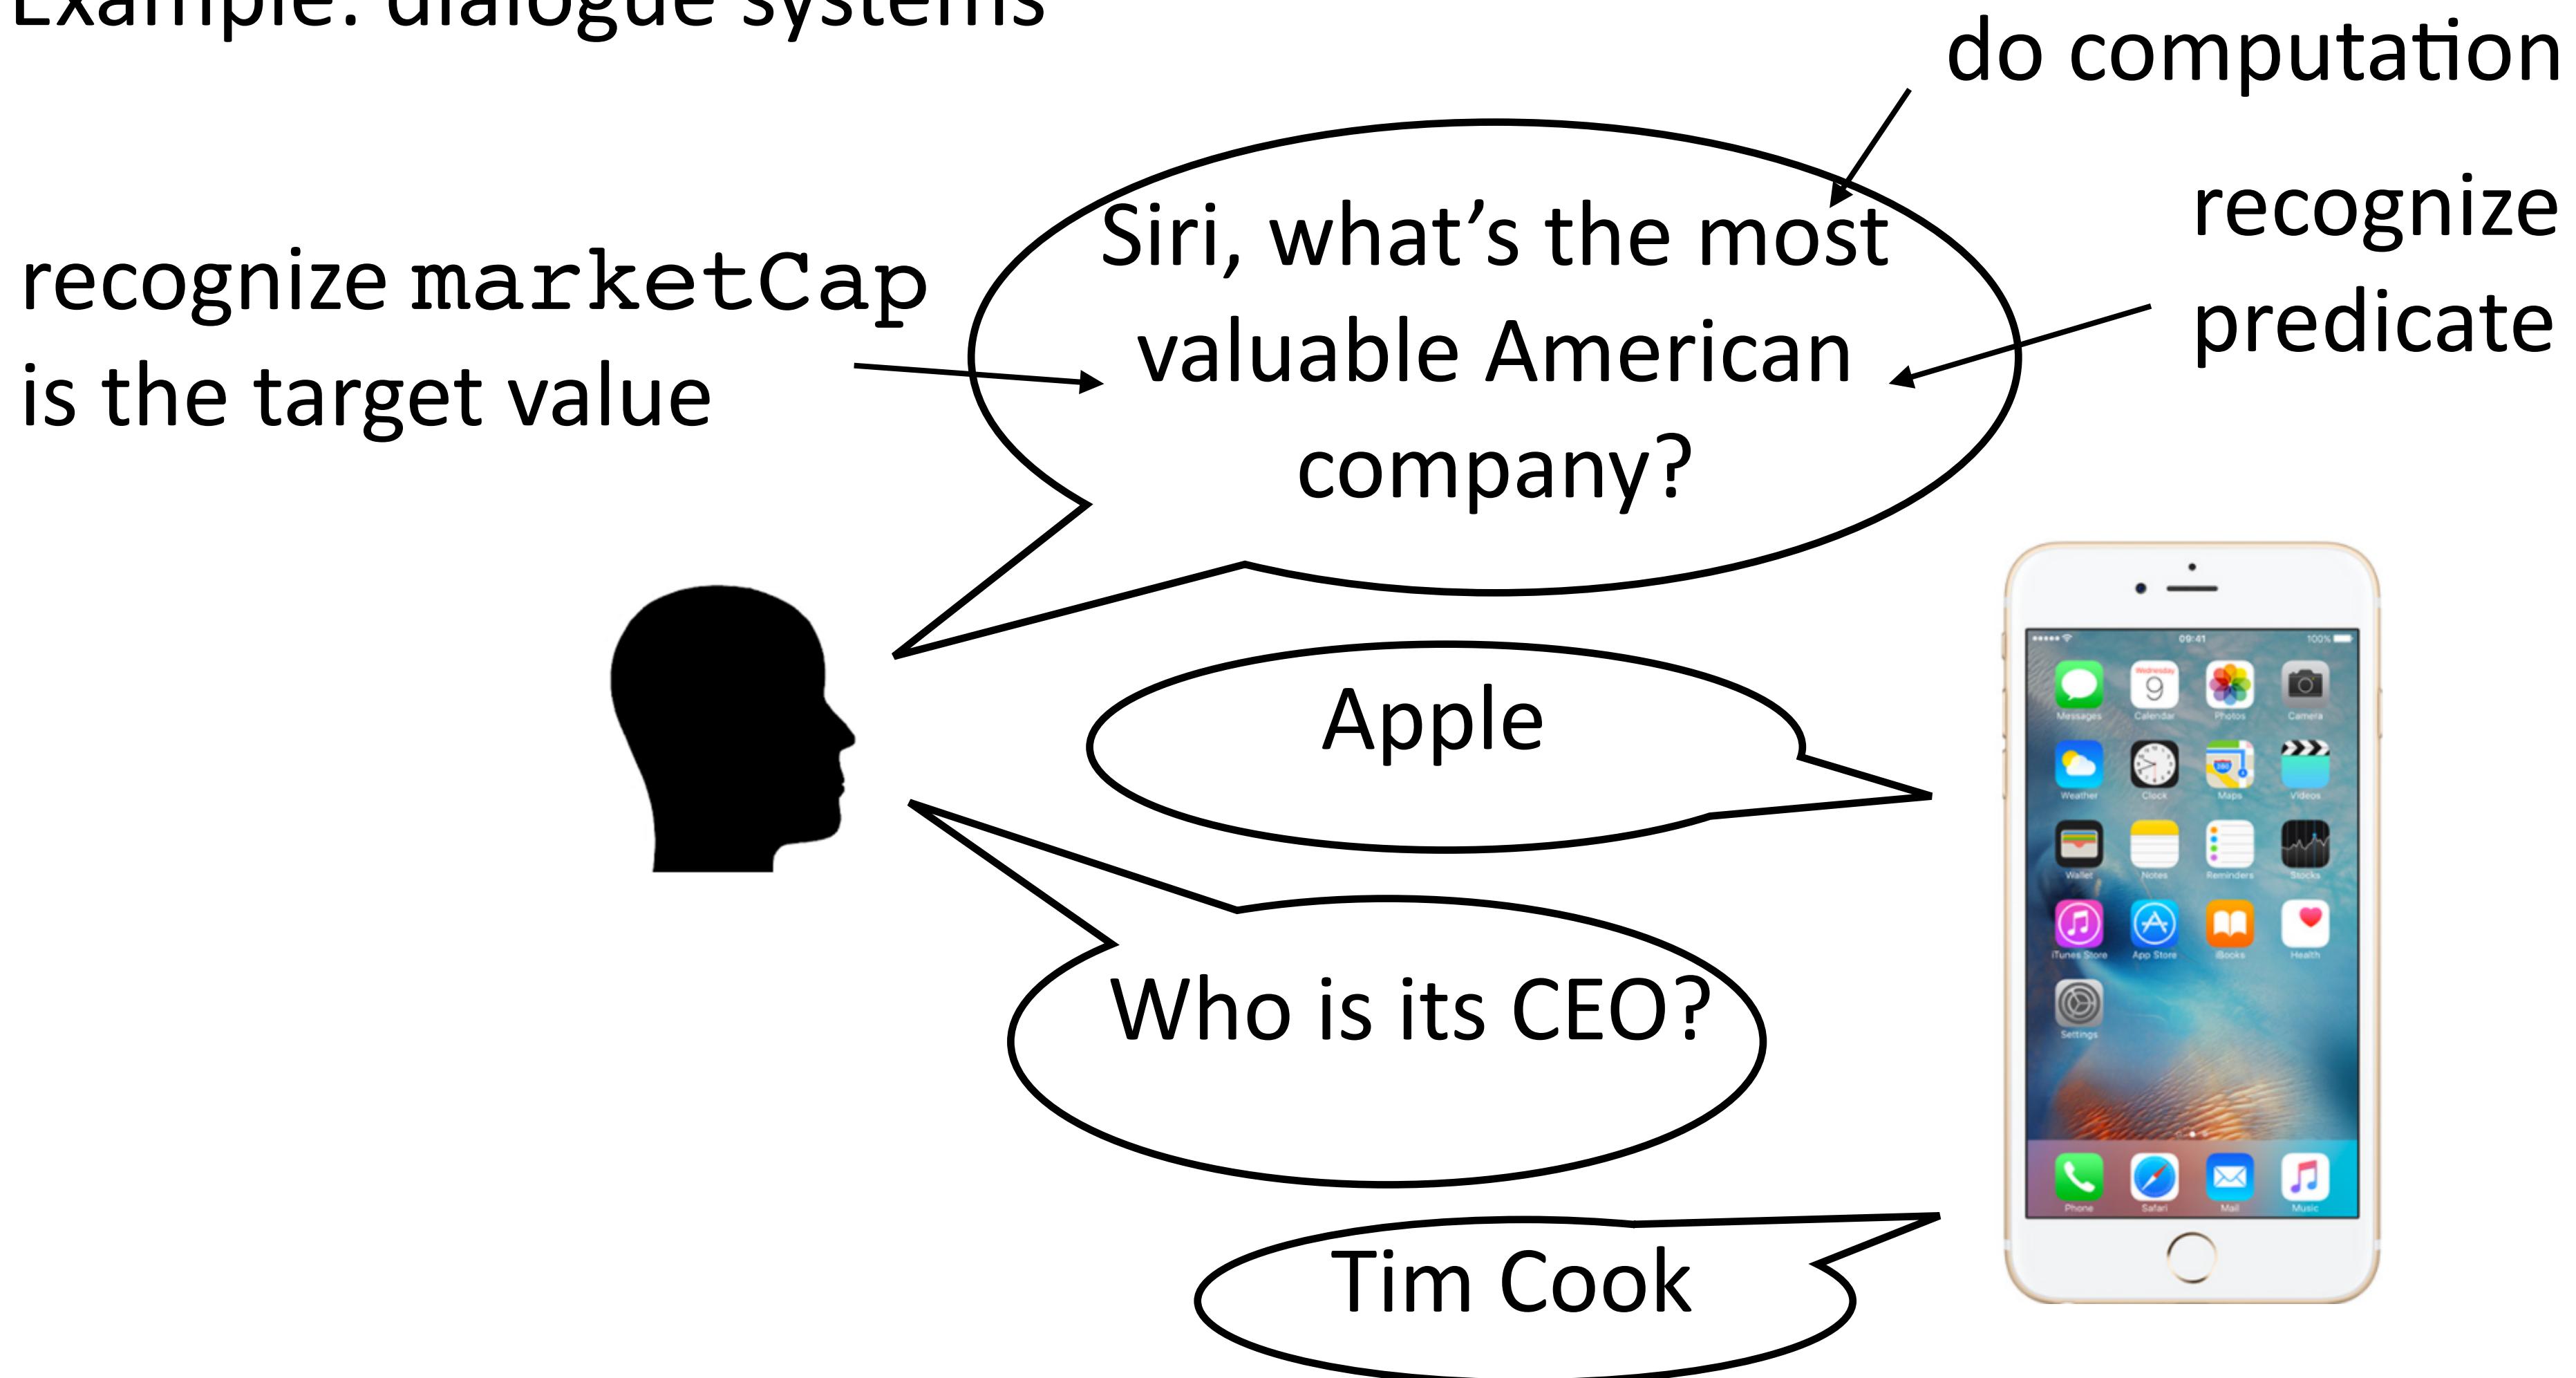
What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



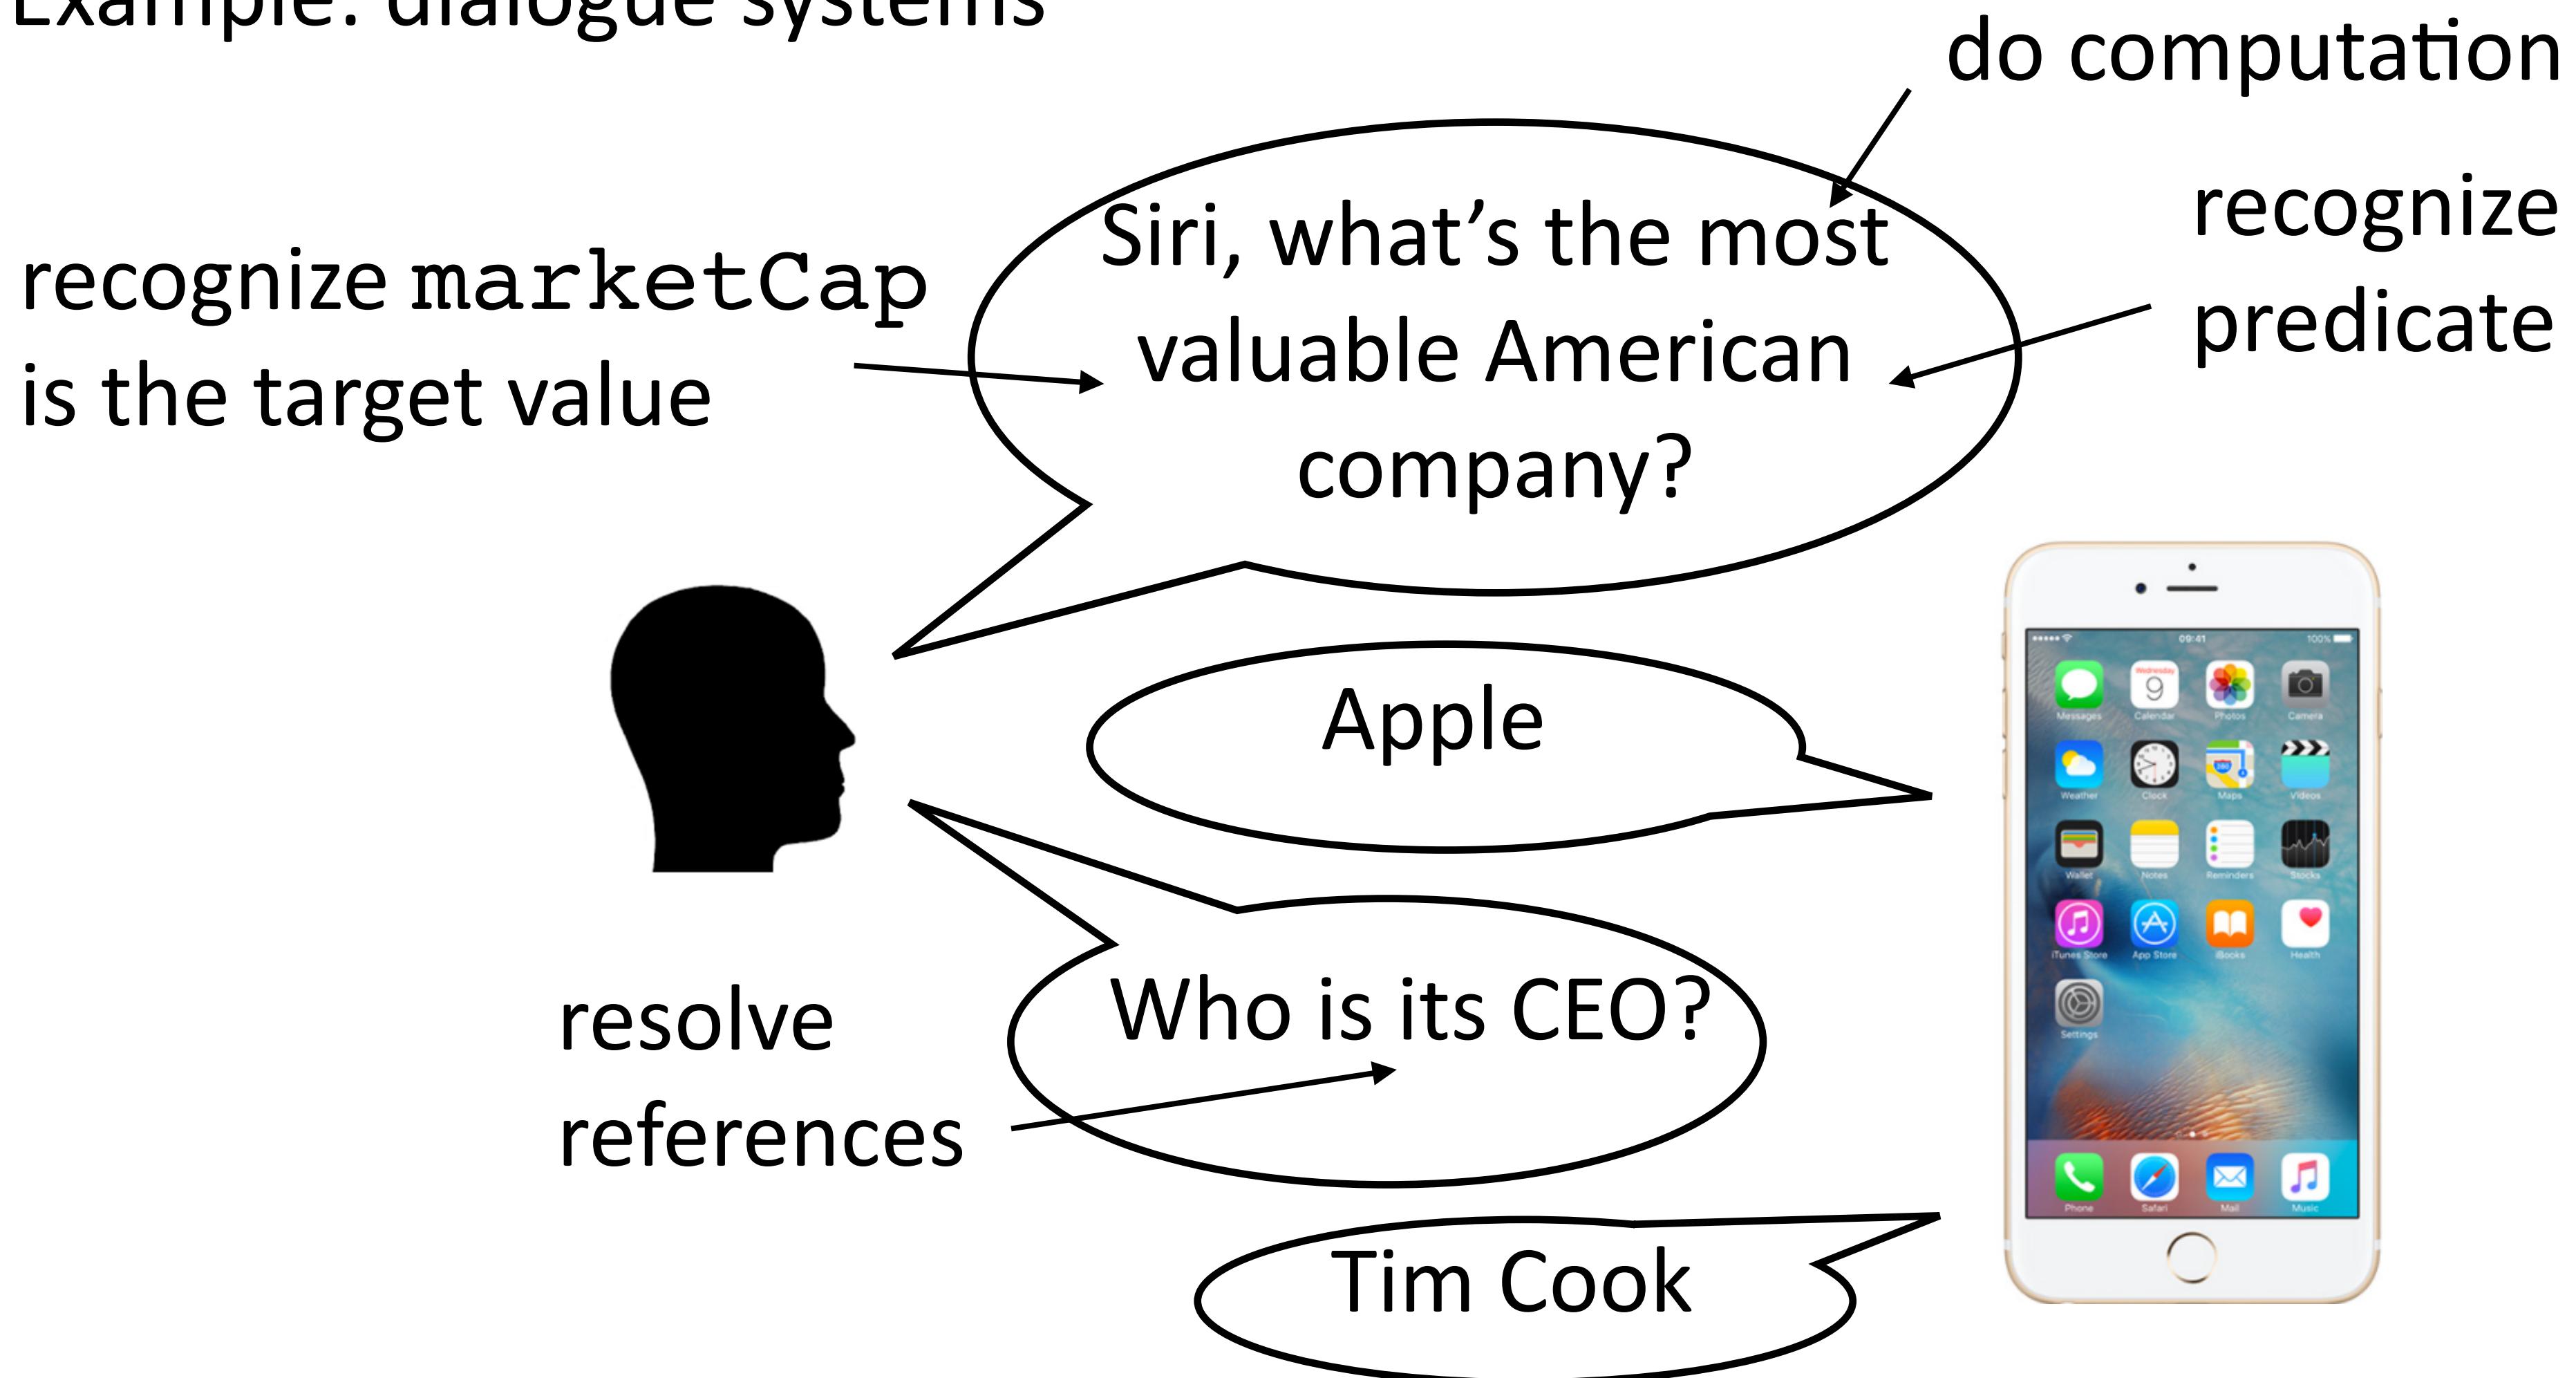
What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



What's the goal of NLP?

- ▶ Be able to solve problems that require deep understanding of text
- ▶ Example: dialogue systems



Automatic Summarization

Automatic Summarization

POLITICS

Google Critic Ousted From Think Tank Funded by the Tech Giant

WASHINGTON — In the hours after European antitrust regulators levied a record [\\$2.7 billion fine](#) against Google in late June, an influential Washington think tank learned what can happen when a tech giant that shapes public policy debates with its enormous wealth is criticized.

• • •

But not long after one of New America's scholars [posted a statement](#) on the think tank's website praising the European Union's penalty against Google, Mr. Schmidt, who had been chairman of New America until 2016, communicated his displeasure with the statement to the group's president, Anne-Marie Slaughter, according to the scholar.

• • •

Ms. Slaughter told Mr. Lynn that “the time has come for Open Markets and New America to part ways,” according to an email from Ms. Slaughter to Mr. Lynn. The email suggested that the entire Open Markets team — nearly 10 full-time employees and unpaid fellows — would be exiled from New America.

Automatic Summarization

POLITICS

Google Critic Ousted From Think Tank Funded by the Tech Giant

WASHINGTON — In the hours after European antitrust regulators levied a record [\\$2.7 billion fine](#) against Google in late June, an influential Washington think tank learned what can happen when a tech giant that shapes public policy debates with its enormous wealth is criticized.

• • •

But not long after one of New America's scholars [posted a statement](#) on the think tank's website praising the European Union's penalty against Google, Mr. Schmidt, who had been chairman of New America until 2016, communicated his displeasure with the statement to the group's president, Anne-Marie Slaughter, according to the scholar.

• • •

Ms. Slaughter told Mr. Lynn that “the time has come for Open Markets and New America to part ways,” according to an email from Ms. Slaughter to Mr. Lynn. The email suggested that the entire Open Markets team — nearly 10 full-time employees and unpaid fellows — would be exiled from New America.

One of New America's writers posted a statement critical of Google. Eric Schmidt, Google's CEO, was displeased.

The writer and his team were dismissed.

Automatic Summarization

POLITICS

Google Critic Ousted From Think Tank Funded by the Tech Giant

WASHINGTON — In the hours after European antitrust regulators levied a record [\\$2.7 billion fine](#) against Google in late June, an influential Washington think tank learned what can happen when a tech giant that shapes public policy debates with its enormous wealth is criticized.

• • •

But not long after one of New America's scholars [posted a statement](#) on the think tank's website praising the European Union's penalty against Google, Mr. Schmidt, who had been chairman of New America until 2016, communicated his displeasure with the statement to the group's president,

Anne-Marie Slaughter, according to the scholar.

• • •

Ms. Slaughter told Mr. Lynn that “the time has come for Open Markets and New America to part ways,” according to an email from Ms. Slaughter to Mr. Lynn. The email suggested that the entire Open Markets team — nearly 10 full-time employees and unpaid fellows — would be exiled from New America.

compress
text

One of New America's writers posted a statement critical of Google. Eric Schmidt, Google's CEO, was displeased.

The writer and his team were dismissed.

Automatic Summarization

POLITICS

Google Critic Ousted From Think Tank Funded by the Tech Giant

WASHINGTON — In the hours after European antitrust regulators levied a record [\\$2.7 billion fine](#) against Google in late June, an influential Washington think tank learned what can happen when a tech giant that shapes public policy debates with its enormous wealth is criticized.

•••

But not long after one of New America's scholars [posted a statement](#) on the think tank's website praising the European Union's penalty against Google, Mr. Schmidt, who had been chairman of New America until 2016, communicated his displeasure with the statement to the group's president,

Anne-Marie Slaughter, according to the scholar.

•••

Ms. Slaughter told Mr. Lynn that “the time has come for Open Markets and New America to part ways,” according to an email from Ms. Slaughter to Mr. Lynn. The email suggested that the entire Open Markets team — nearly 10 full-time employees and unpaid fellows — would be exiled from New America.

compress
text

provide missing
context

One of New America's writers posted a statement critical of Google. Eric Schmidt, [Google's CEO](#), was displeased.

The writer and his team were dismissed.

Automatic Summarization

POLITICS

Google Critic Ousted From Think Tank Funded by the Tech Giant

WASHINGTON — In the hours after European antitrust regulators levied a record [\\$2.7 billion fine](#) against Google in late June, an influential Washington think tank learned what can happen when a tech giant that shapes public policy debates with its enormous wealth is criticized.

•••

But not long after one of New America's scholars [posted a statement](#) on the think tank's website praising the European Union's penalty against Google, Mr. Schmidt, who had been chairman of New America until 2016, communicated his displeasure with the statement to the group's president,

Anne-Marie Slaughter, according to the scholar.

•••

Ms. Slaughter told Mr. Lynn that “the time has come for Open Markets and New America to part ways,” according to an email from Ms. Slaughter to Mr. Lynn. The email suggested that the entire Open Markets team — nearly 10 full-time employees and unpaid fellows — would be [exiled](#) from New America.

compress
text

provide missing
context

One of New America's writers posted a statement critical of Google. Eric Schmidt, [Google's CEO](#), was displeased.

The writer and his team were dismissed.

paraphrase to provide clarity

Machine Translation



THE WALL STREET JOURNAL.

美国众议院议长选举大戏落幕，
共和党议员重点转向支出及中国
问题

7 小时前

Machine Translation



THE WALL STREET JOURNAL.

美国众议院议长选举大戏落幕，
共和党议员重点转向支出及中国
问题

7 小时前

CHINESE (SIMPLIFIED) - DETECTED

CHINESE (SIMPLIFIED)

HINDI

FRE ▾



ENGLISH

SPANISH

ARABIC ▾

美国众议院议长选举大戏落幕，共和党议员重点转向支出及中国问题



U.S. House speaker race ends as Republican lawmakers focus on spending, China



Machine Translation



THE WALL STREET JOURNAL.

美国众议院议长选举大戏落幕，
共和党议员重点转向支出及中国
问题

7小时前



CHINESE (SIMPLIFIED) - DETECTED

CHINESE (SIMPLIFIED)

HINDI

FRE ▾



ENGLISH

SPANISH

ARABIC ▾

美国众议院议长选举大戏落幕，共和党议员重点转向支出及中国问题



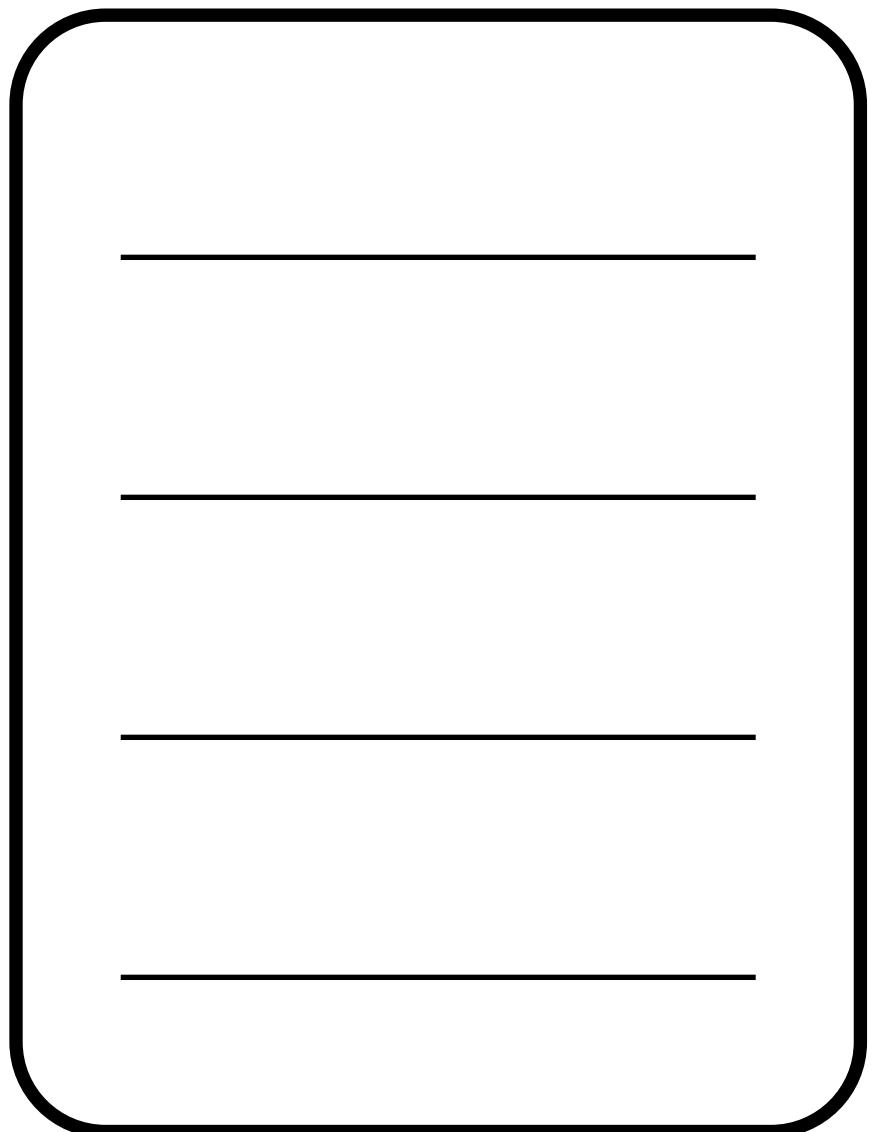
U.S. House speaker race ends as Republican lawmakers focus on spending, China



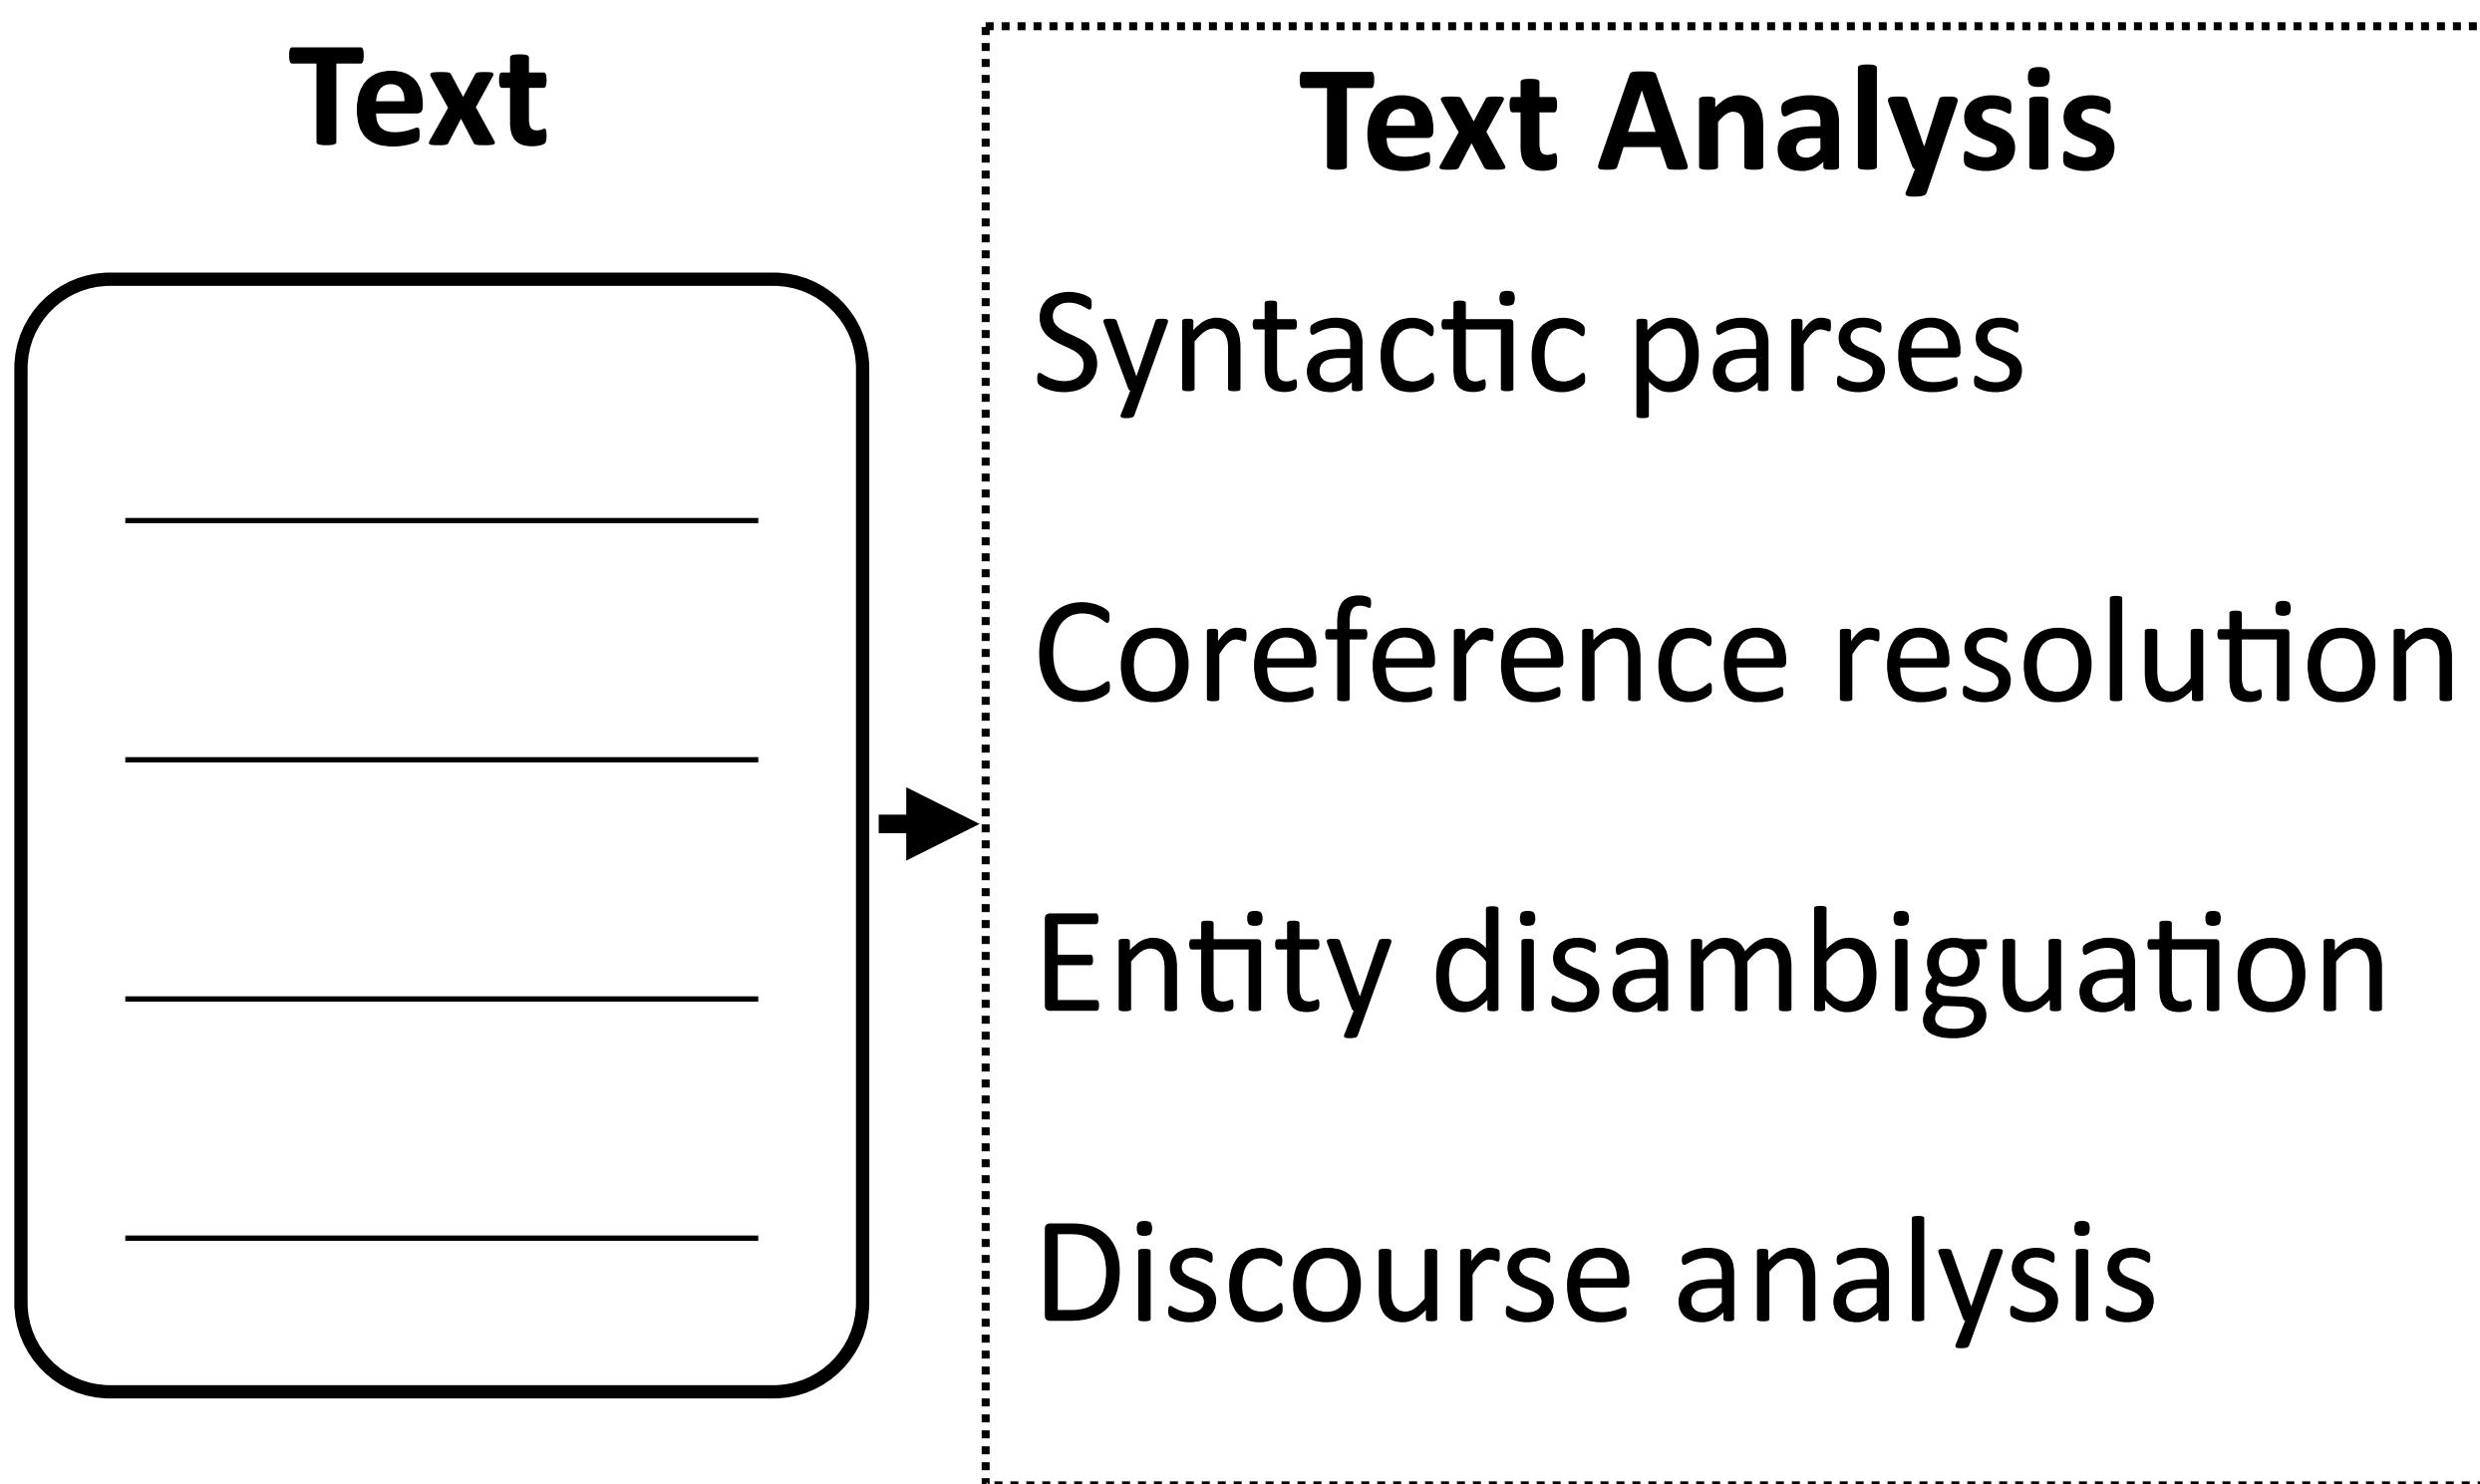
NLP Analysis Pipeline

NLP Analysis Pipeline

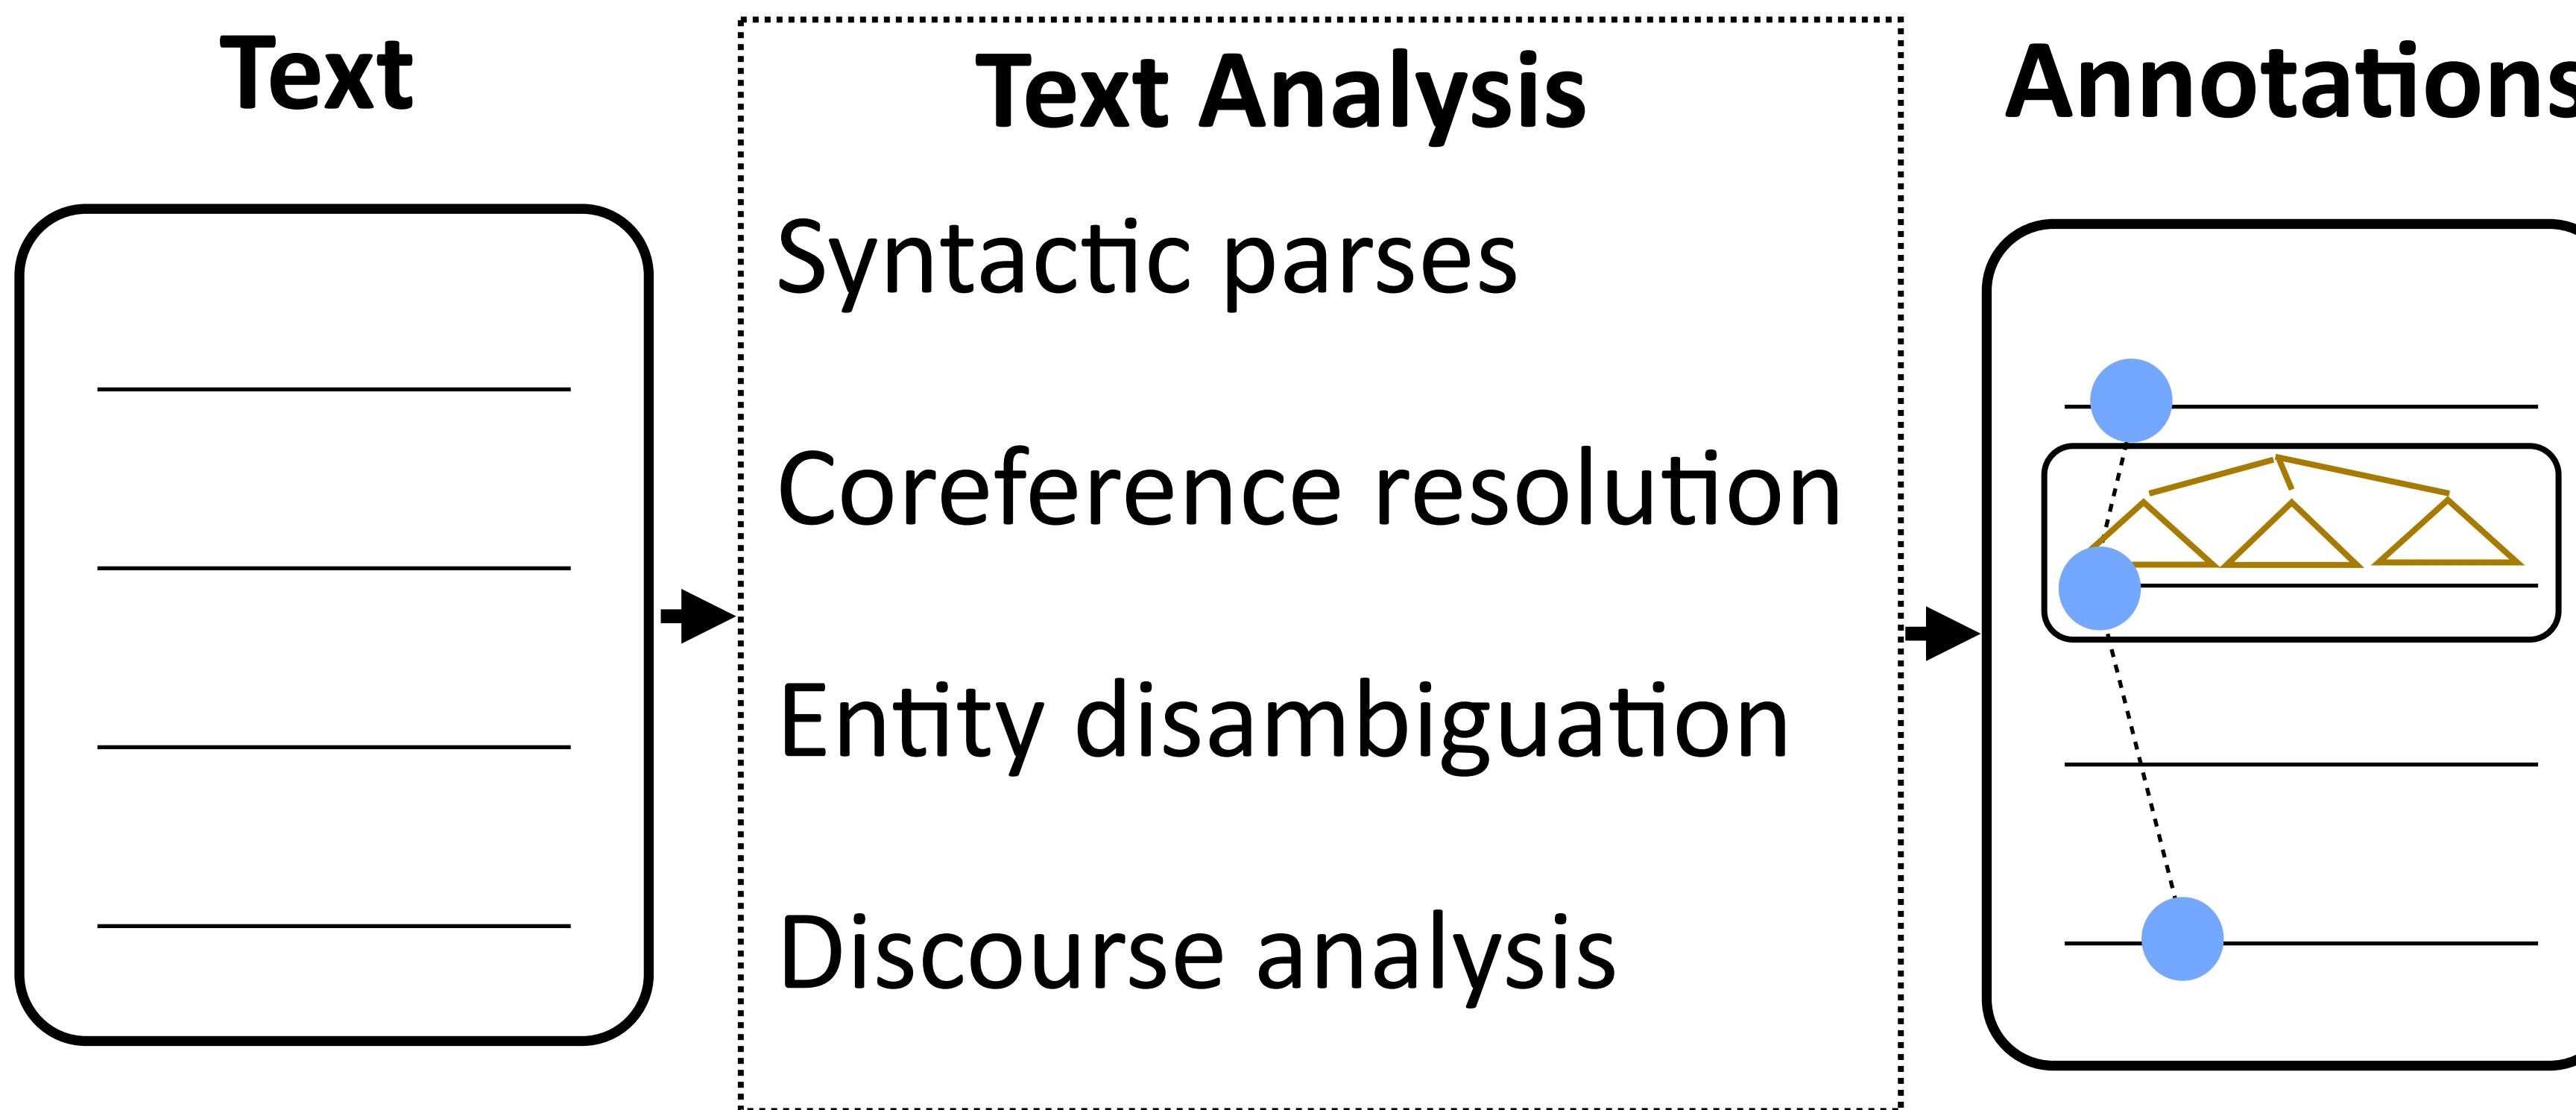
Text



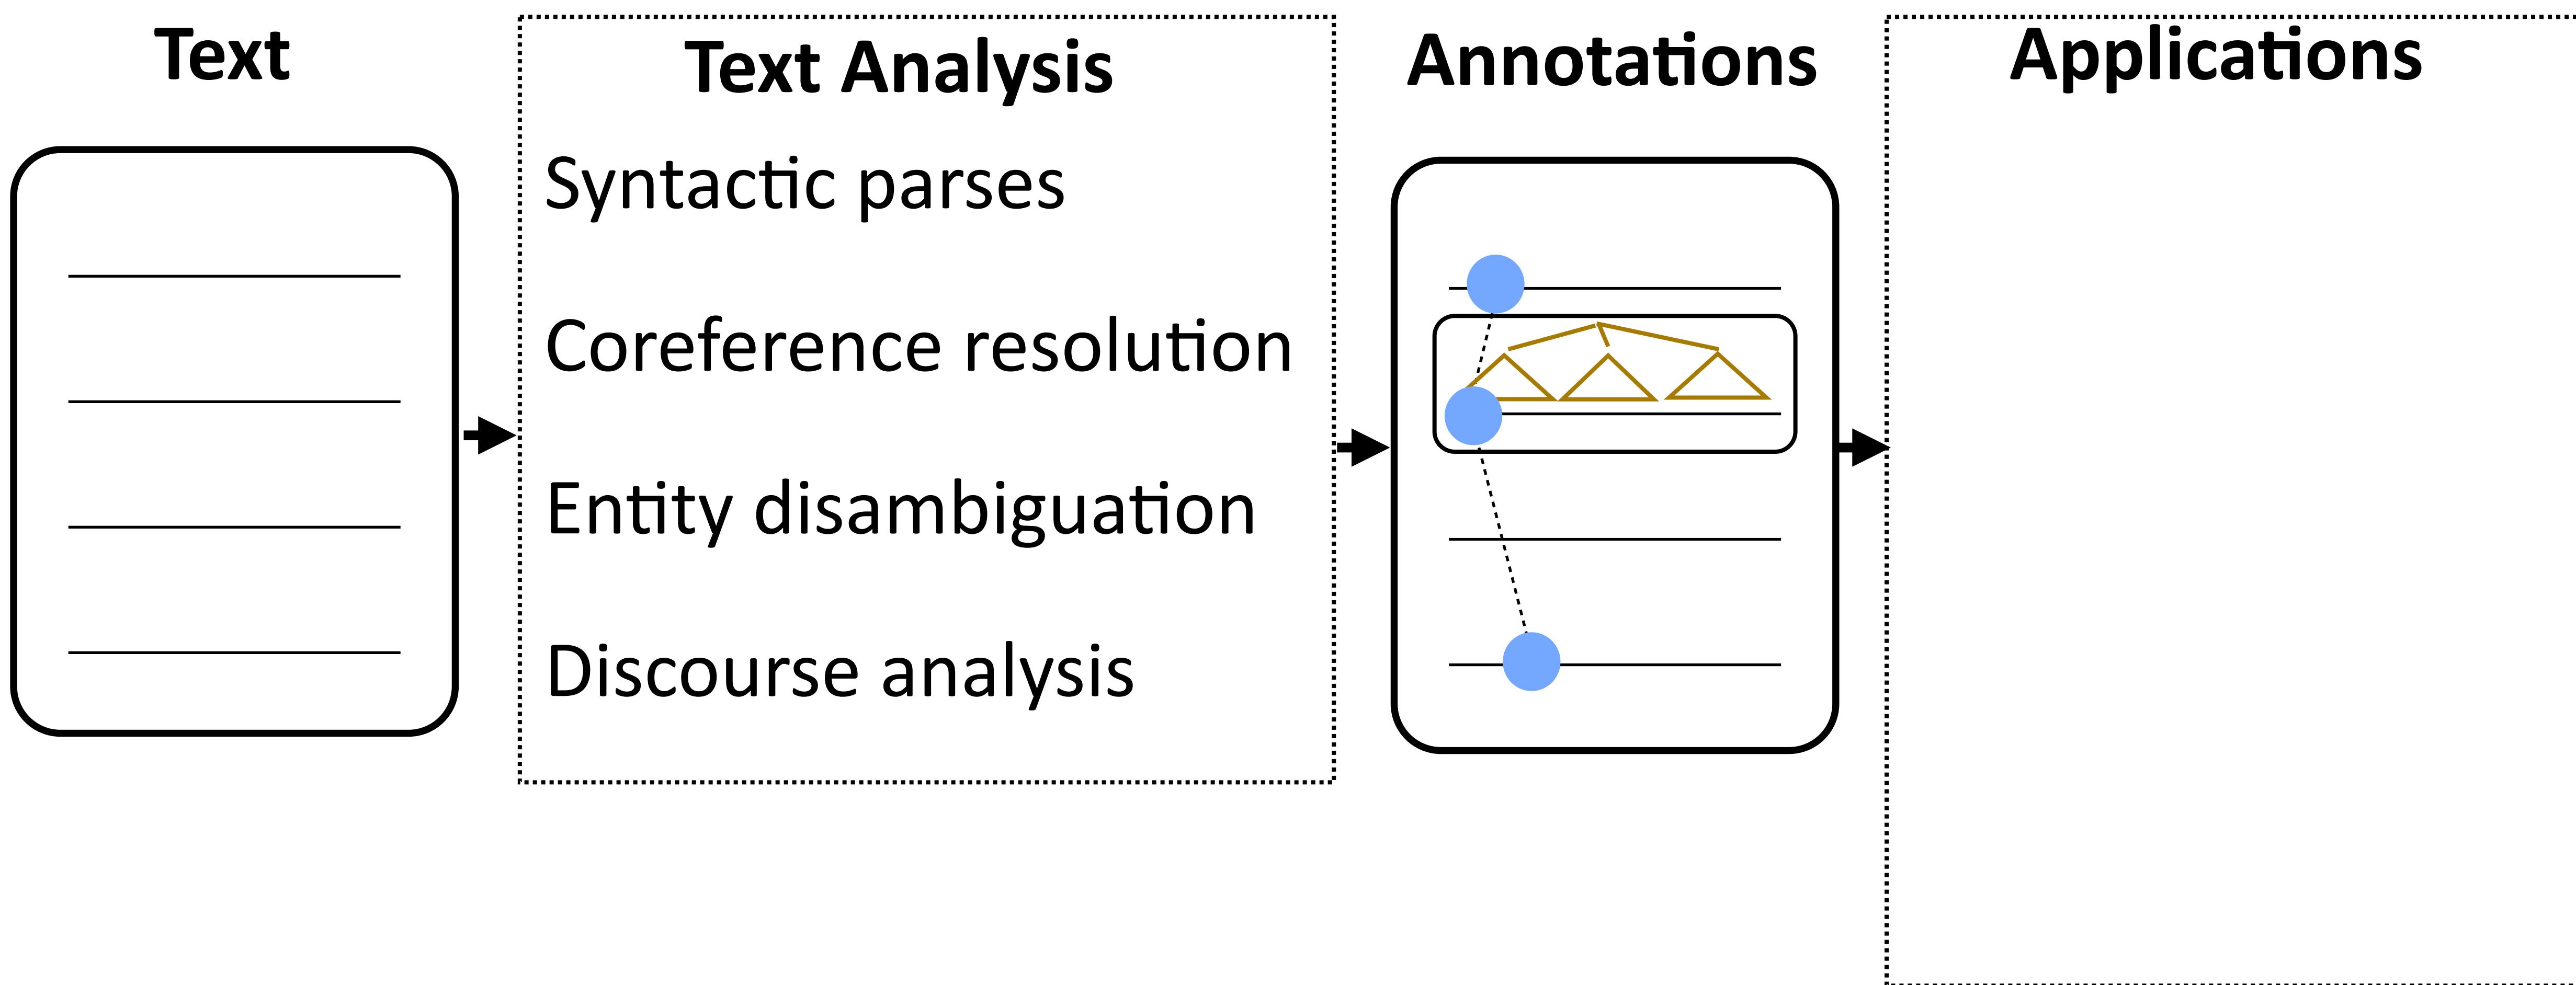
NLP Analysis Pipeline



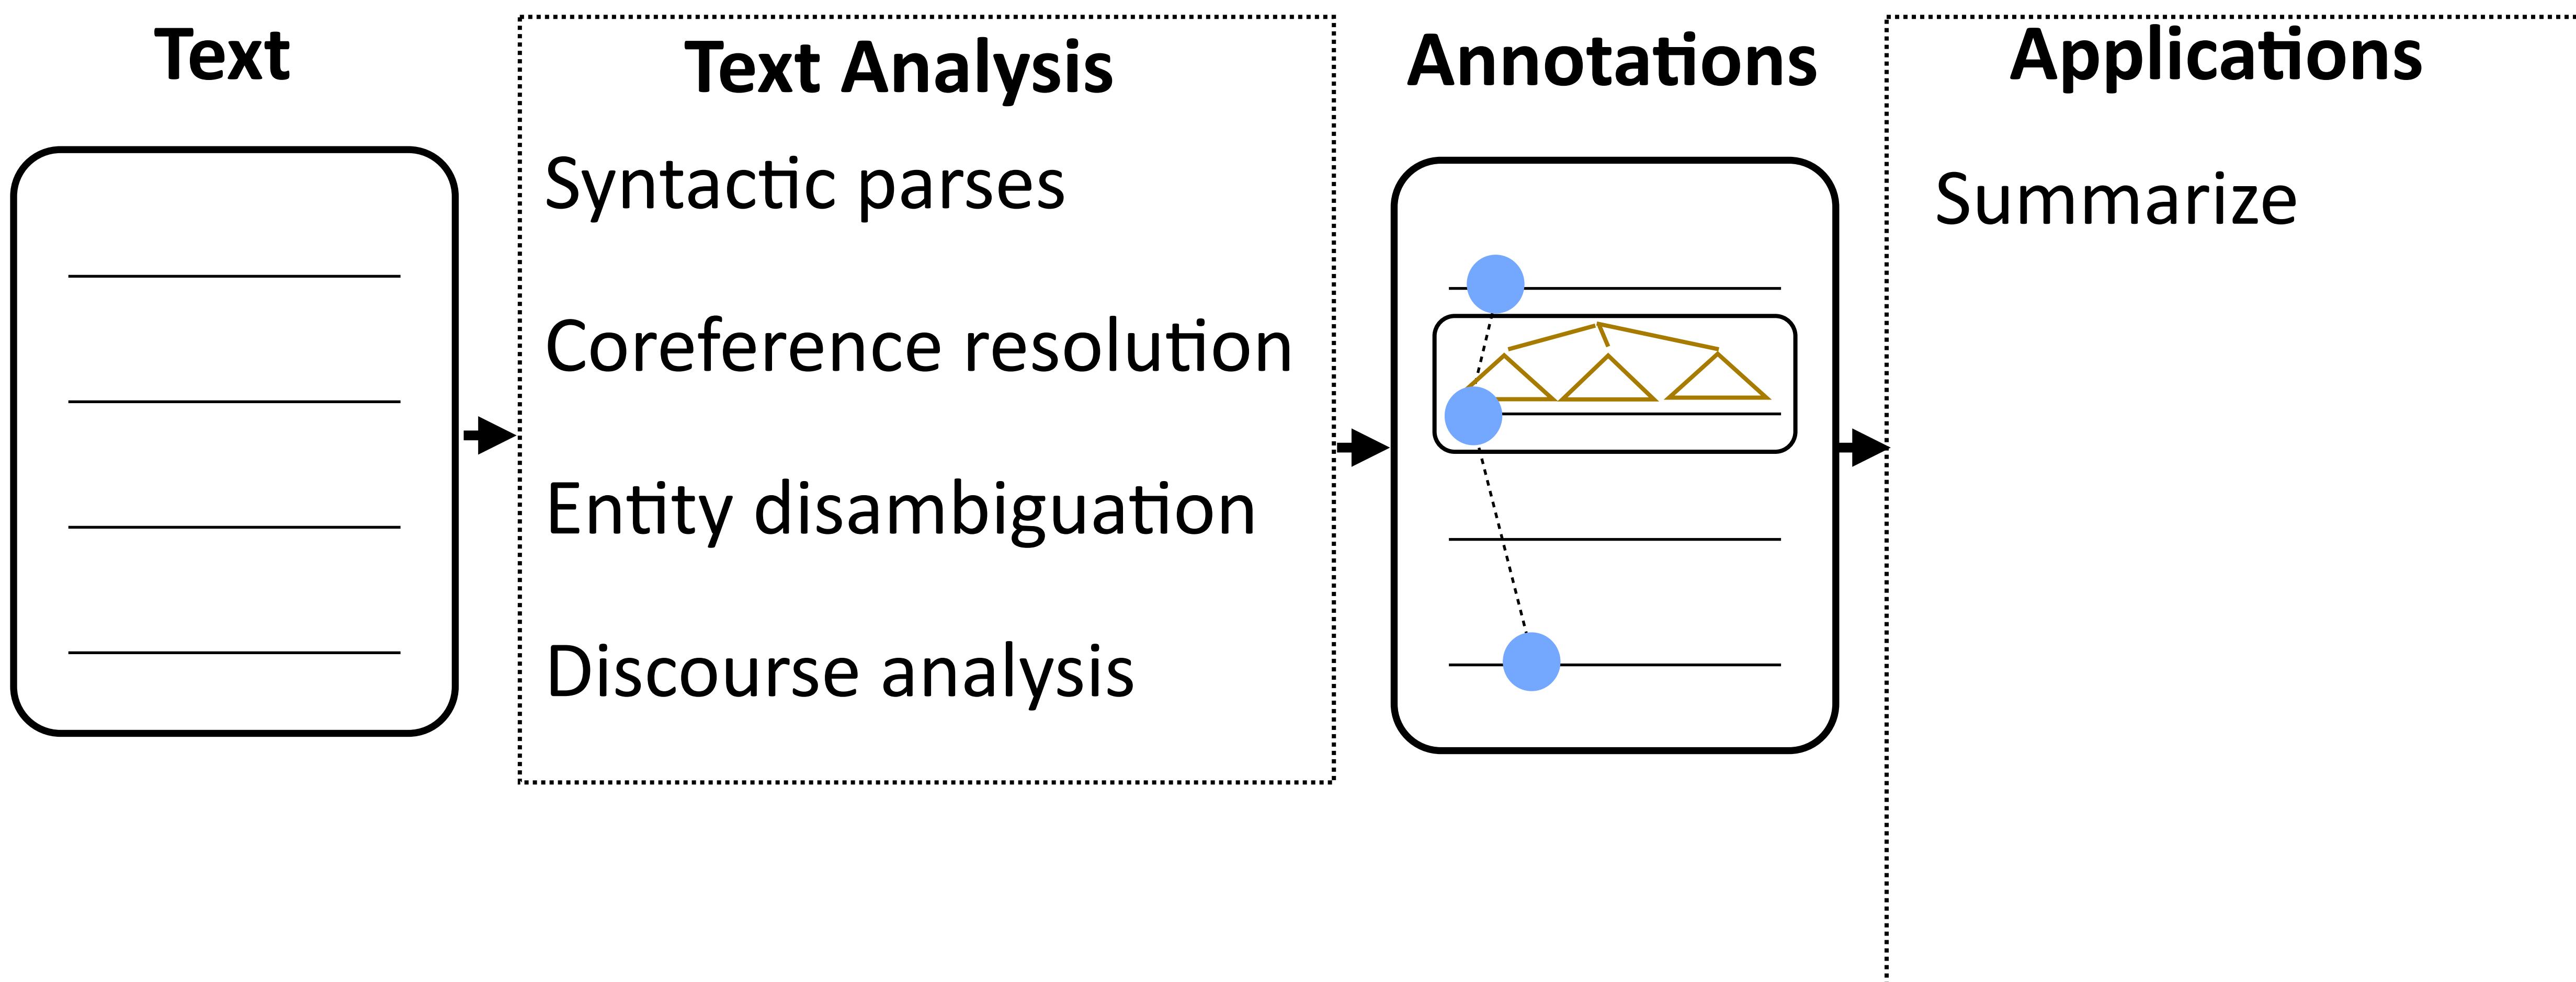
NLP Analysis Pipeline



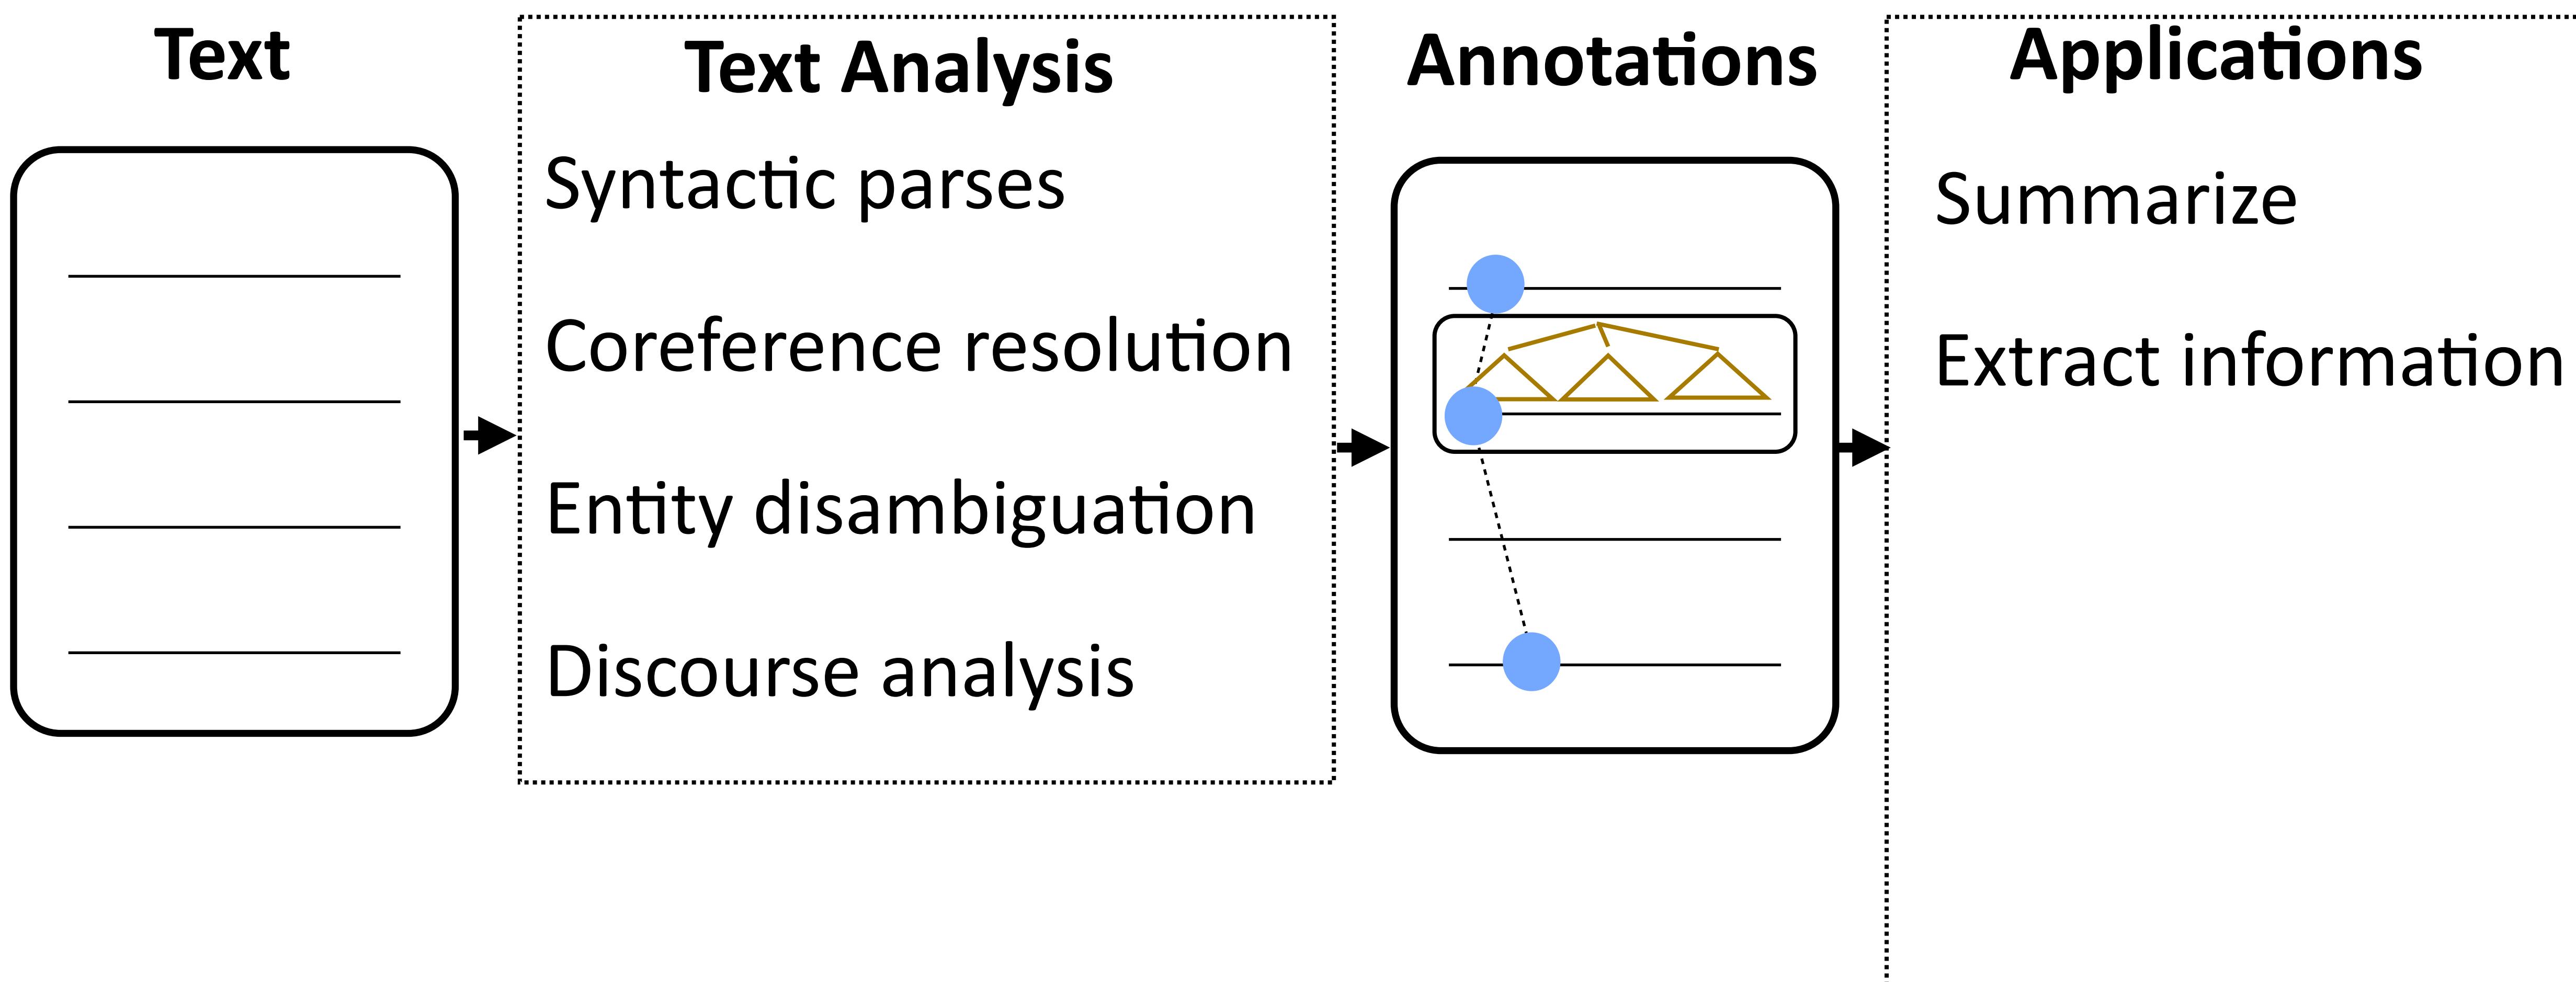
NLP Analysis Pipeline



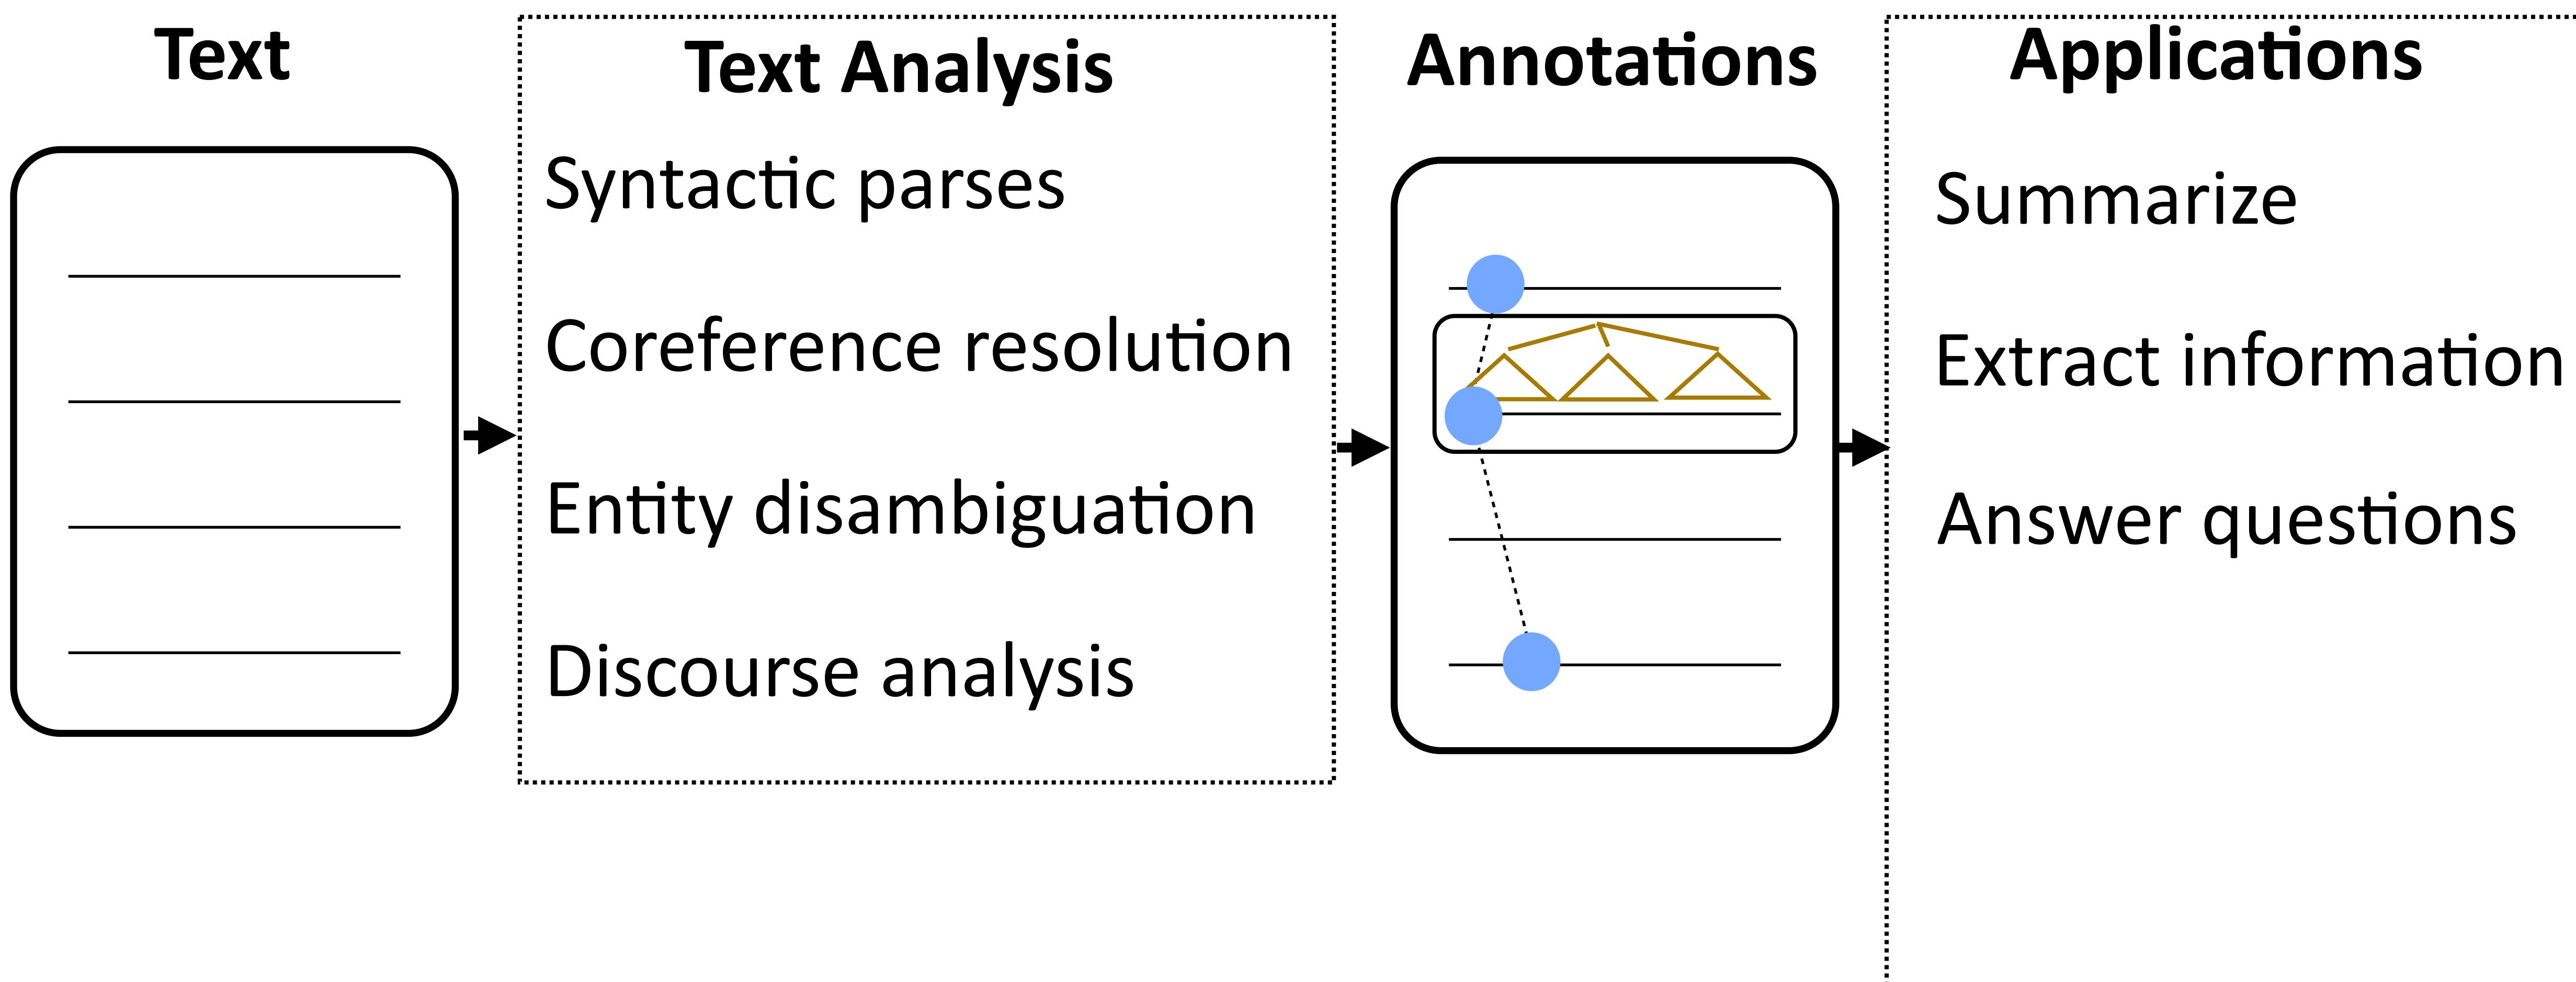
NLP Analysis Pipeline



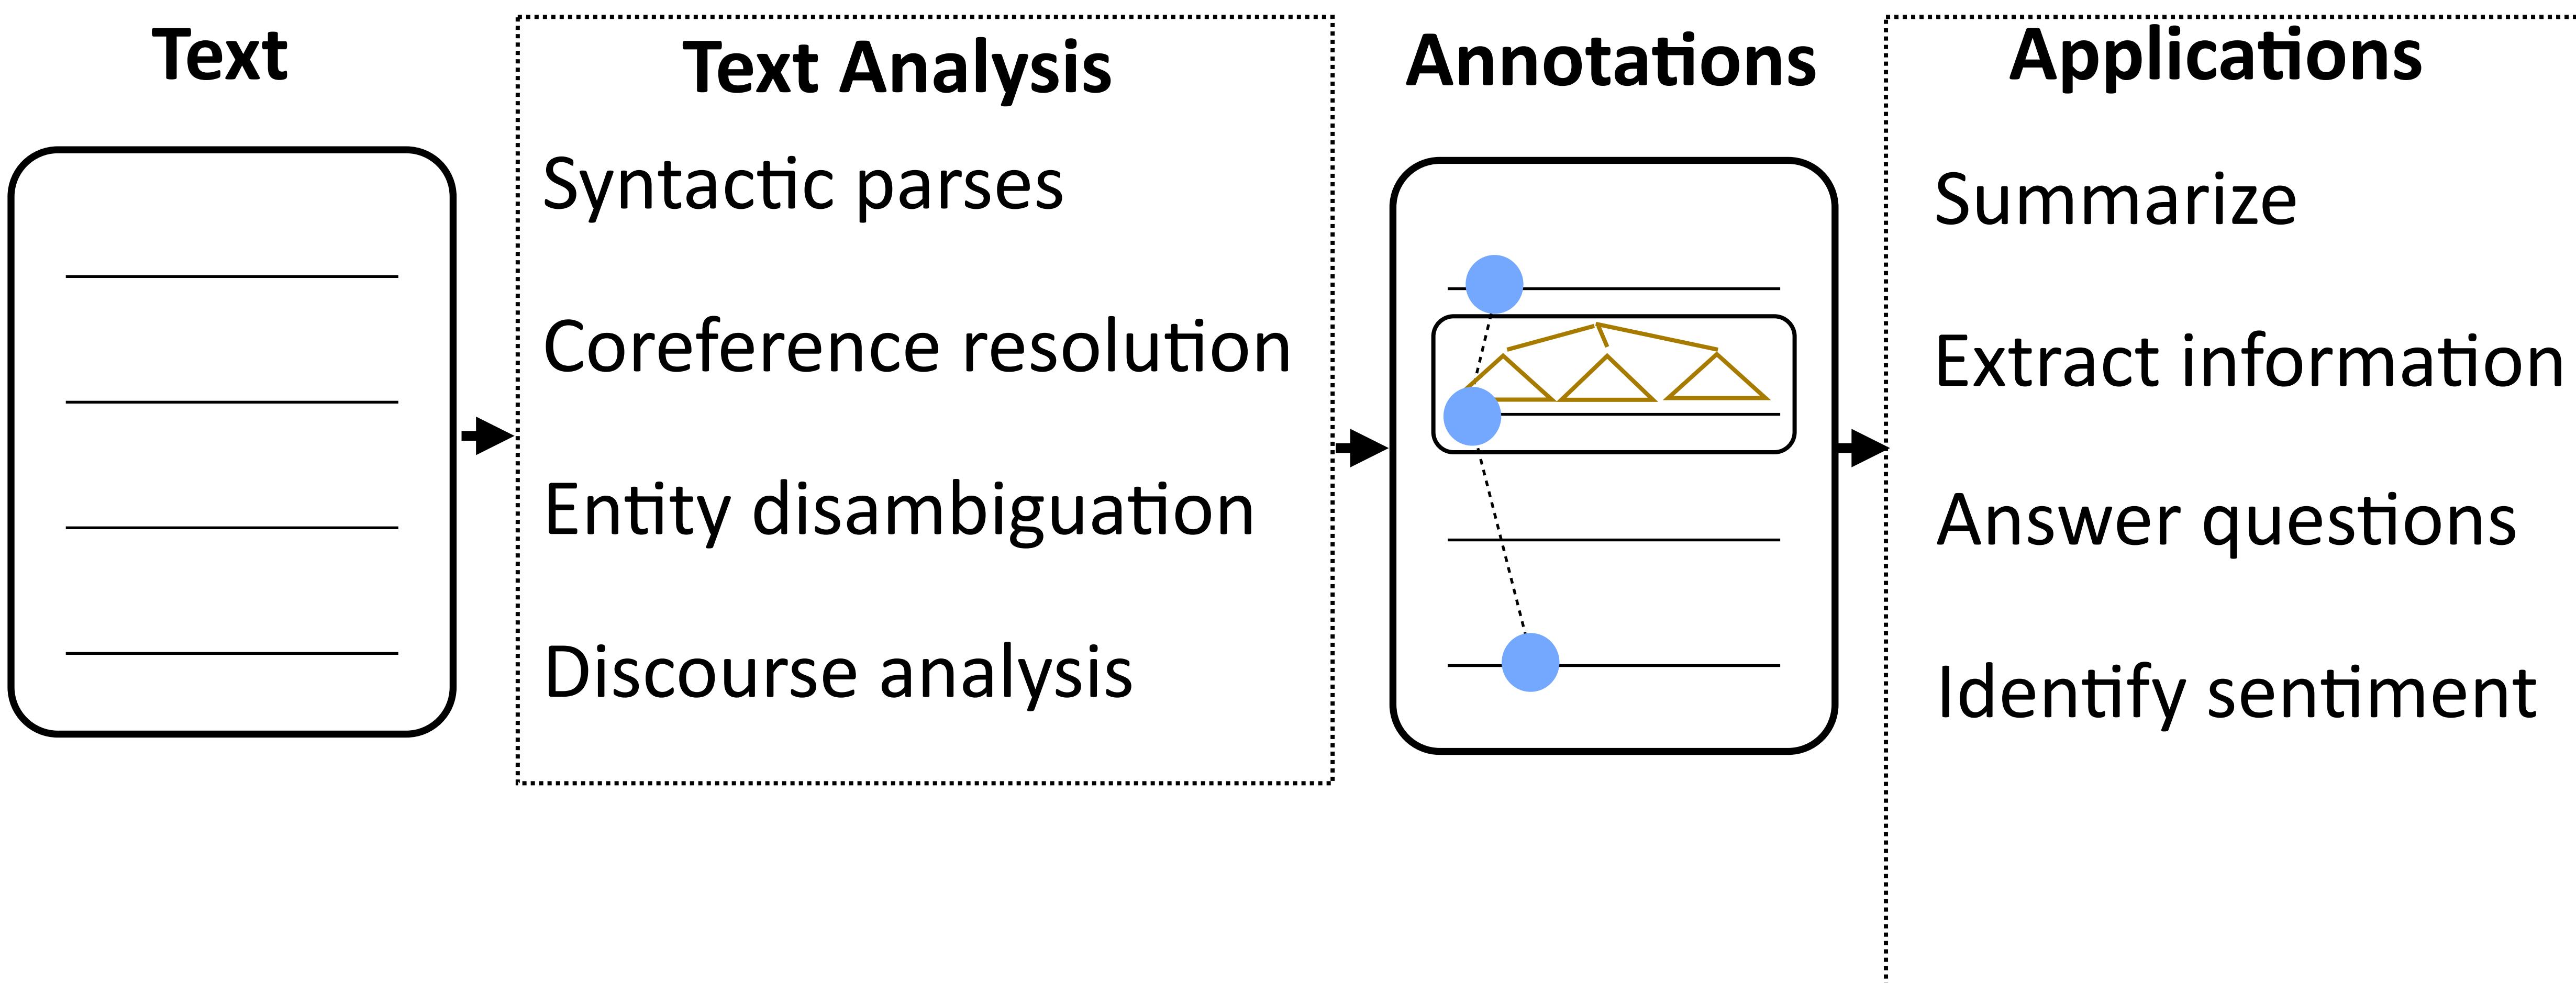
NLP Analysis Pipeline



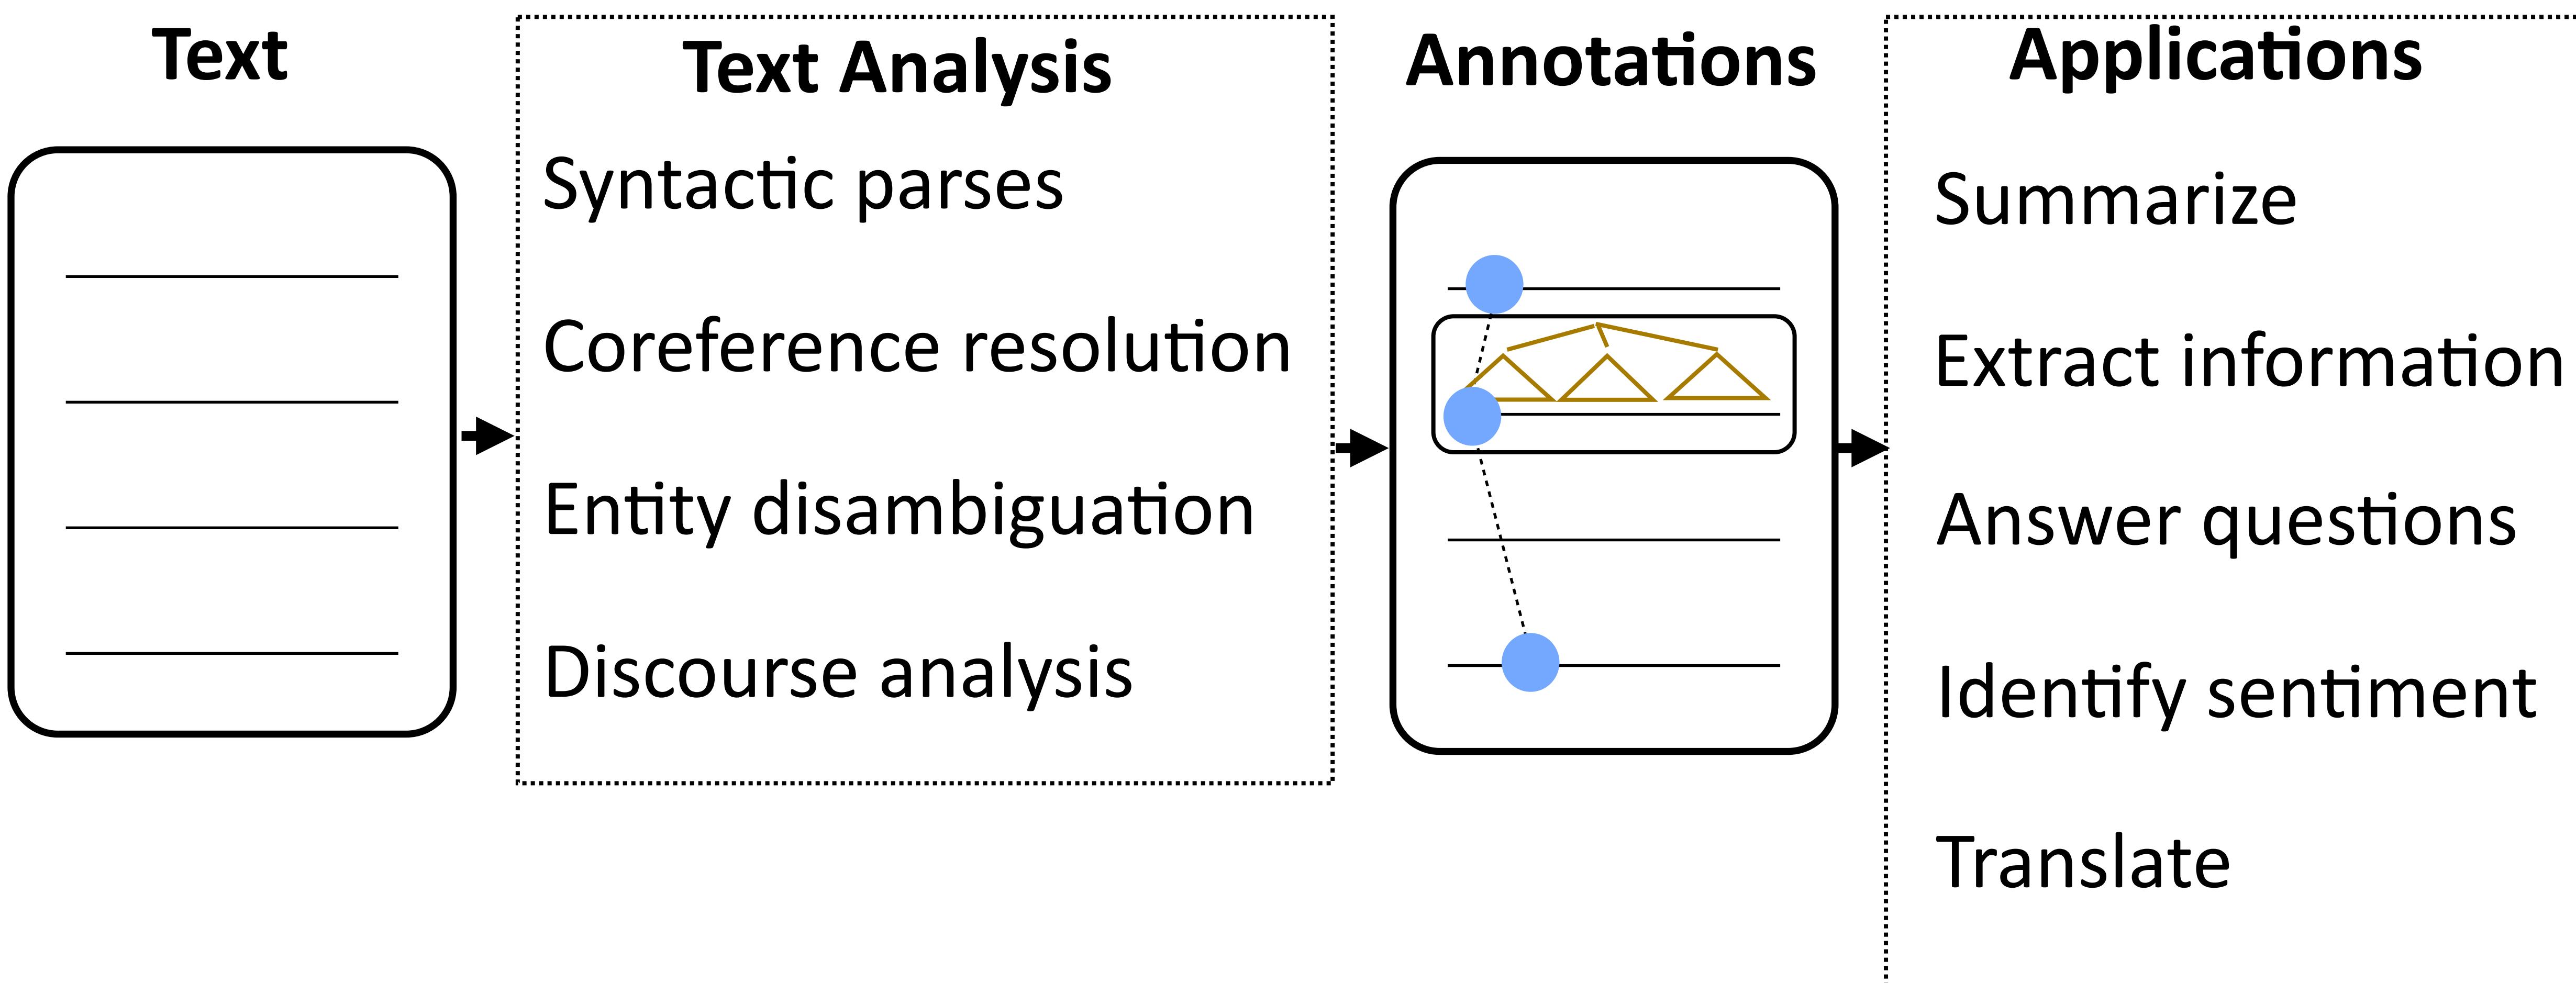
NLP Analysis Pipeline



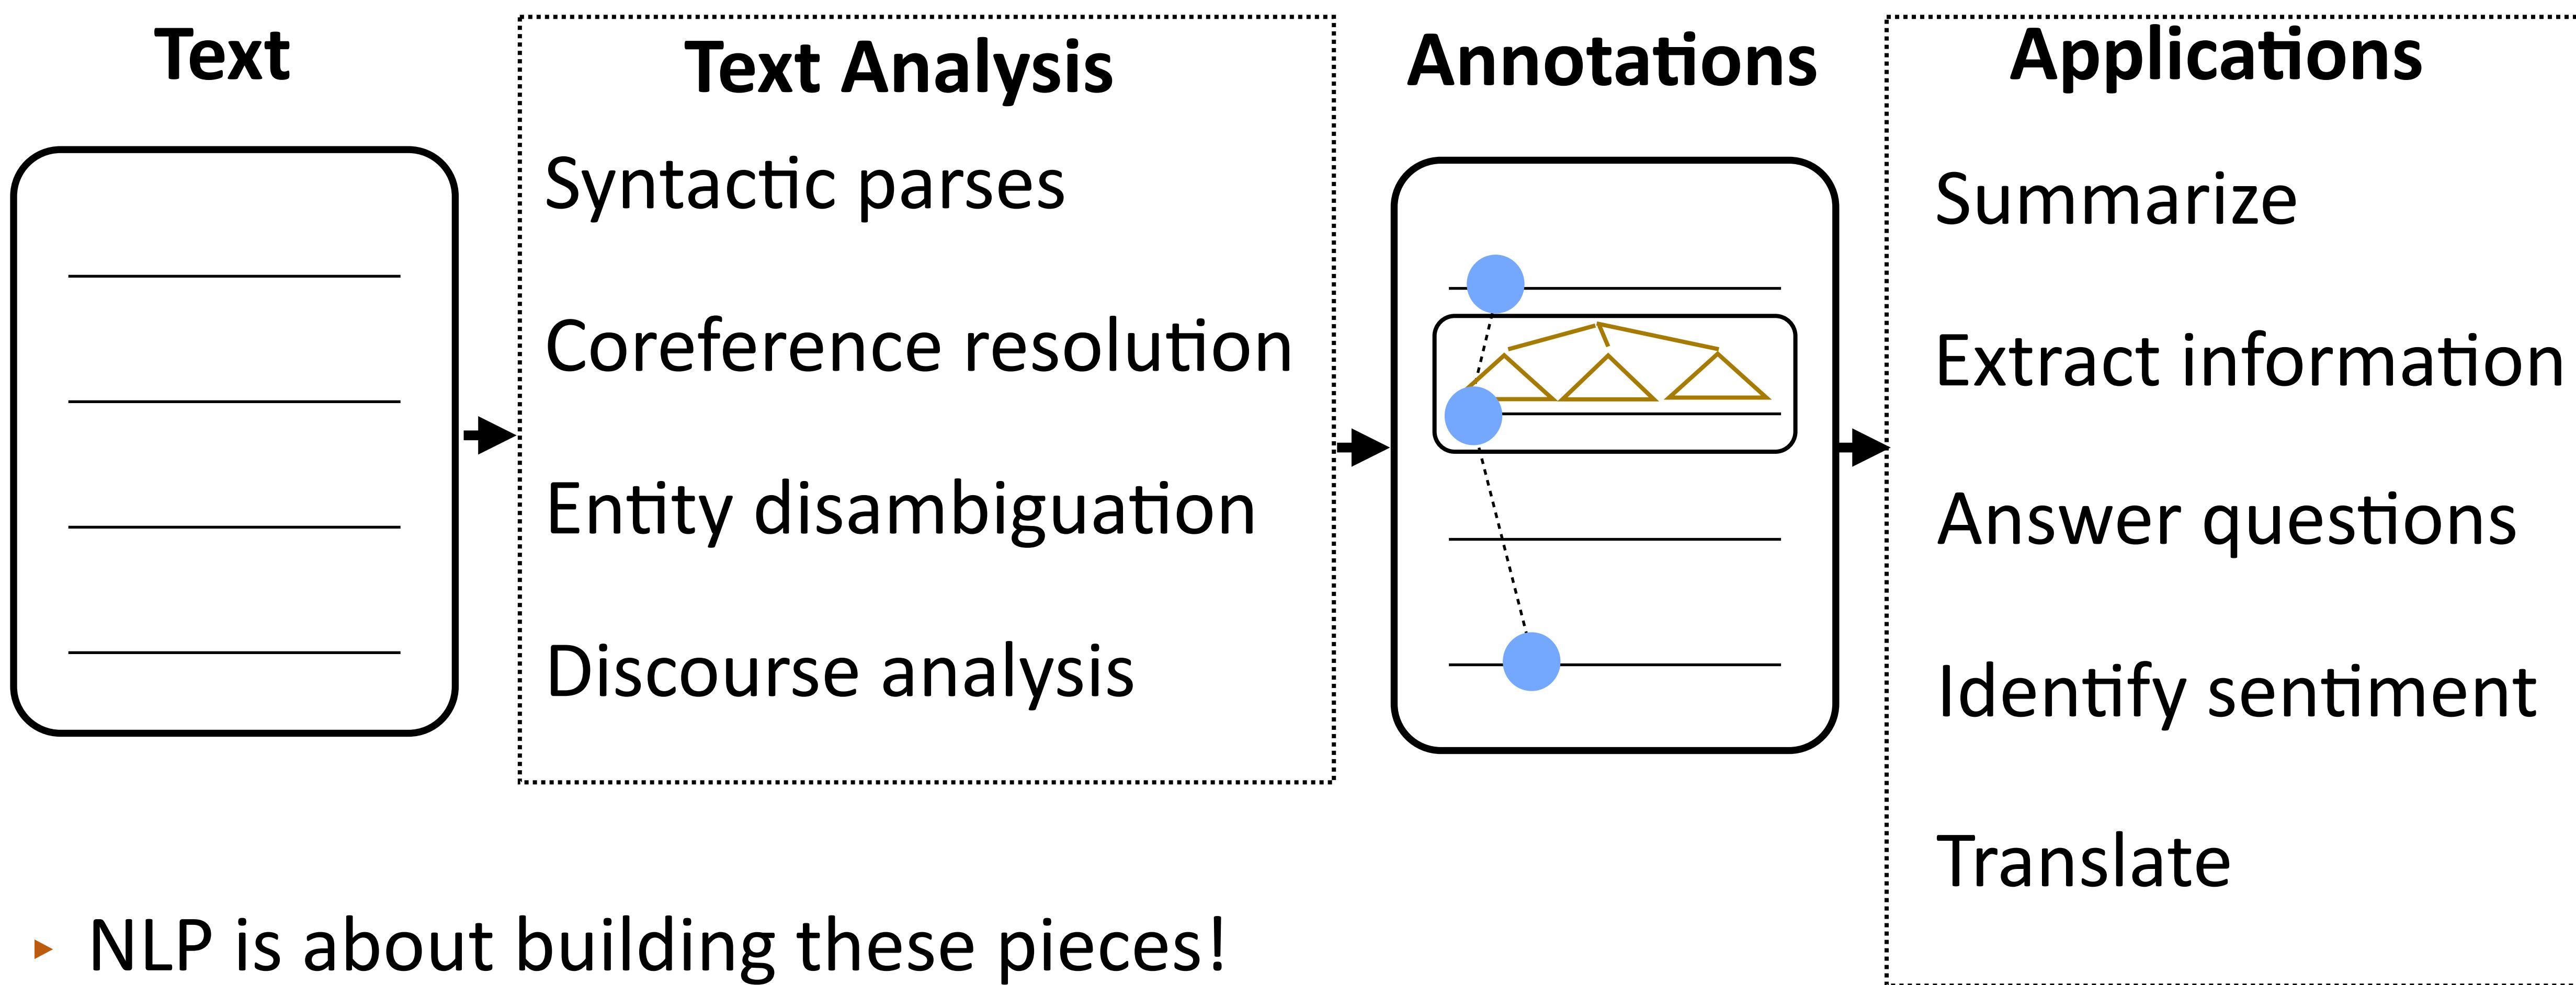
NLP Analysis Pipeline



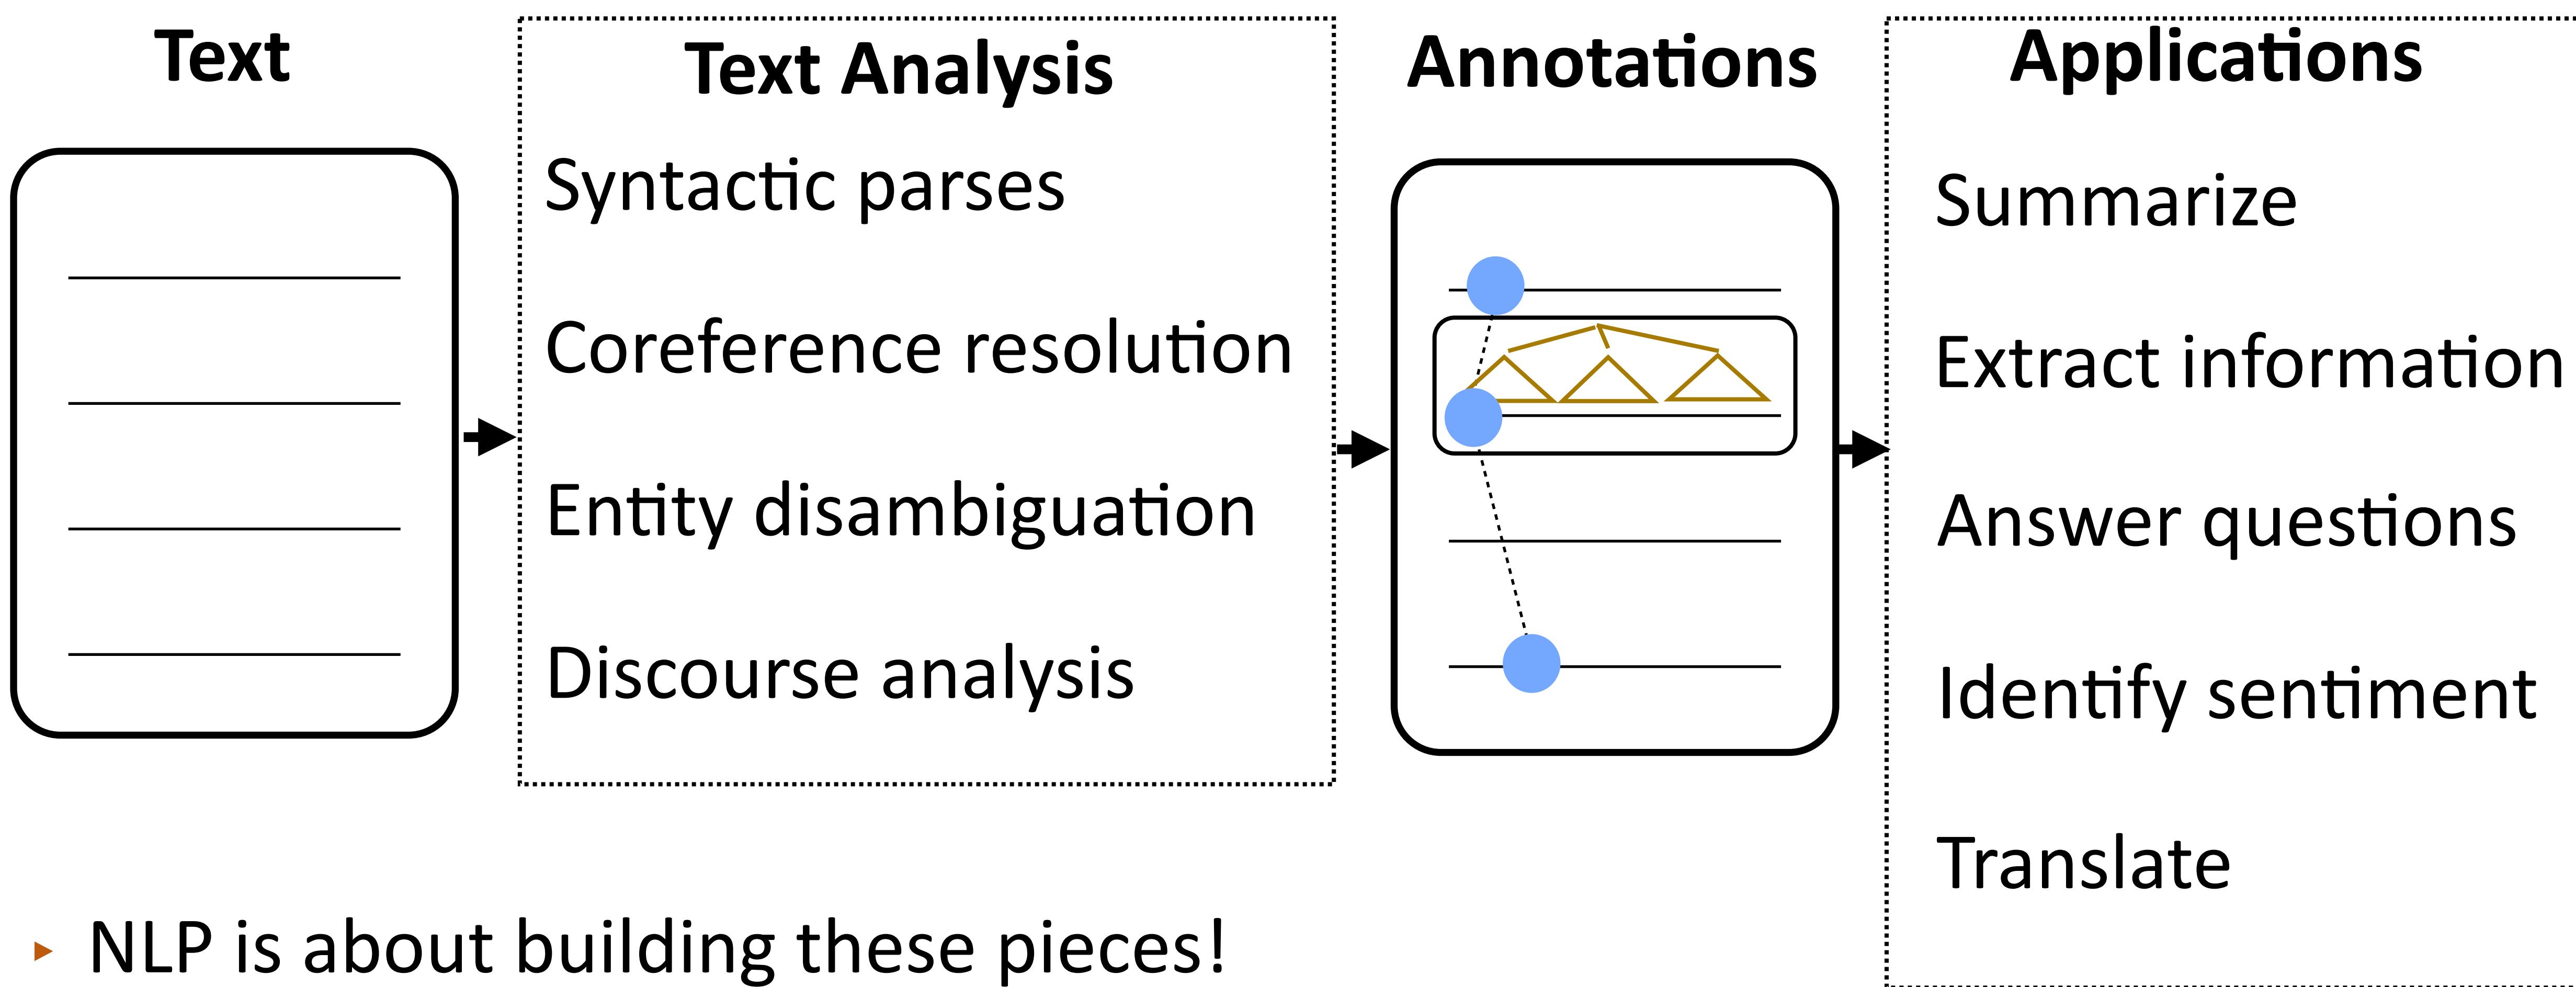
NLP Analysis Pipeline



NLP Analysis Pipeline



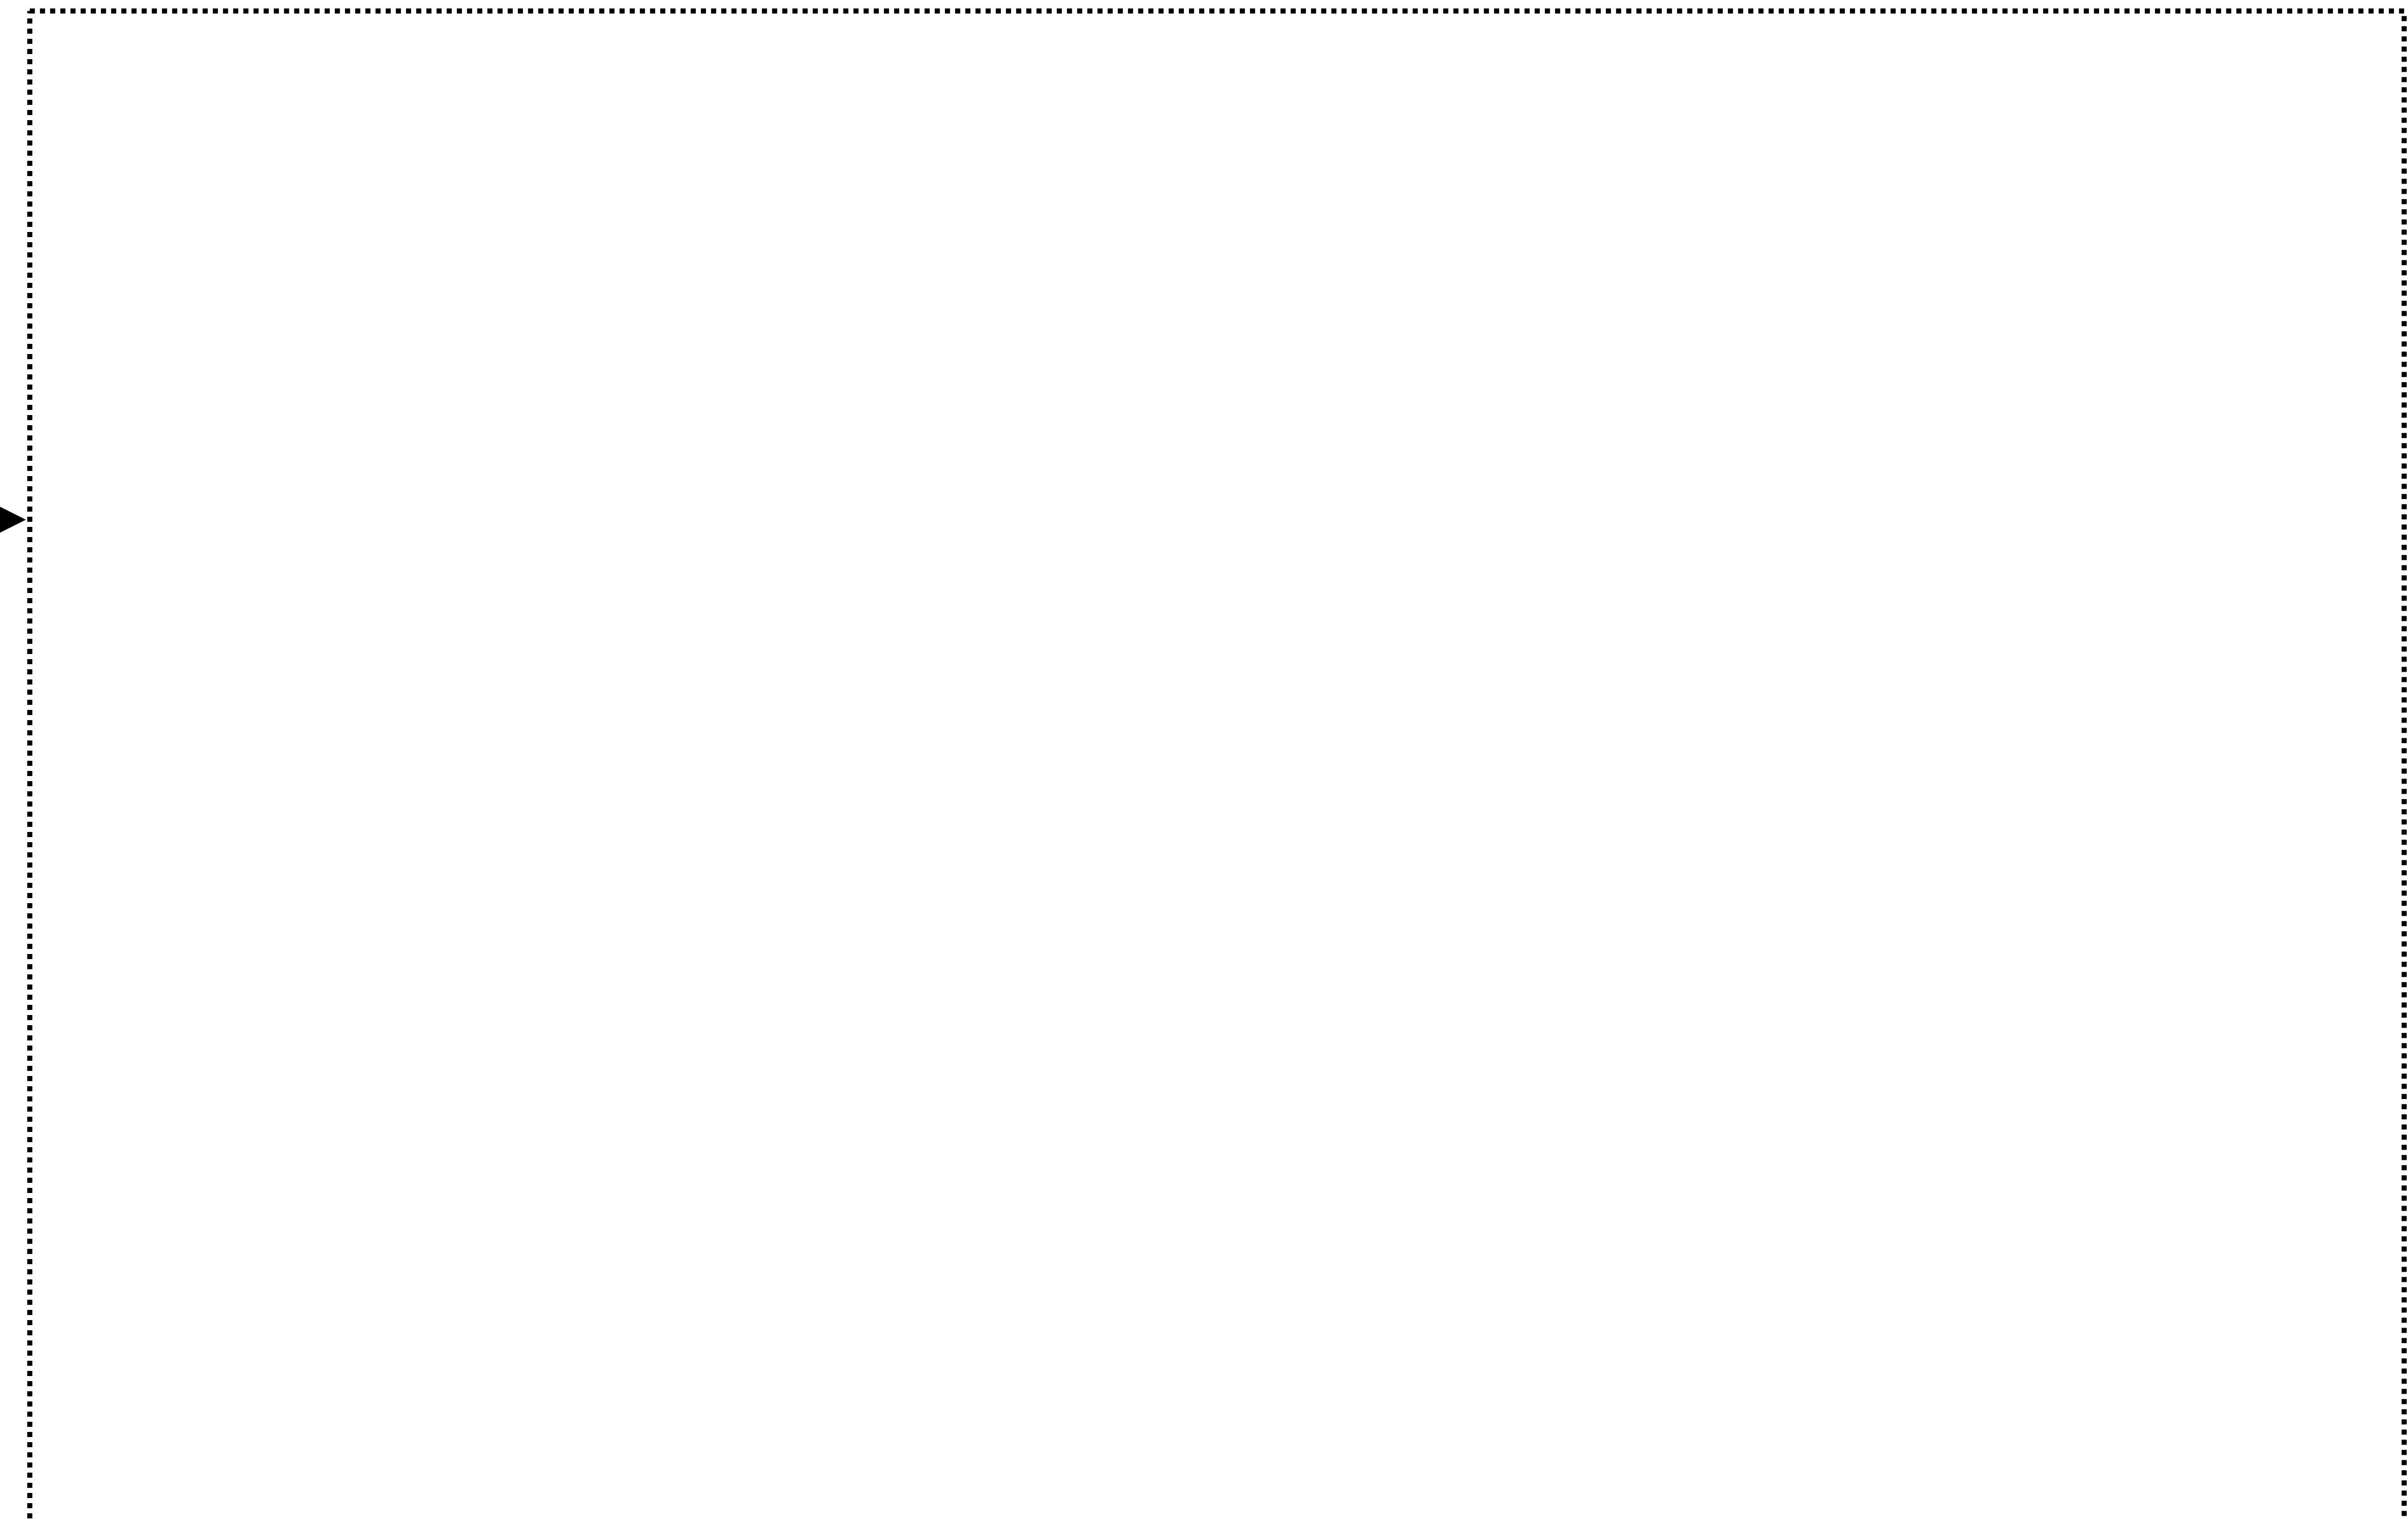
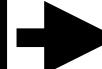
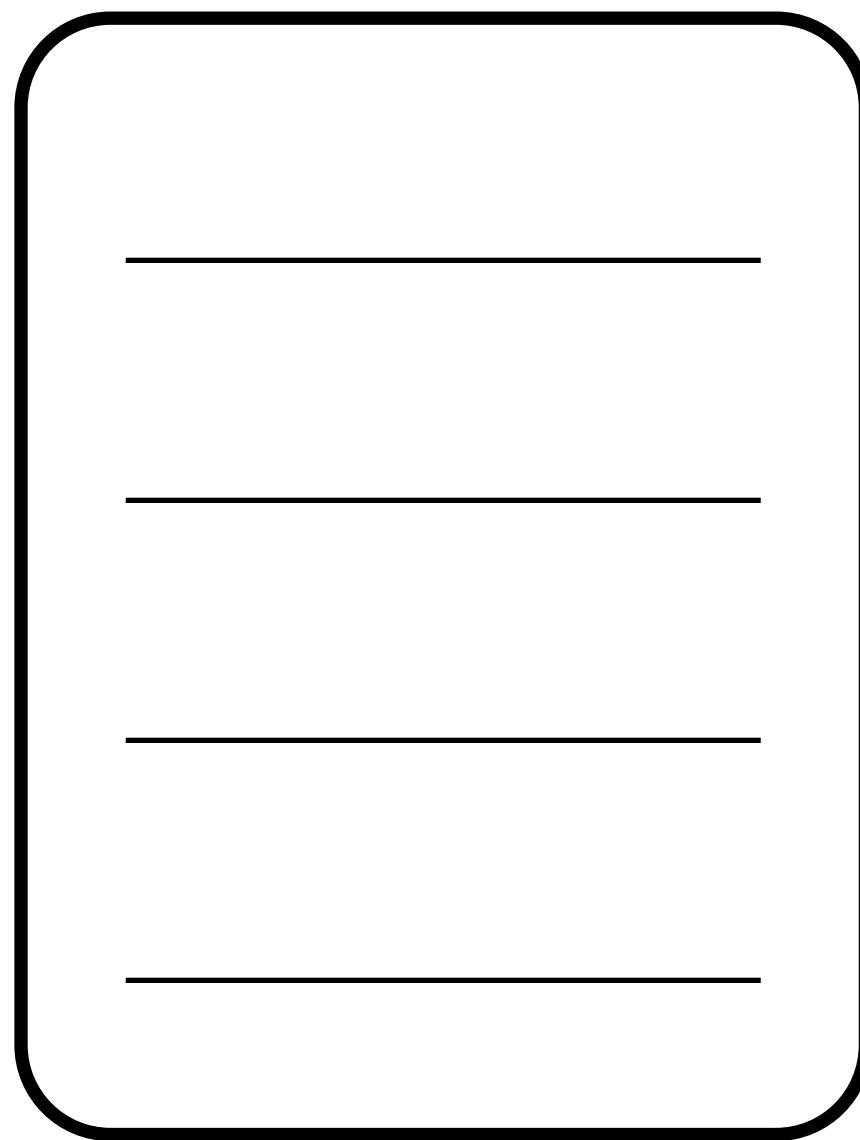
NLP Analysis Pipeline



- ▶ NLP is about building these pieces!
- ▶ All of these components are modeled with statistical approaches trained with machine learning

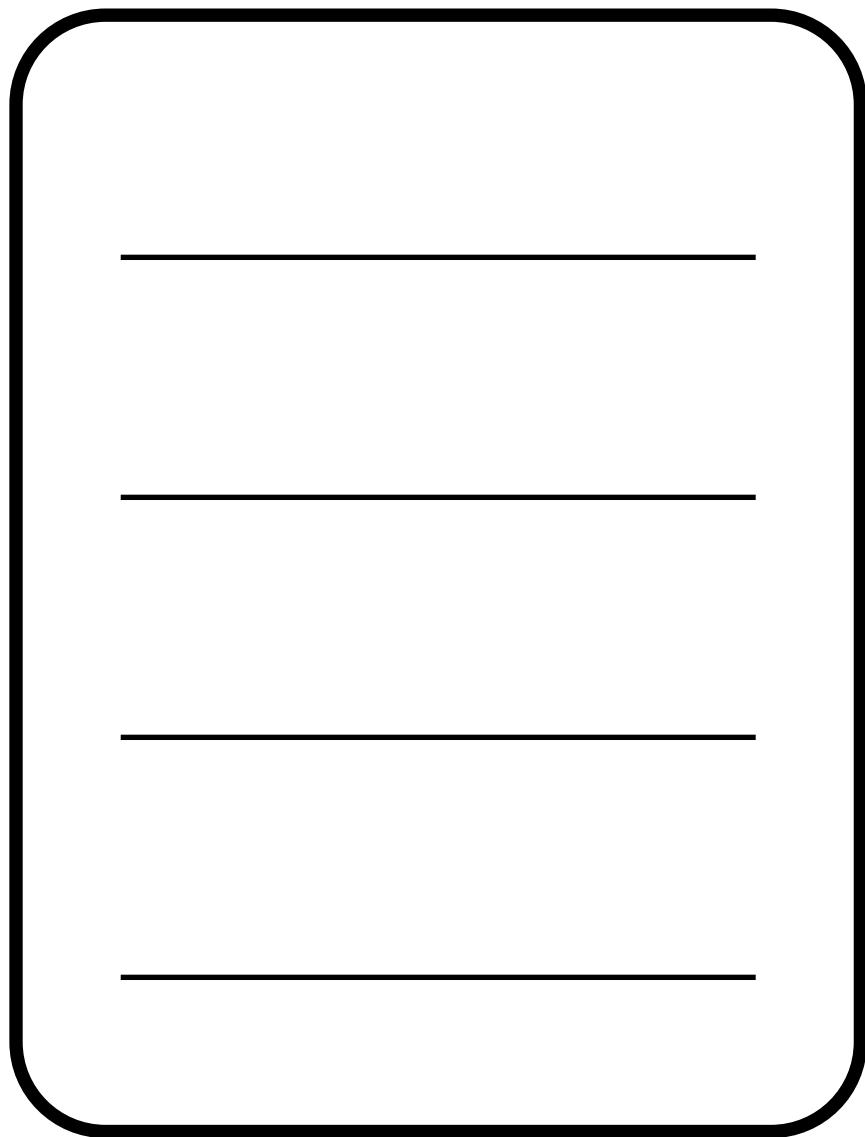
How do we represent language?

Text



How do we represent language?

Text

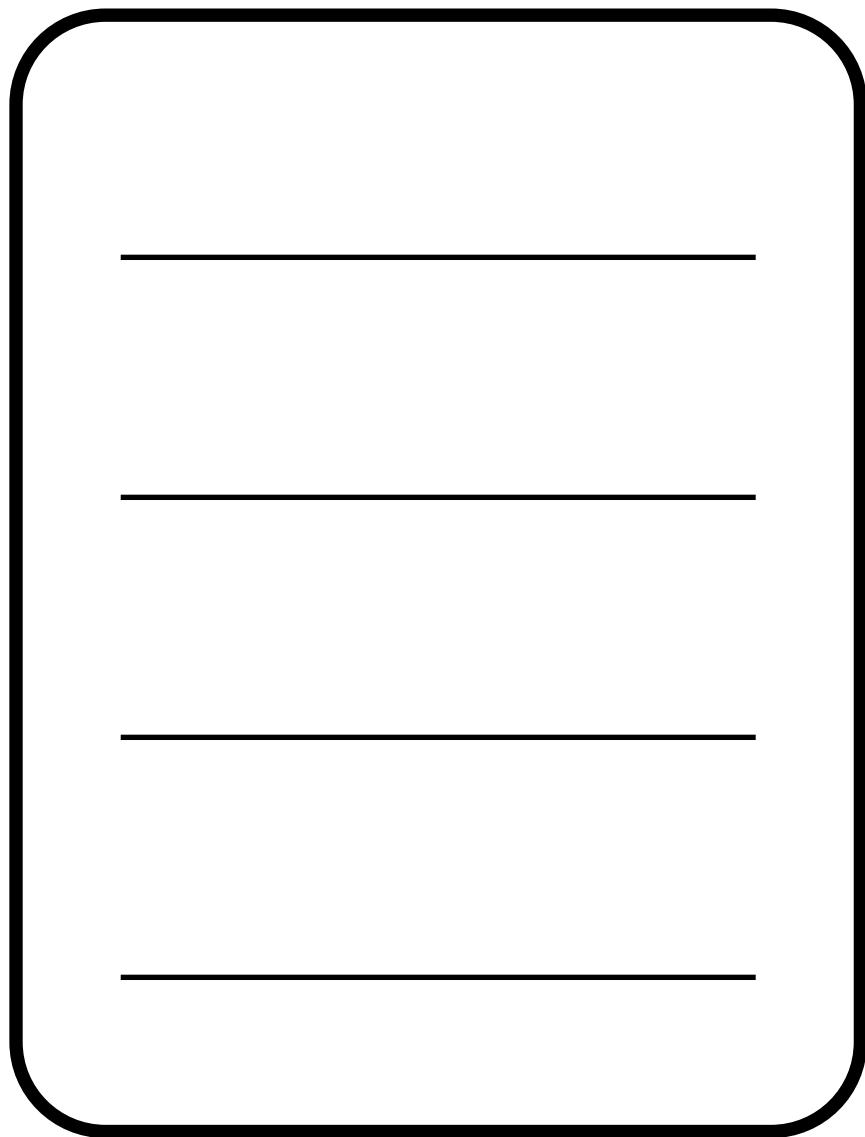


Labels



How do we represent language?

Text

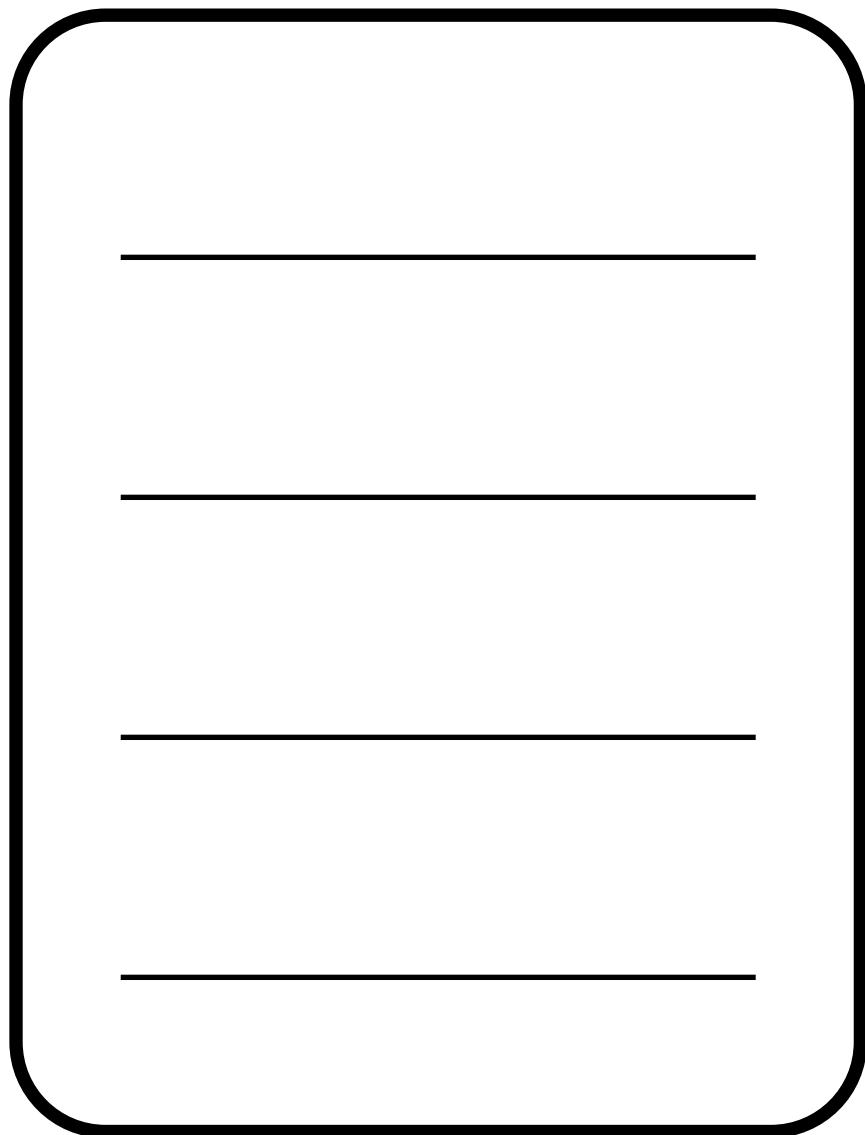


Labels

the movie was good +

How do we represent language?

Text



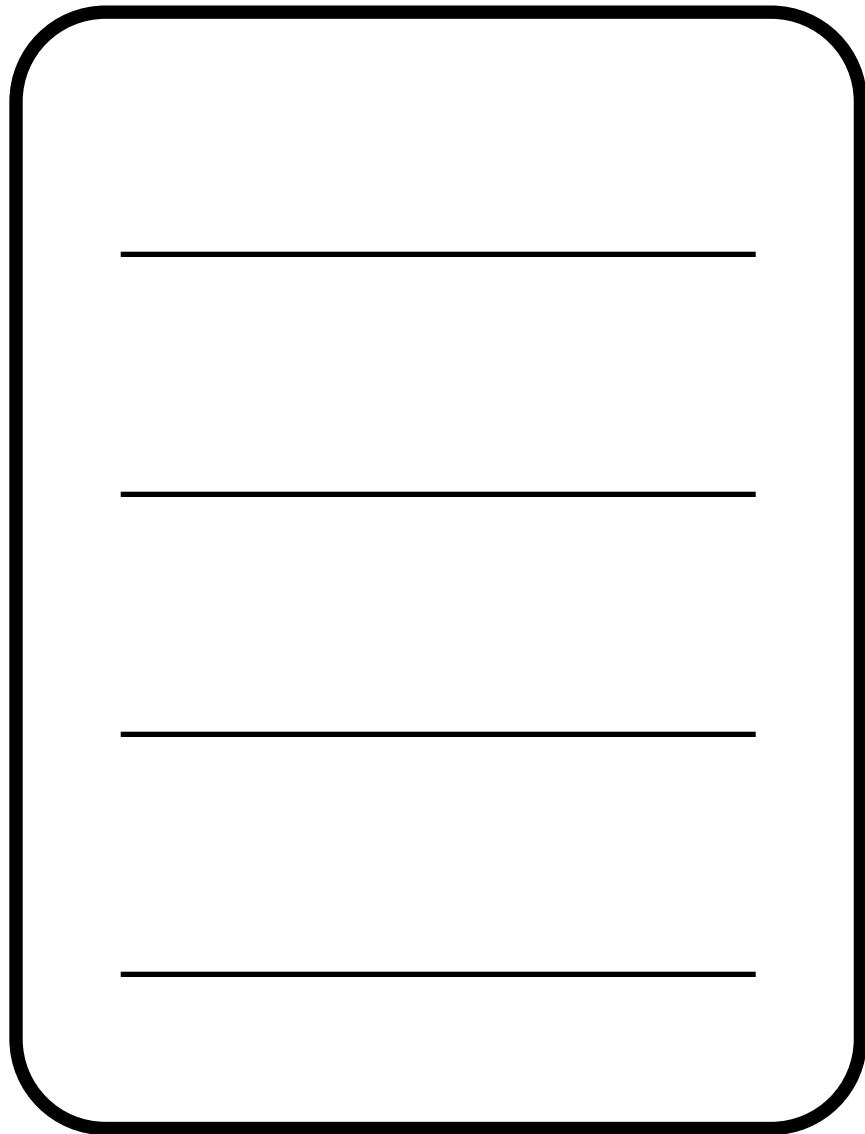
Labels

the movie was good +

Beyoncé had one of the best videos of all time subjective

How do we represent language?

Text



Labels

the movie was good +

Beyoncé had one of the best videos of all time subjective

Sequences/tags

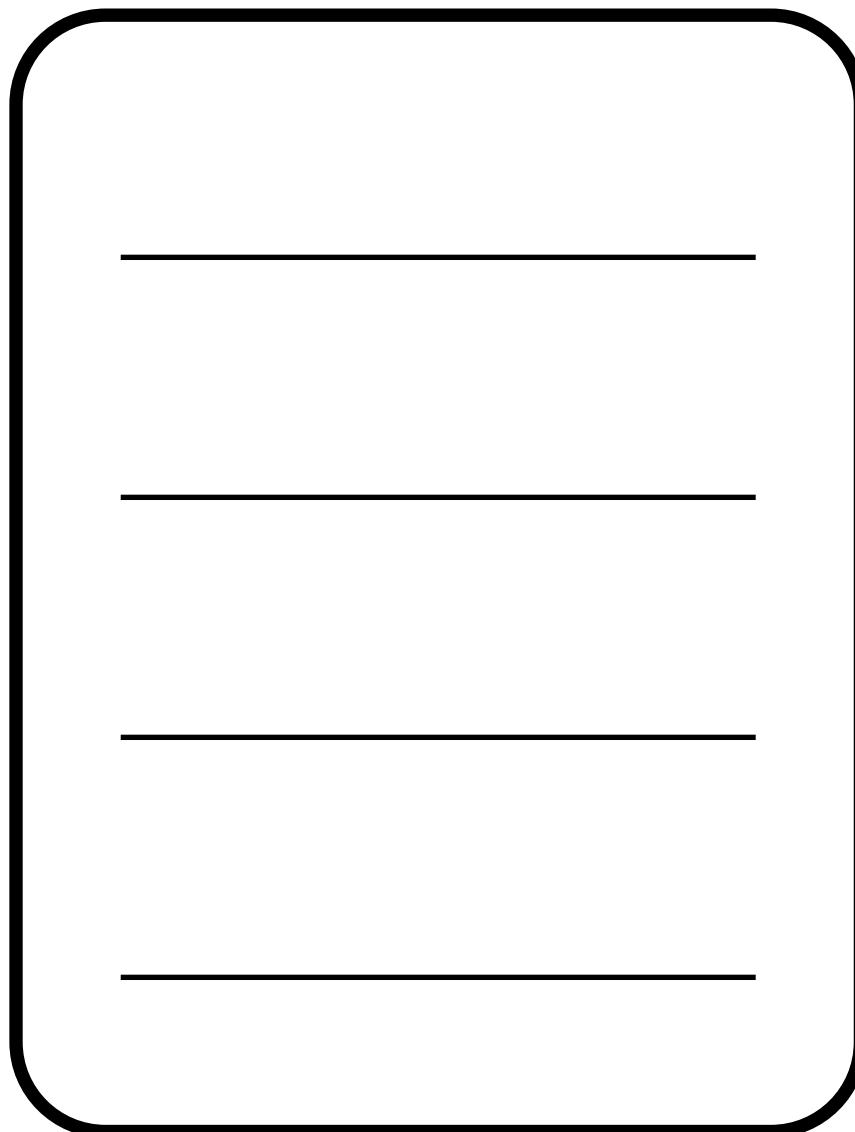
PERSON

Tom Cruise stars in the new Mission Impossible film

MOVIE

How do we represent language?

Text



Labels

the movie was good +

Beyoncé had one of the best videos of all time subjective

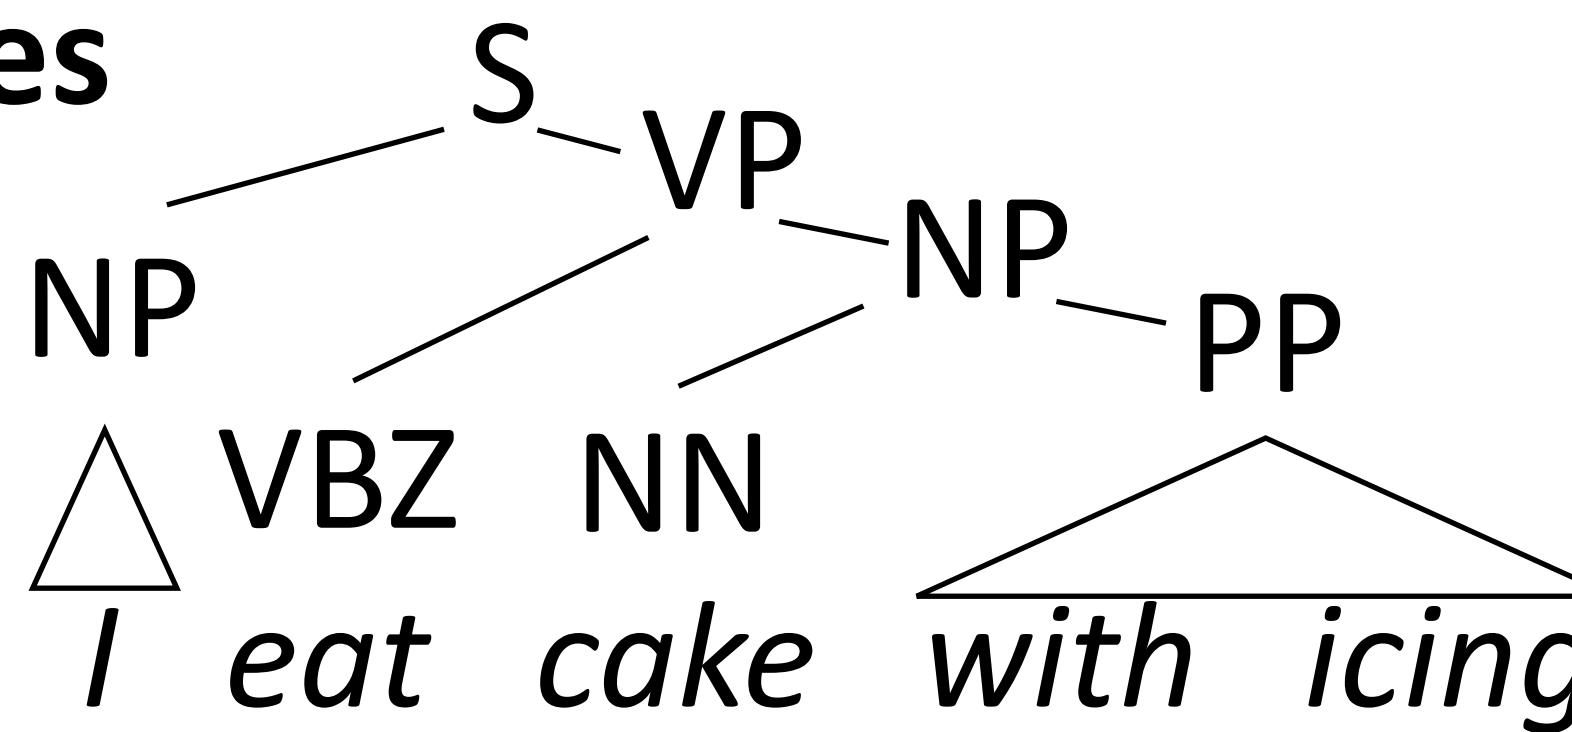
Sequences/tags

PERSON

Tom Cruise stars in the new Mission Impossible film

MOVIE

Trees



How do we represent language?

Text

Labels

the movie was good +

Beyoncé had one of the best videos of all time subjective

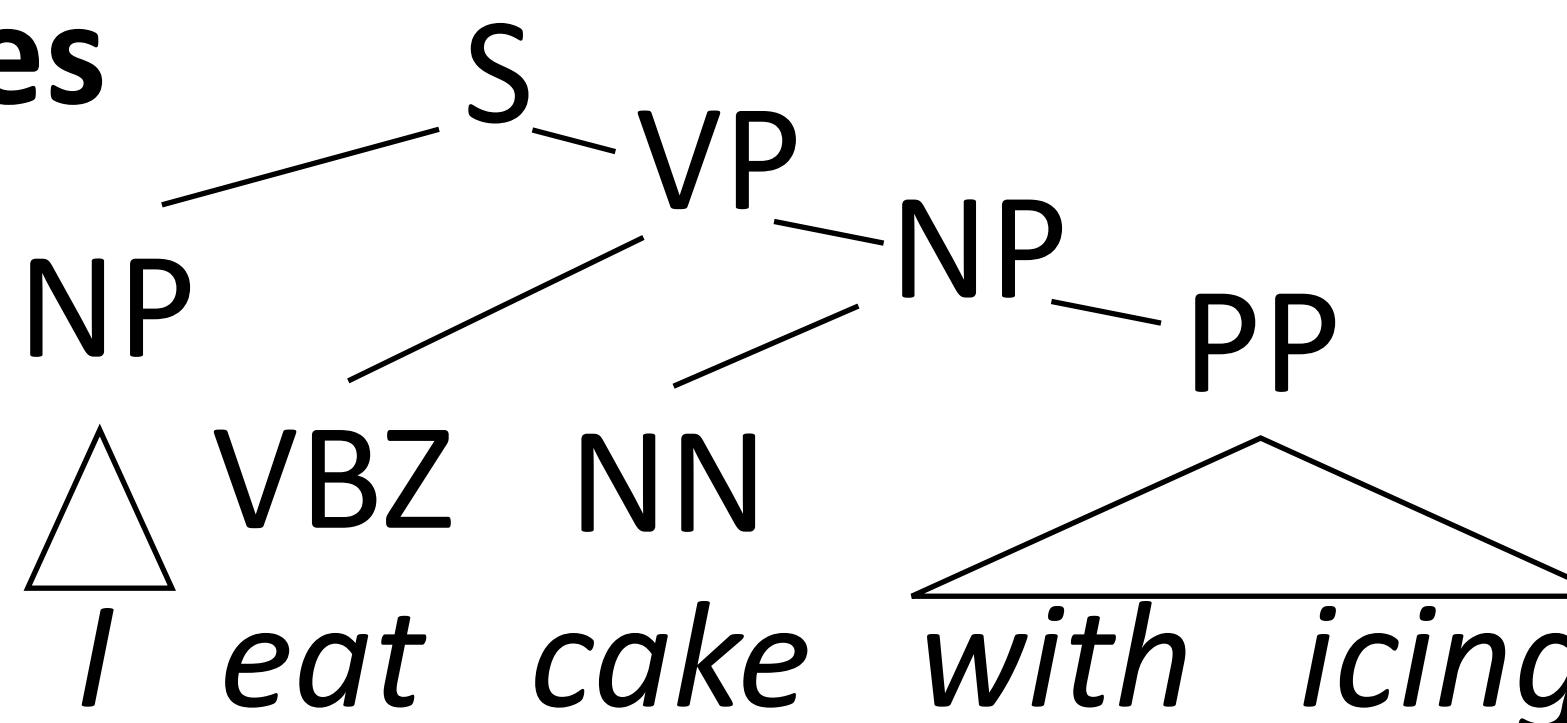
Sequences/tags

PERSON

Tom Cruise stars in the new Mission Impossible film

MOVIE

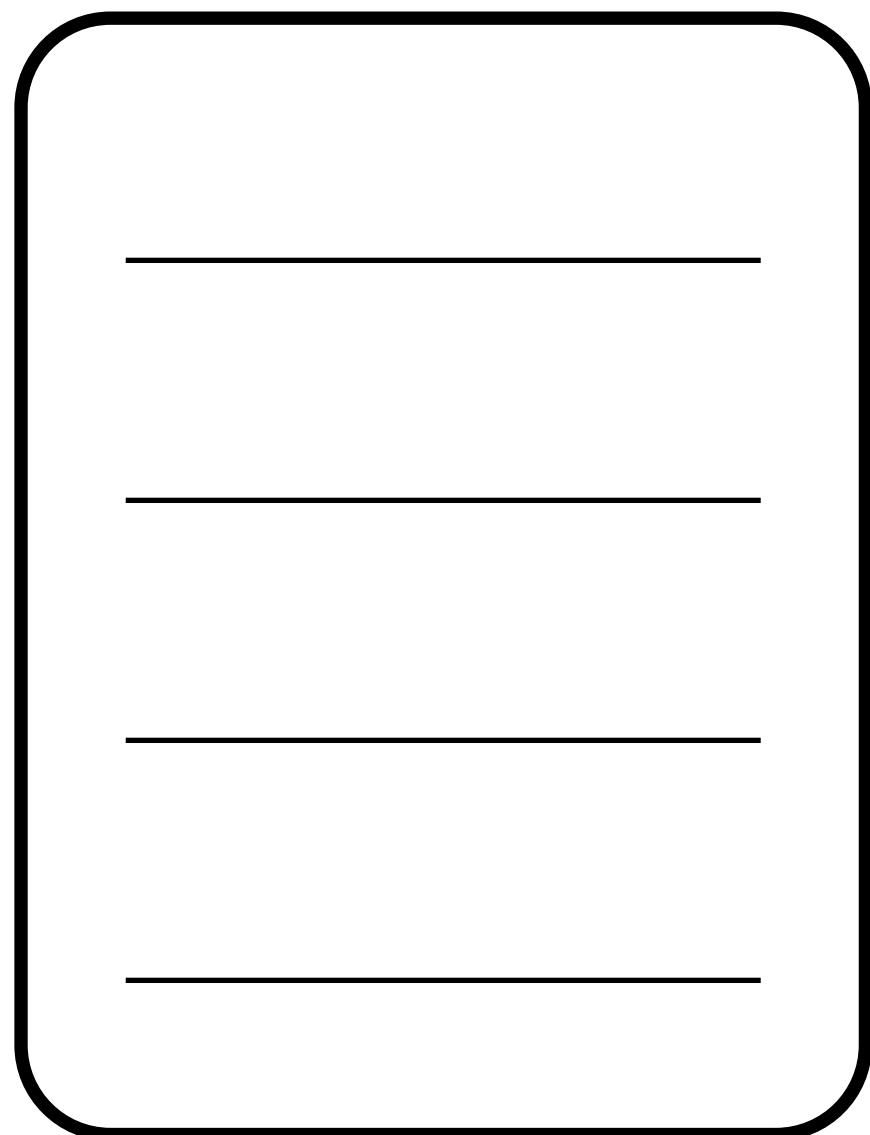
Trees



$\lambda x. \text{flight}(x) \wedge \text{dest}(x) = \text{Miami}$
flights to Miami

How do we use these representations?

Text



Text Analysis

Labels

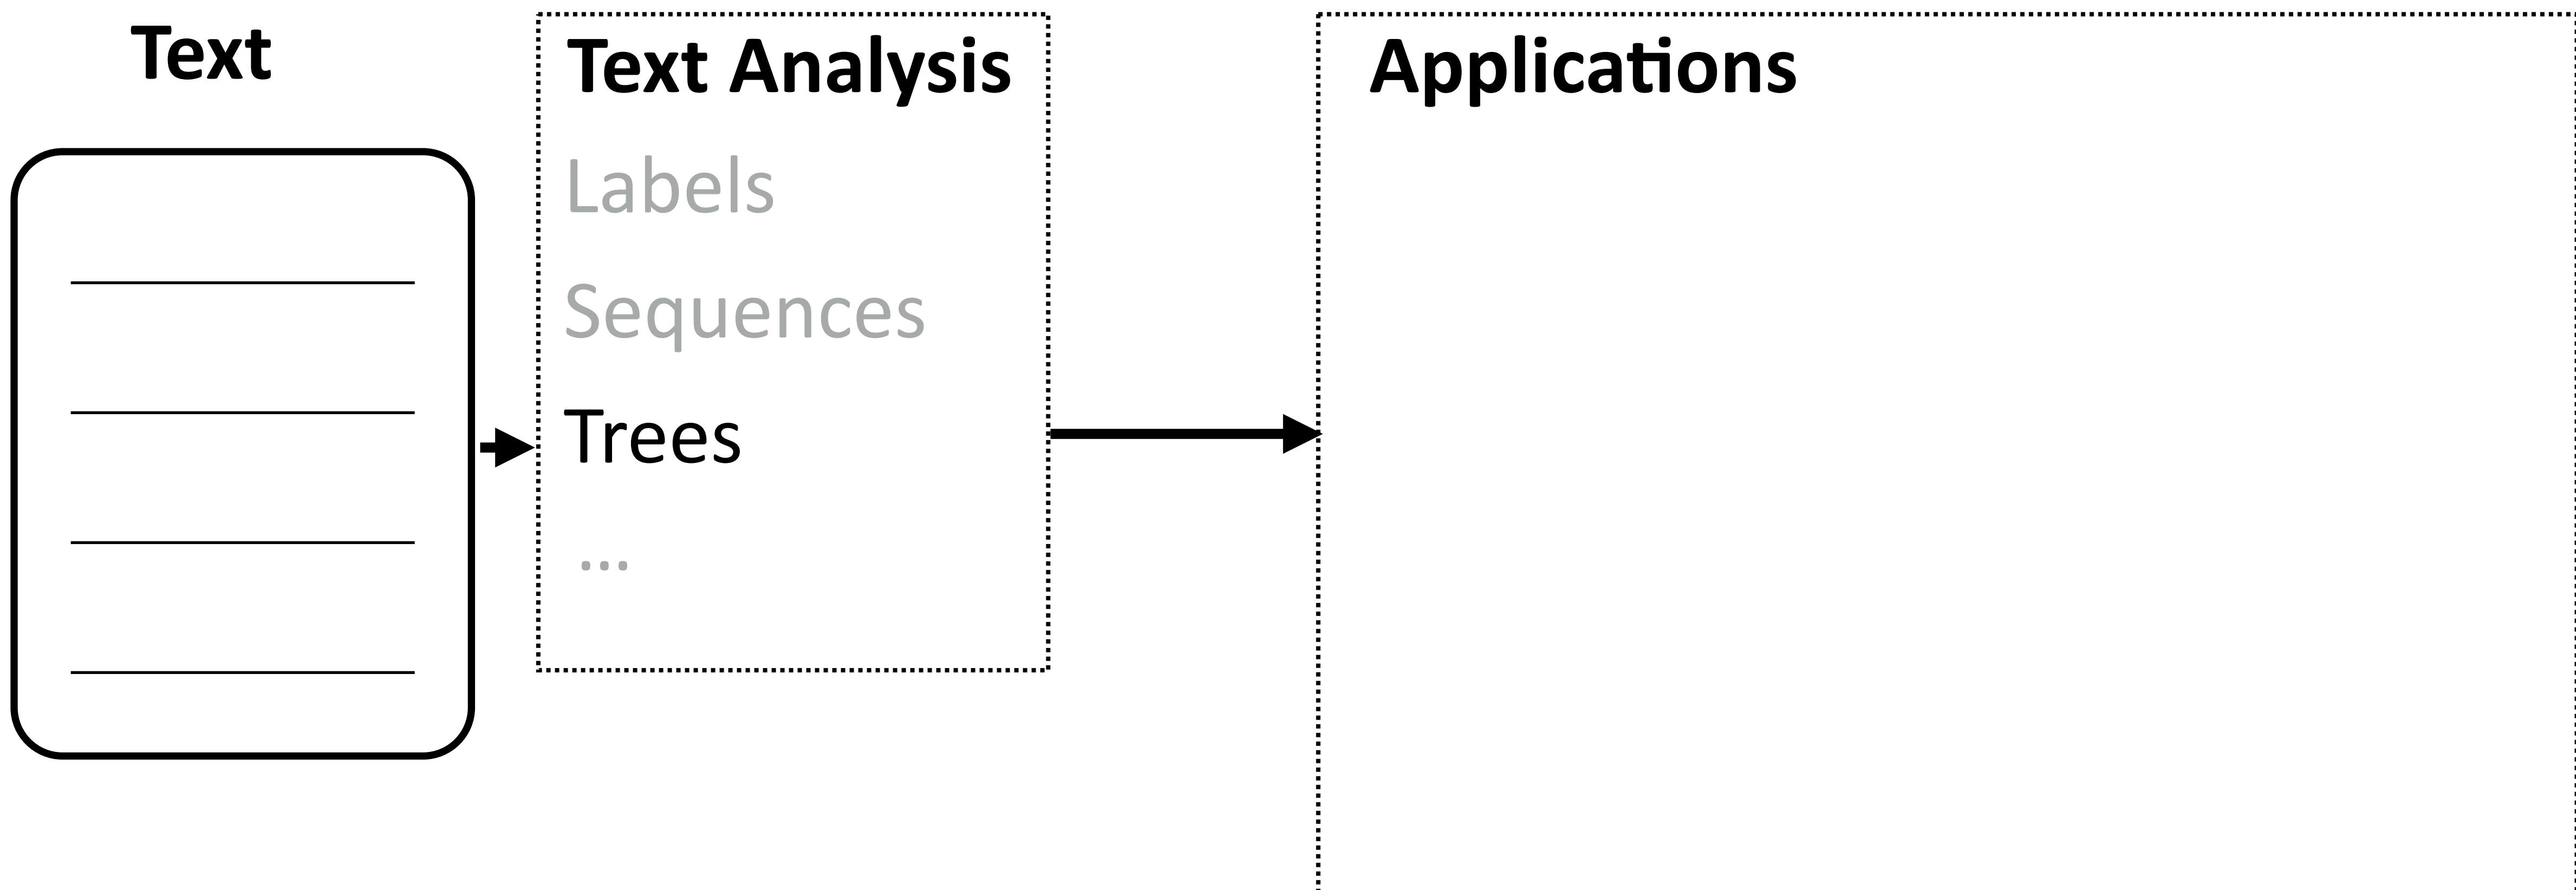
Sequences

Trees

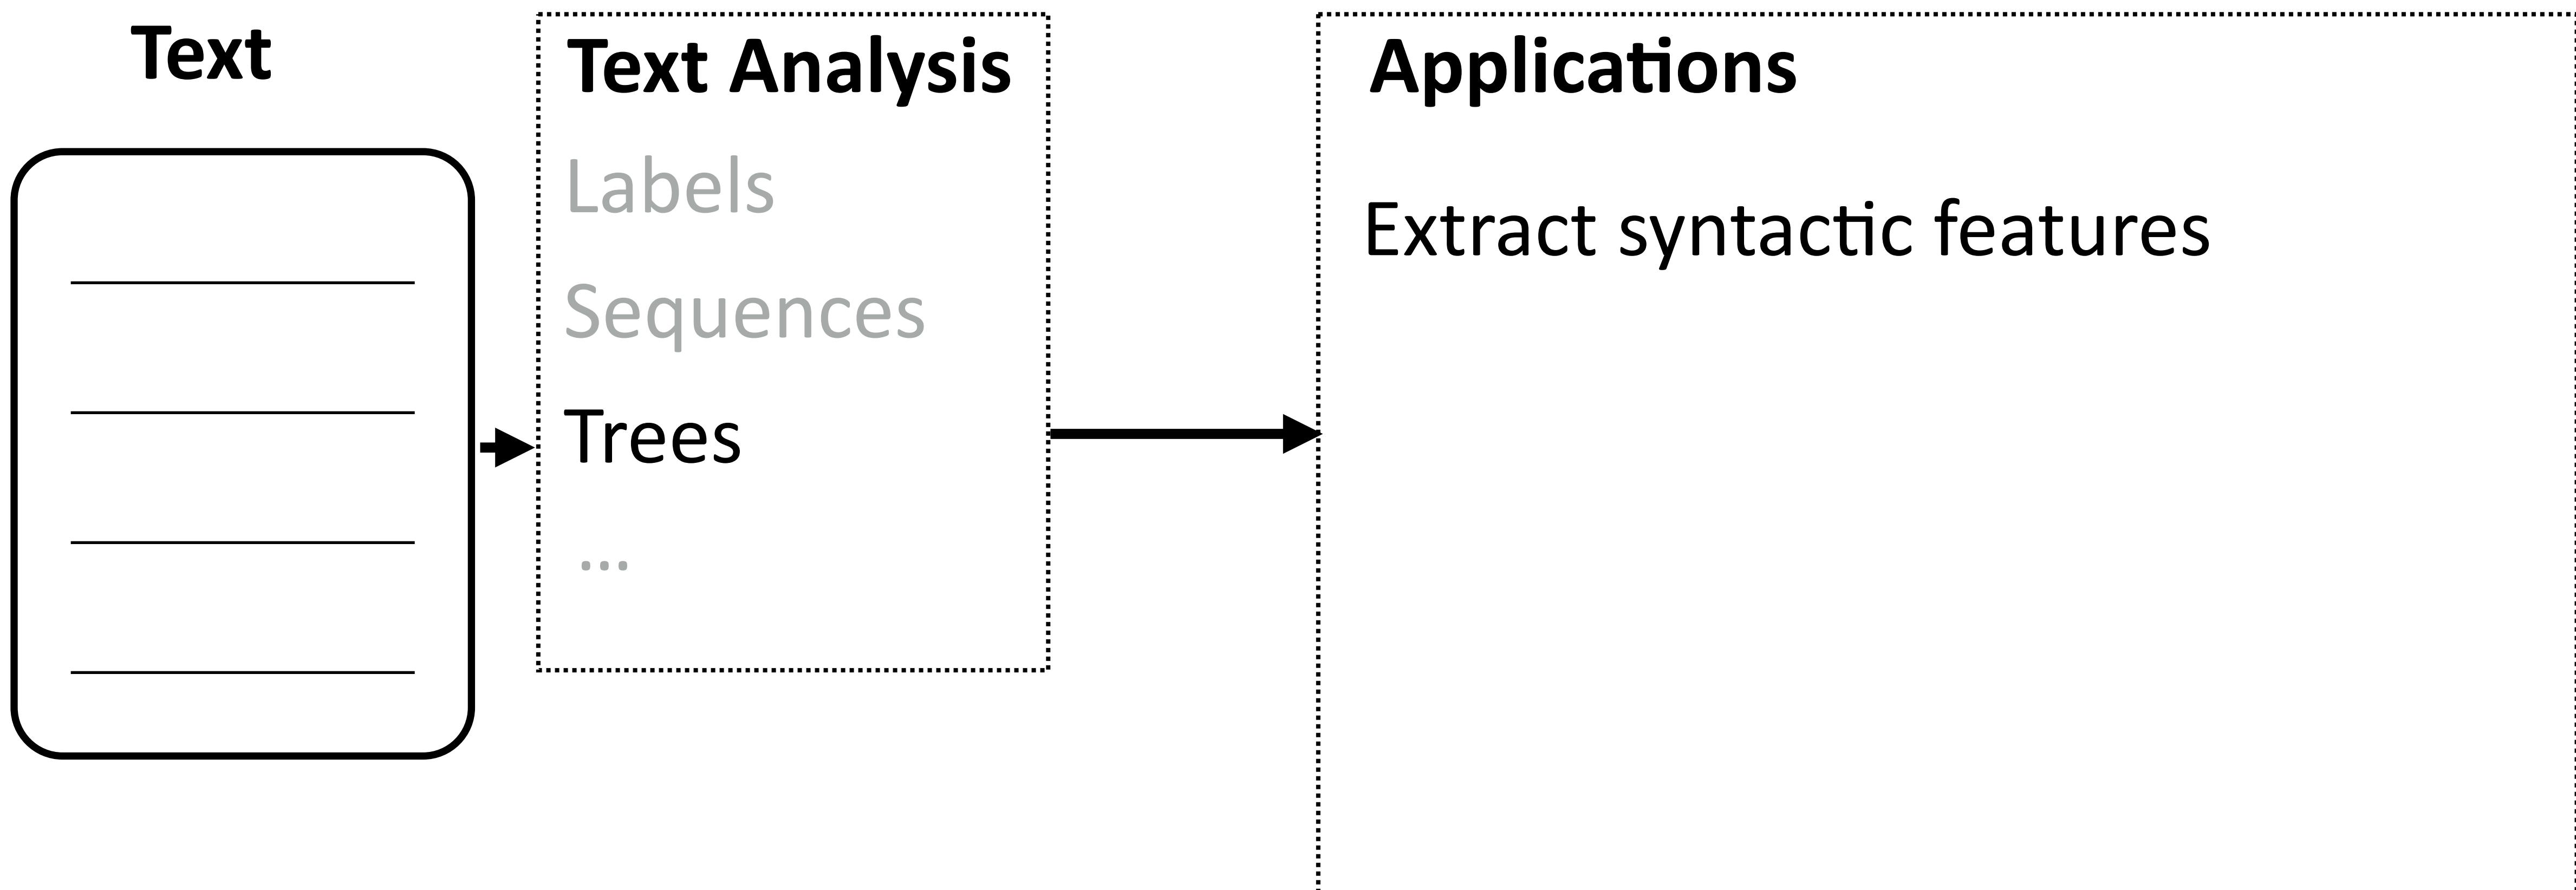
...



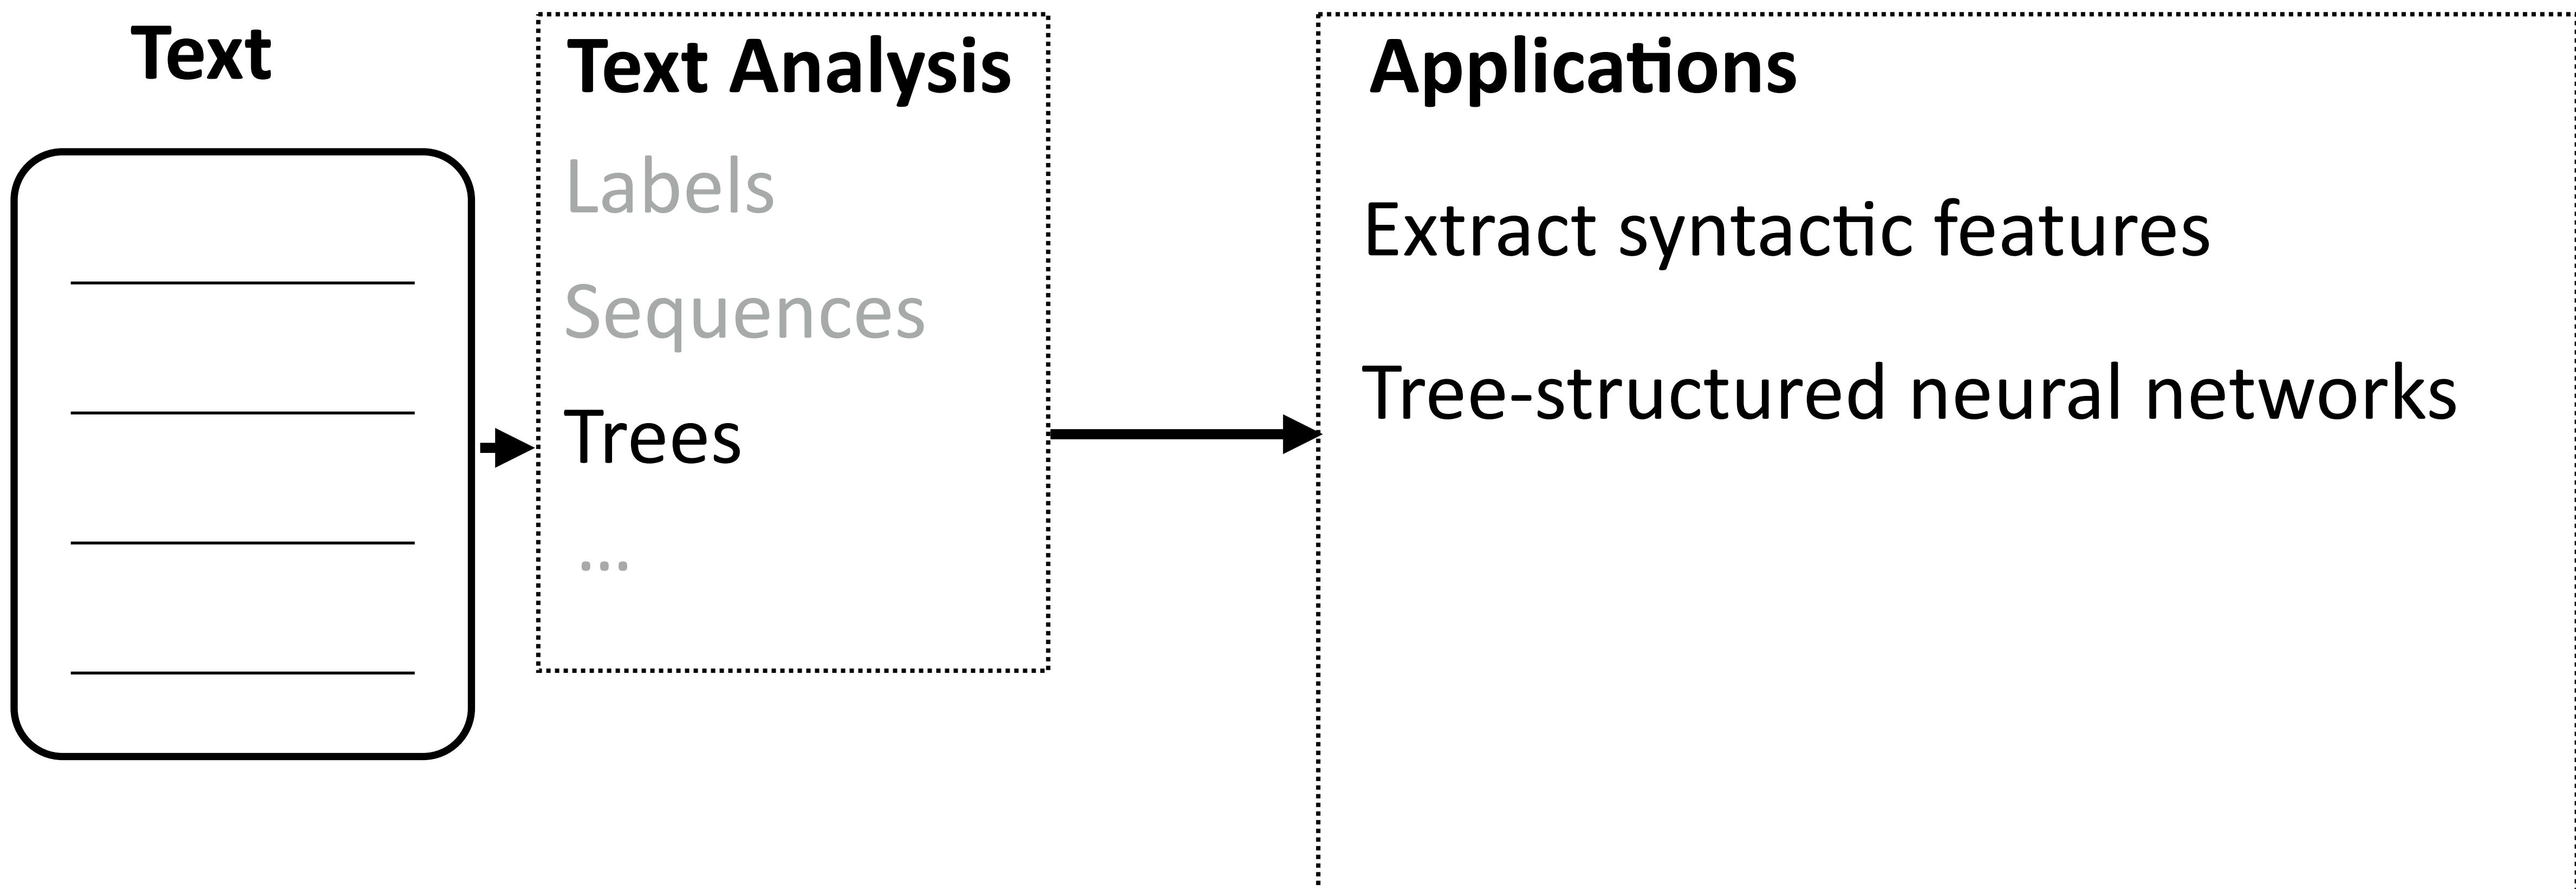
How do we use these representations?



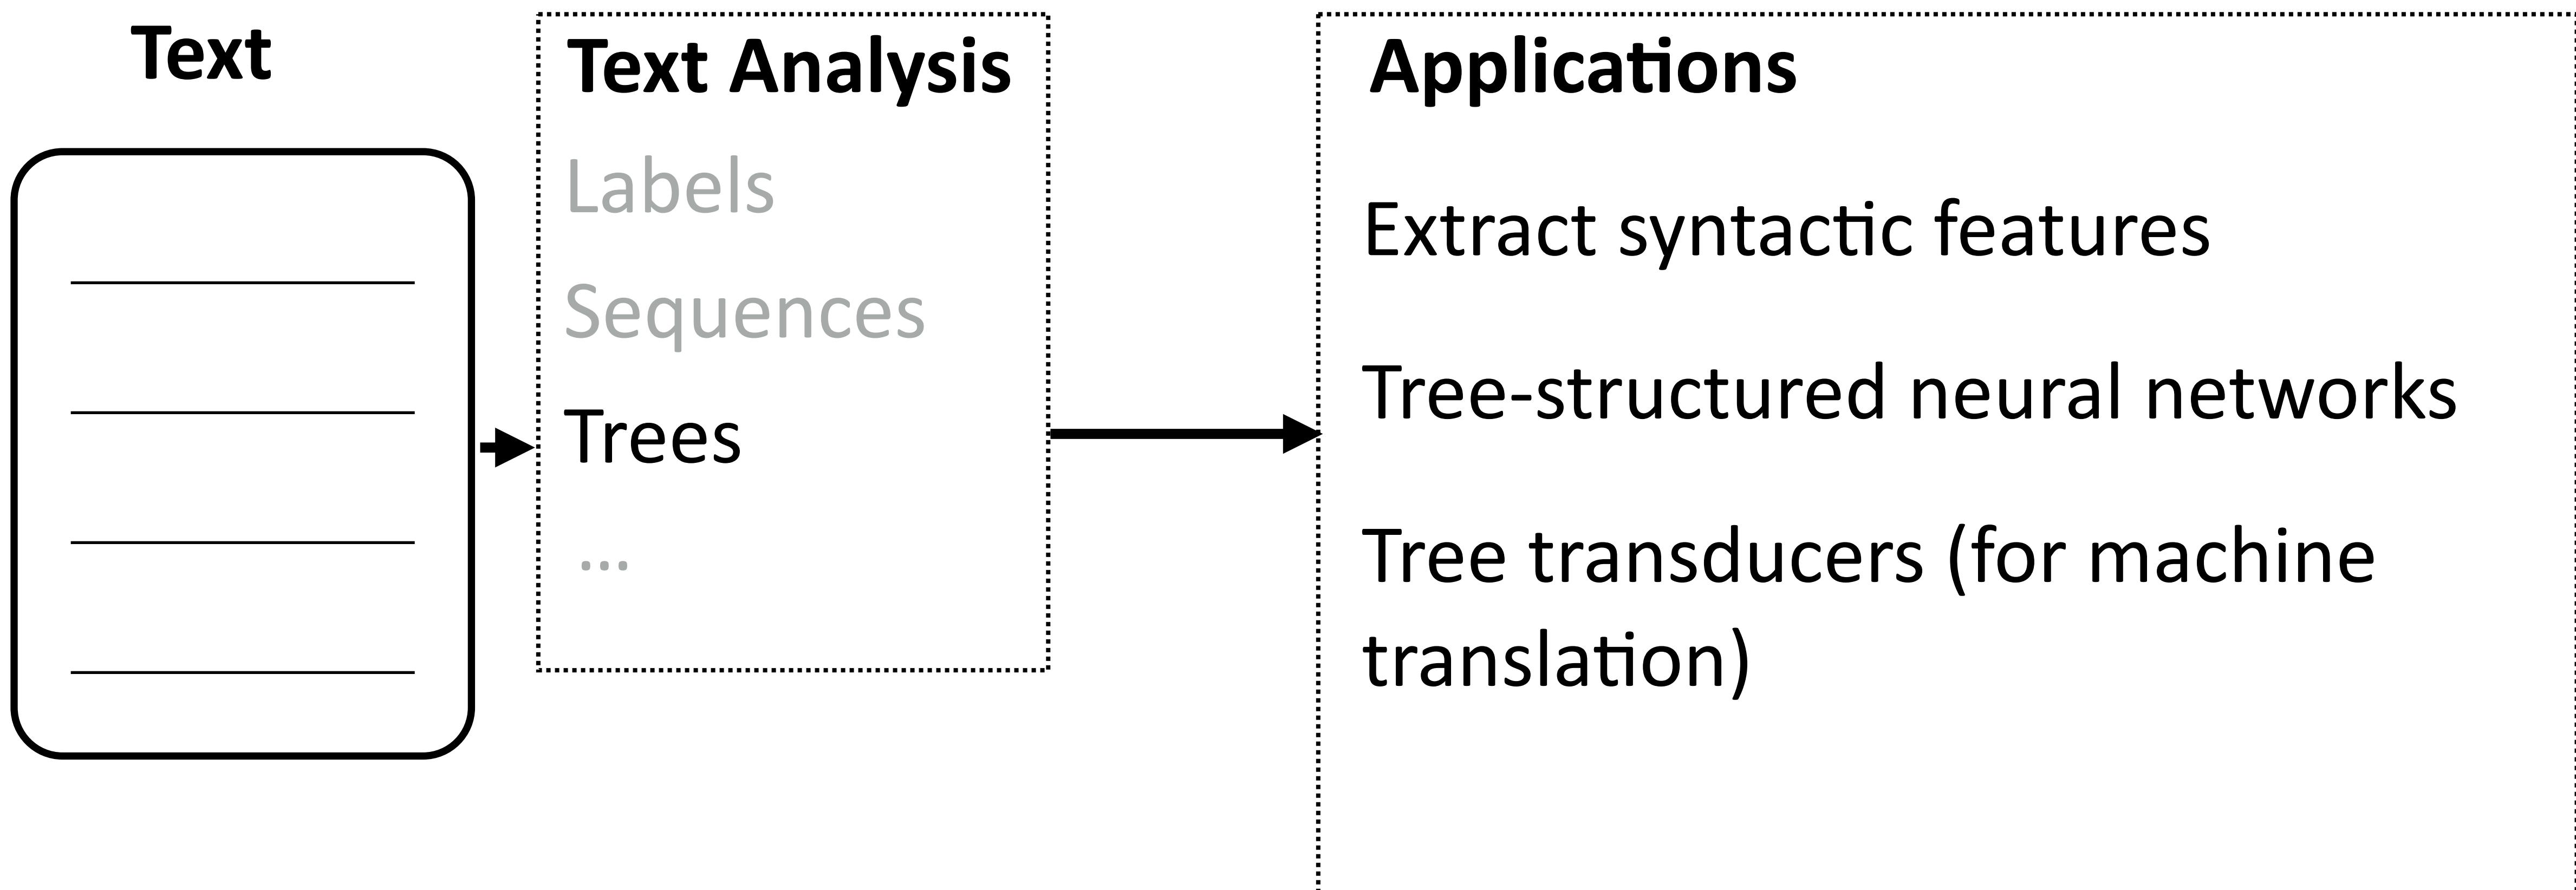
How do we use these representations?



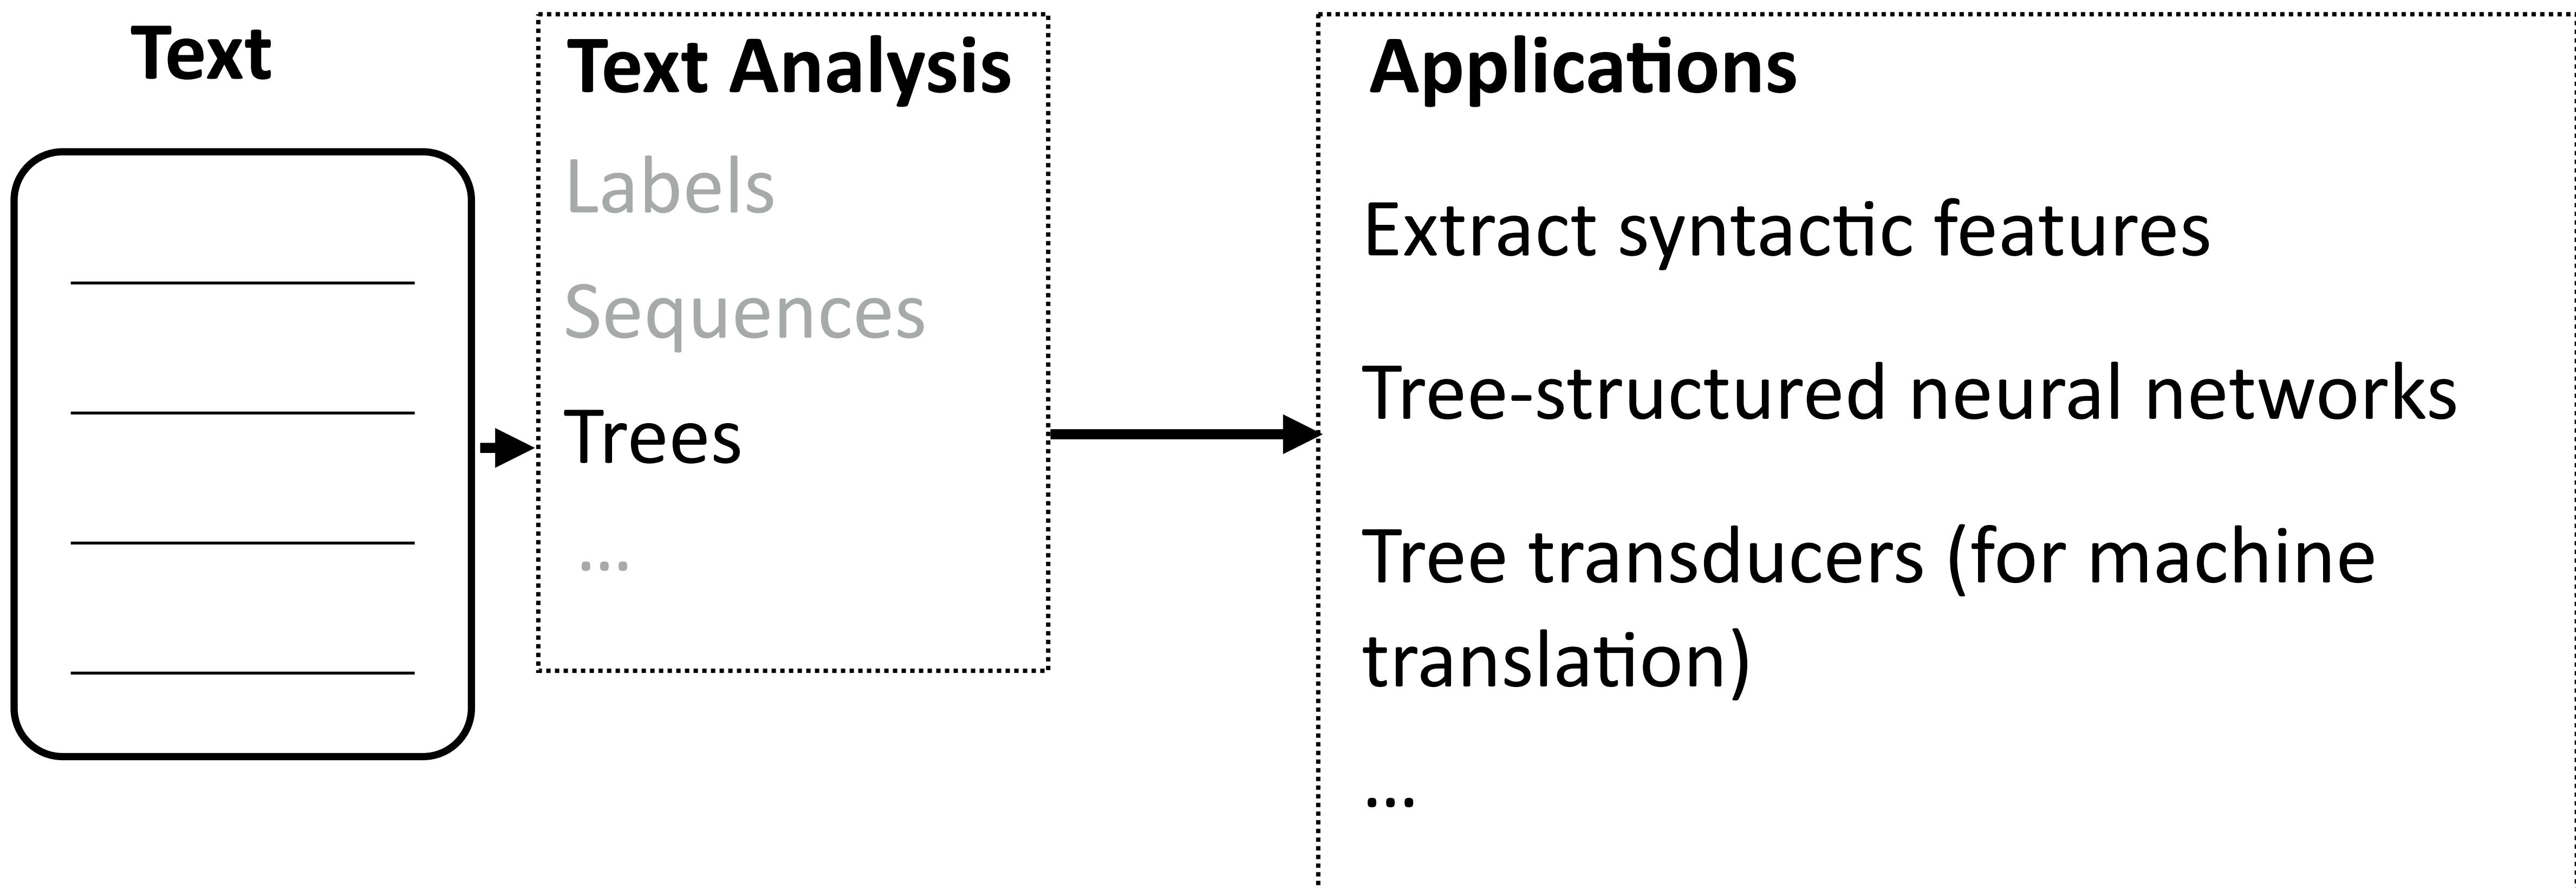
How do we use these representations?



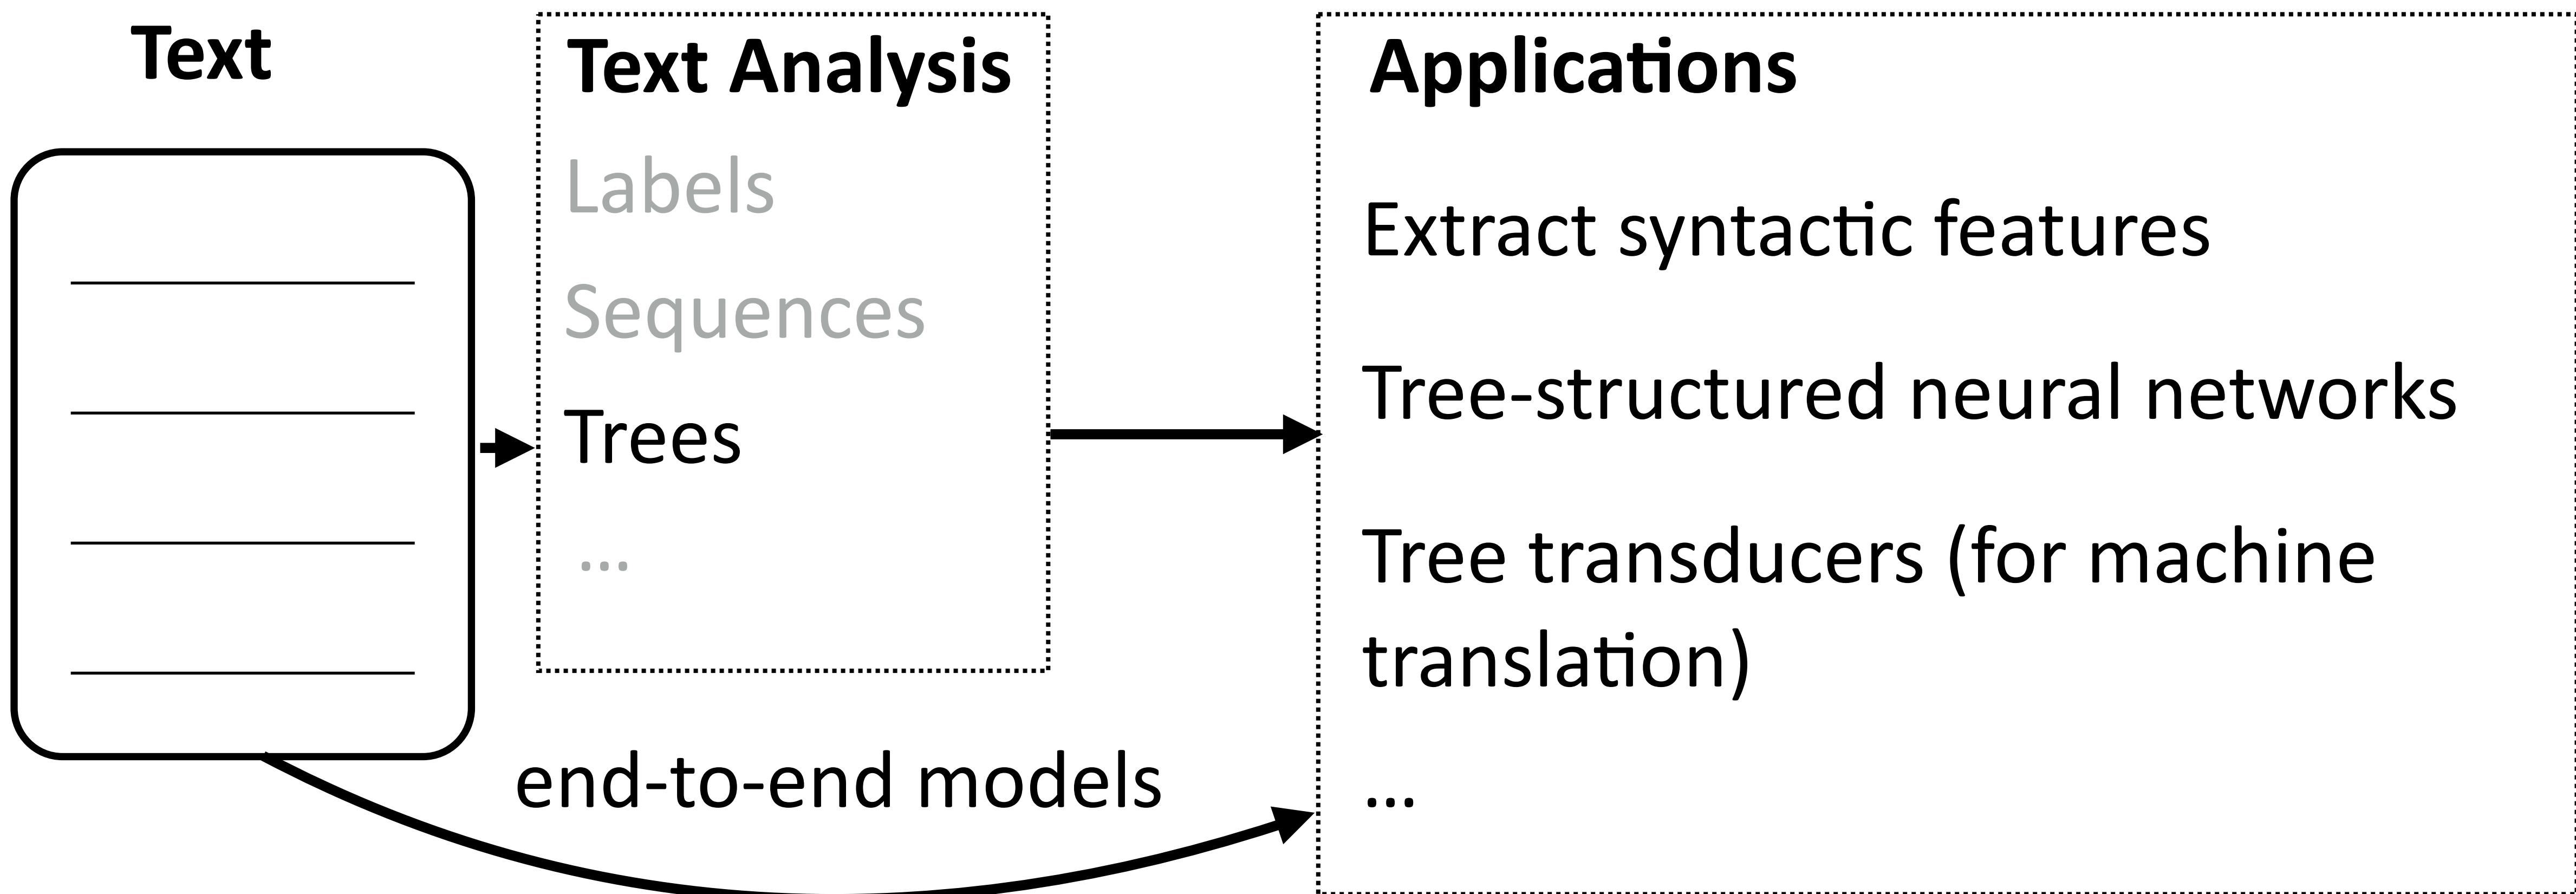
How do we use these representations?



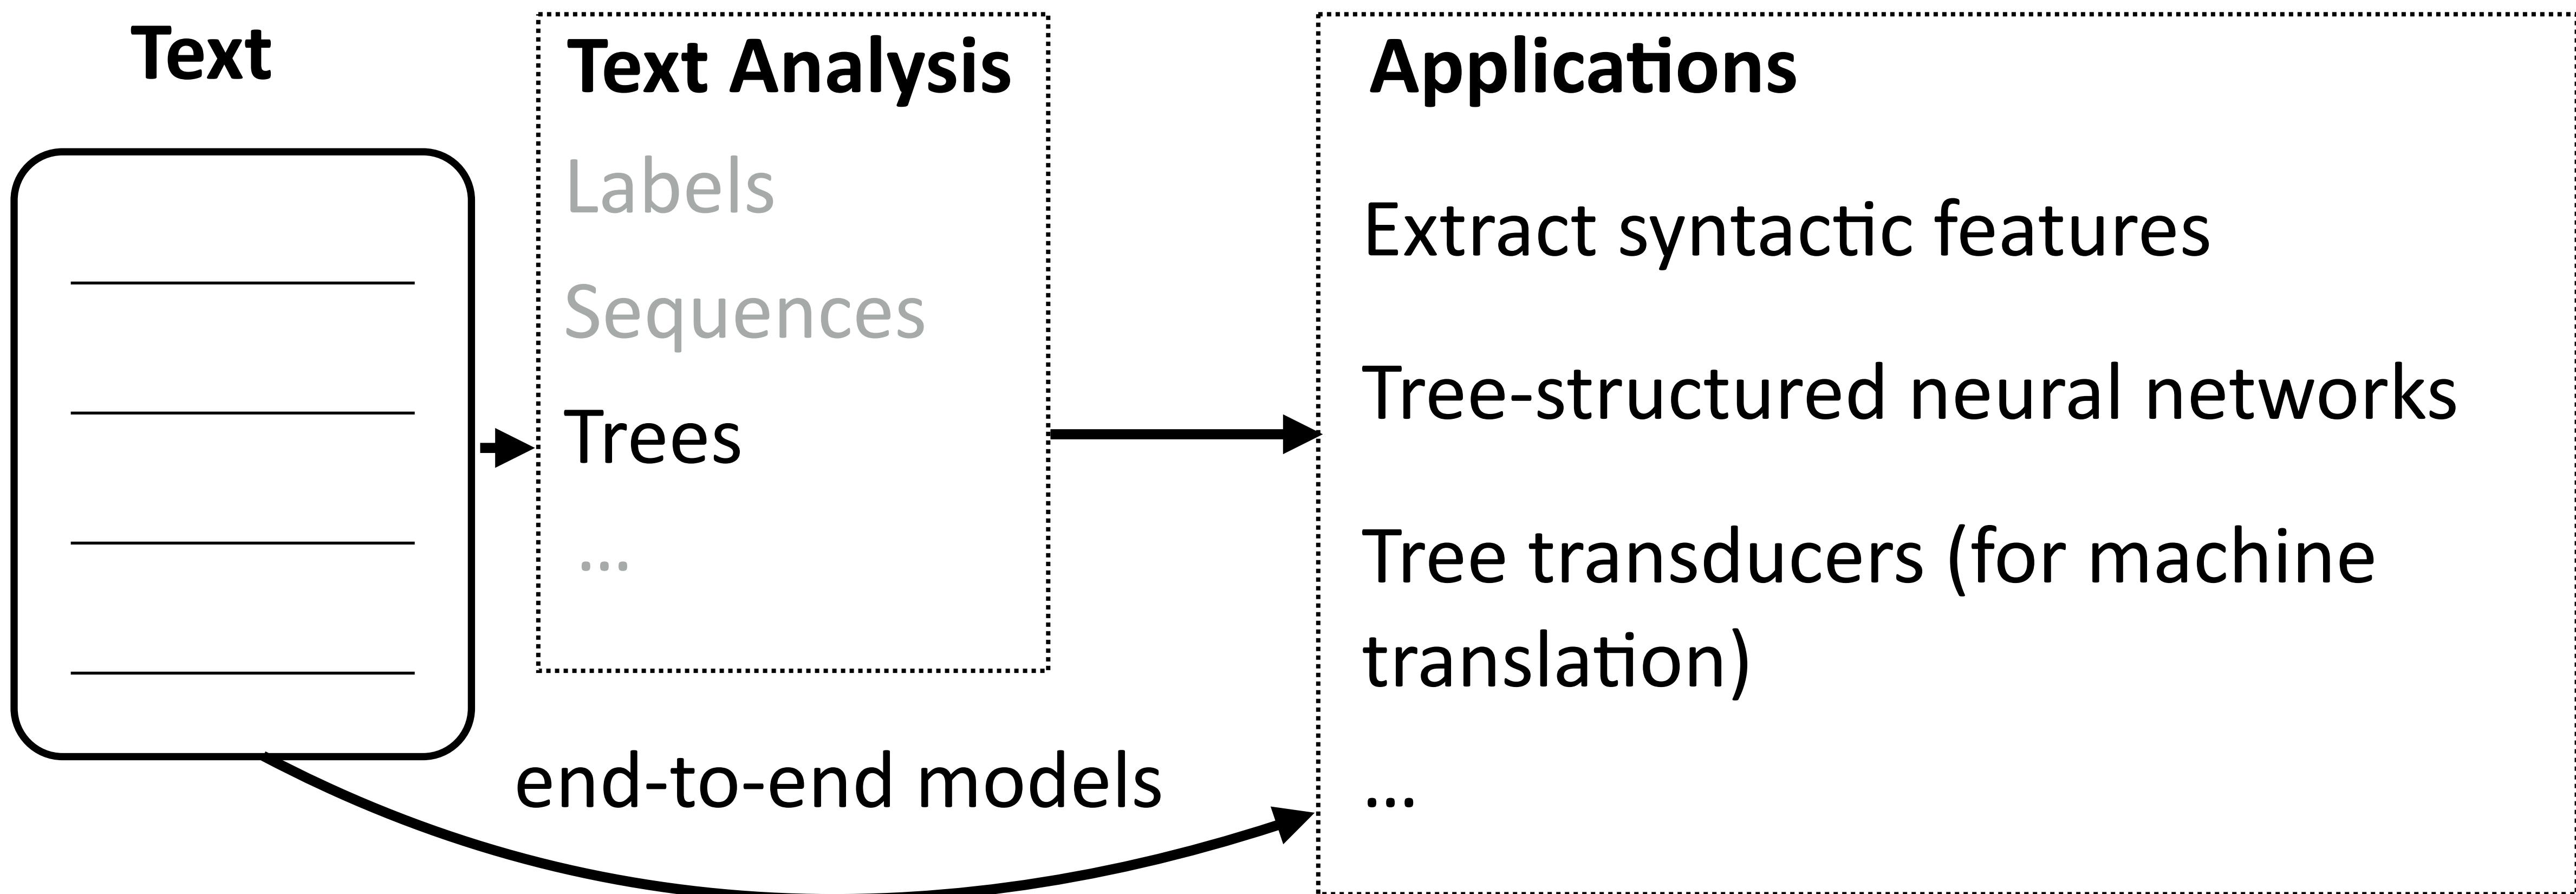
How do we use these representations?



How do we use these representations?

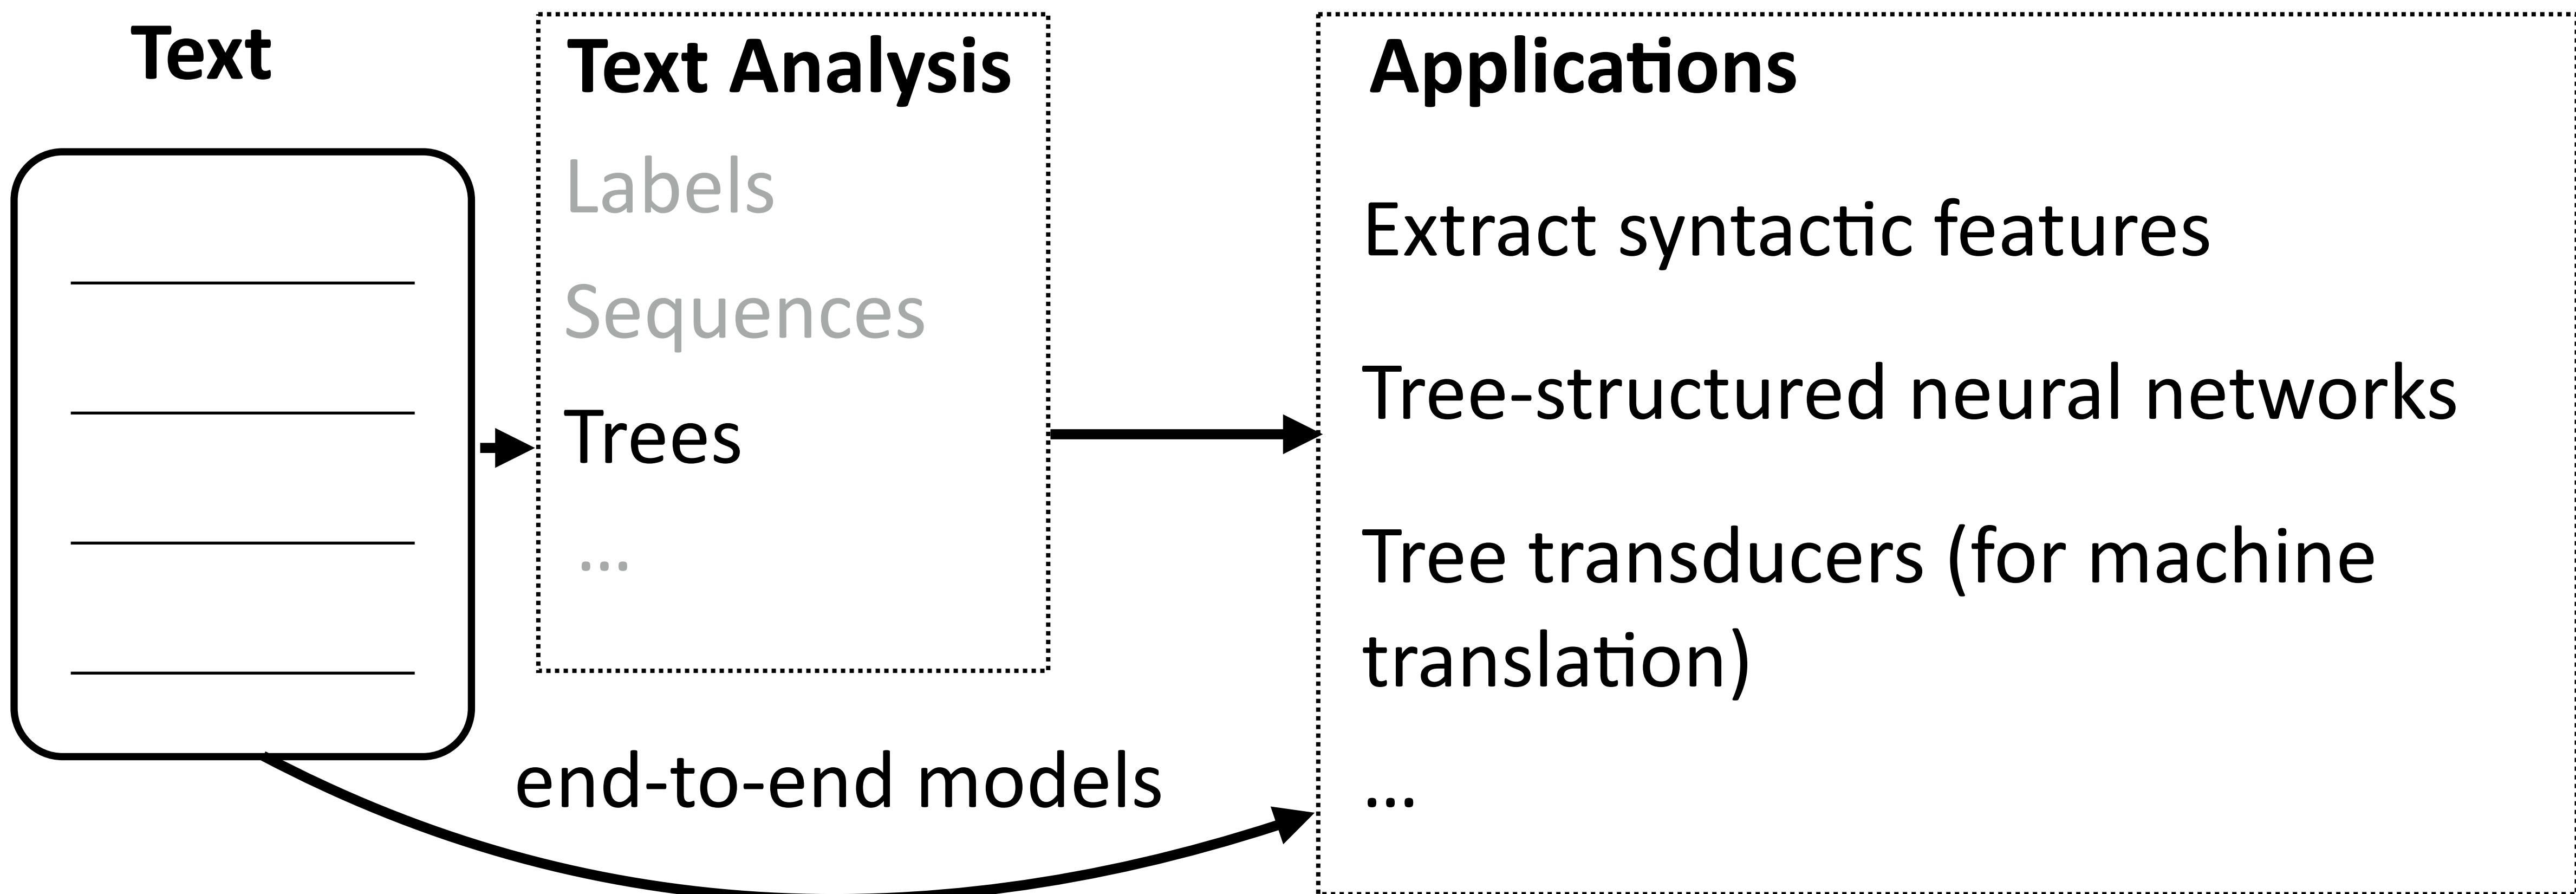


How do we use these representations?



- ▶ Main question: What representations do we need for language? What do we want to know about it?

How do we use these representations?



- ▶ Main question: What representations do we need for language? What do we want to know about it?
- ▶ Boils down to: what ambiguities do we need to resolve?

Why is language hard?
(and how can we handle that?)

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

The city council refused the demonstrators a permit because they _____ violence

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

The city council refused the demonstrators a permit because they _____ violence

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

The city council refused the demonstrators a permit because they _____ violence

they advocated

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

The city council refused the demonstrators a permit because they _____ violence

they advocated

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

The city council refused the demonstrators a permit because they _____ violence

they advocated

they feared

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

The city council refused the demonstrators a permit because they _____ violence

they advocated

they feared

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

The city council refused the demonstrators a permit because they _____ violence

they advocated

they feared

- ▶ This is so complicated that it's an AI challenge problem! (AI-complete)

Language is Ambiguous!

- ▶ Hector Levesque (2011): “Winograd schema challenge” (named after Terry Winograd, the creator of SHRDLU)

The city council refused the demonstrators a permit because they _____ violence

they advocated

they feared

- ▶ This is so complicated that it's an AI challenge problem! (AI-complete)
- ▶ Referential/semantic ambiguity

Language is Ambiguous!

Language is Ambiguous!

- ▶ Ambiguous News Headlines:

Language is Ambiguous!

- ▶ Ambiguous News Headlines:
 - ▶ Teacher Strikes Idle Kids

Language is Ambiguous!

- ▶ Ambiguous News Headlines:
 - ▶ Teacher Strikes Idle Kids
 - ▶ Hospitals Sued by 7 Foot Doctors

Language is Ambiguous!

- ▶ Ambiguous News Headlines:
 - ▶ Teacher Strikes Idle Kids
 - ▶ Hospitals Sued by 7 Foot Doctors
 - ▶ Ban on Nude Dancing on Governor's Desk

Language is Ambiguous!

- ▶ Ambiguous News Headlines:
 - ▶ Teacher Strikes Idle Kids
 - ▶ Hospitals Sued by 7 Foot Doctors
 - ▶ Ban on Nude Dancing on Governor's Desk
 - ▶ Iraqi Head Seeks Arms

Language is Ambiguous!

- ▶ Ambiguous News Headlines:
 - ▶ Teacher Strikes Idle Kids
 - ▶ Hospitals Sued by 7 Foot Doctors
 - ▶ Ban on Nude Dancing on Governor's Desk
 - ▶ Iraqi Head Seeks Arms
 - ▶ Stolen Painting Found by Tree

Language is Ambiguous!

- ▶ Ambiguous News Headlines:
 - ▶ Teacher Strikes Idle Kids
 - ▶ Hospitals Sued by 7 Foot Doctors
 - ▶ Ban on Nude Dancing on Governor's Desk
 - ▶ Iraqi Head Seeks Arms
 - ▶ Stolen Painting Found by Tree
 - ▶ Kids Make Nutritious Snacks

Language is Ambiguous!

- ▶ Ambiguous News Headlines:
 - ▶ Teacher Strikes Idle Kids
 - ▶ Hospitals Sued by 7 Foot Doctors
 - ▶ Ban on Nude Dancing on Governor's Desk
 - ▶ Iraqi Head Seeks Arms
 - ▶ Stolen Painting Found by Tree
 - ▶ Kids Make Nutritious Snacks
 - ▶ Local HS Dropouts Cut in Half

Language is Ambiguous!

- ▶ Ambiguous News Headlines:
 - ▶ Teacher Strikes Idle Kids
 - ▶ Hospitals Sued by 7 Foot Doctors
 - ▶ Ban on Nude Dancing on Governor's Desk
 - ▶ Iraqi Head Seeks Arms
 - ▶ Stolen Painting Found by Tree
 - ▶ Kids Make Nutritious Snacks
 - ▶ Local HS Dropouts Cut in Half
- ▶ Syntactic/semantic ambiguity: parsing needed to resolve these, but need context to figure out which parse is correct

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

il fait vraiment beau 

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

It is really nice out

il fait vraiment beau



Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

il fait vraiment beau

It is really nice out
It's really nice

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

il fait vraiment beau



It is really nice out
It's really nice
The weather is beautiful

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

il fait vraiment beau



It is really nice out
It's really nice
The weather is beautiful
It is really beautiful outside

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

il fait vraiment beau



It is really nice out
It's really nice
The weather is beautiful
It is really beautiful outside
He makes truly beautiful

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

il fait vraiment beau



It is really nice out
It's really nice
The weather is beautiful
It is really beautiful outside
He makes truly beautiful
He makes truly boyfriend

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

il fait vraiment beau



It is really nice out
It's really nice
The weather is beautiful
It is really beautiful outside
He makes truly beautiful
He makes truly boyfriend
It fact actually handsome

Language is Really Ambiguous!

- ▶ There aren't just one or two possibilities which are resolved pragmatically

il fait vraiment beau



It is really nice out
It's really nice
The weather is beautiful
It is really beautiful outside
He makes truly beautiful
He makes truly boyfriend
It fact actually handsome

- ▶ Combinatorially many possibilities, many you won't even register as ambiguities, but systems still have to resolve them

What do we need to understand language?

- ▶ Lots of data!

SOURCE	Cela constituerait une solution transitoire qui permettrait de conduire à terme à une charte à valeur contraignante.
HUMAN	That would be an interim solution which would make it possible to work towards a binding charter in the long term .
1x DATA	[this] [constituerait] [assistance] [transitoire] [who] [permettrait] [licences] [to] [terme] [to] [a] [charter] [to] [value] [contraignante] [.]
10x DATA	[it] [would] [a solution] [transitional] [which] [would] [of] [lead] [to] [term] [to a] [charter] [to] [value] [binding] [.]
100x DATA	[this] [would be] [a transitional solution] [which would] [lead to] [a charter] [legally binding] [.]
1000x DATA	[that would be] [a transitional solution] [which would] [eventually lead to] [a binding charter] [.]

What do we need to understand language?

- ▶ World knowledge: have access to information beyond the training data

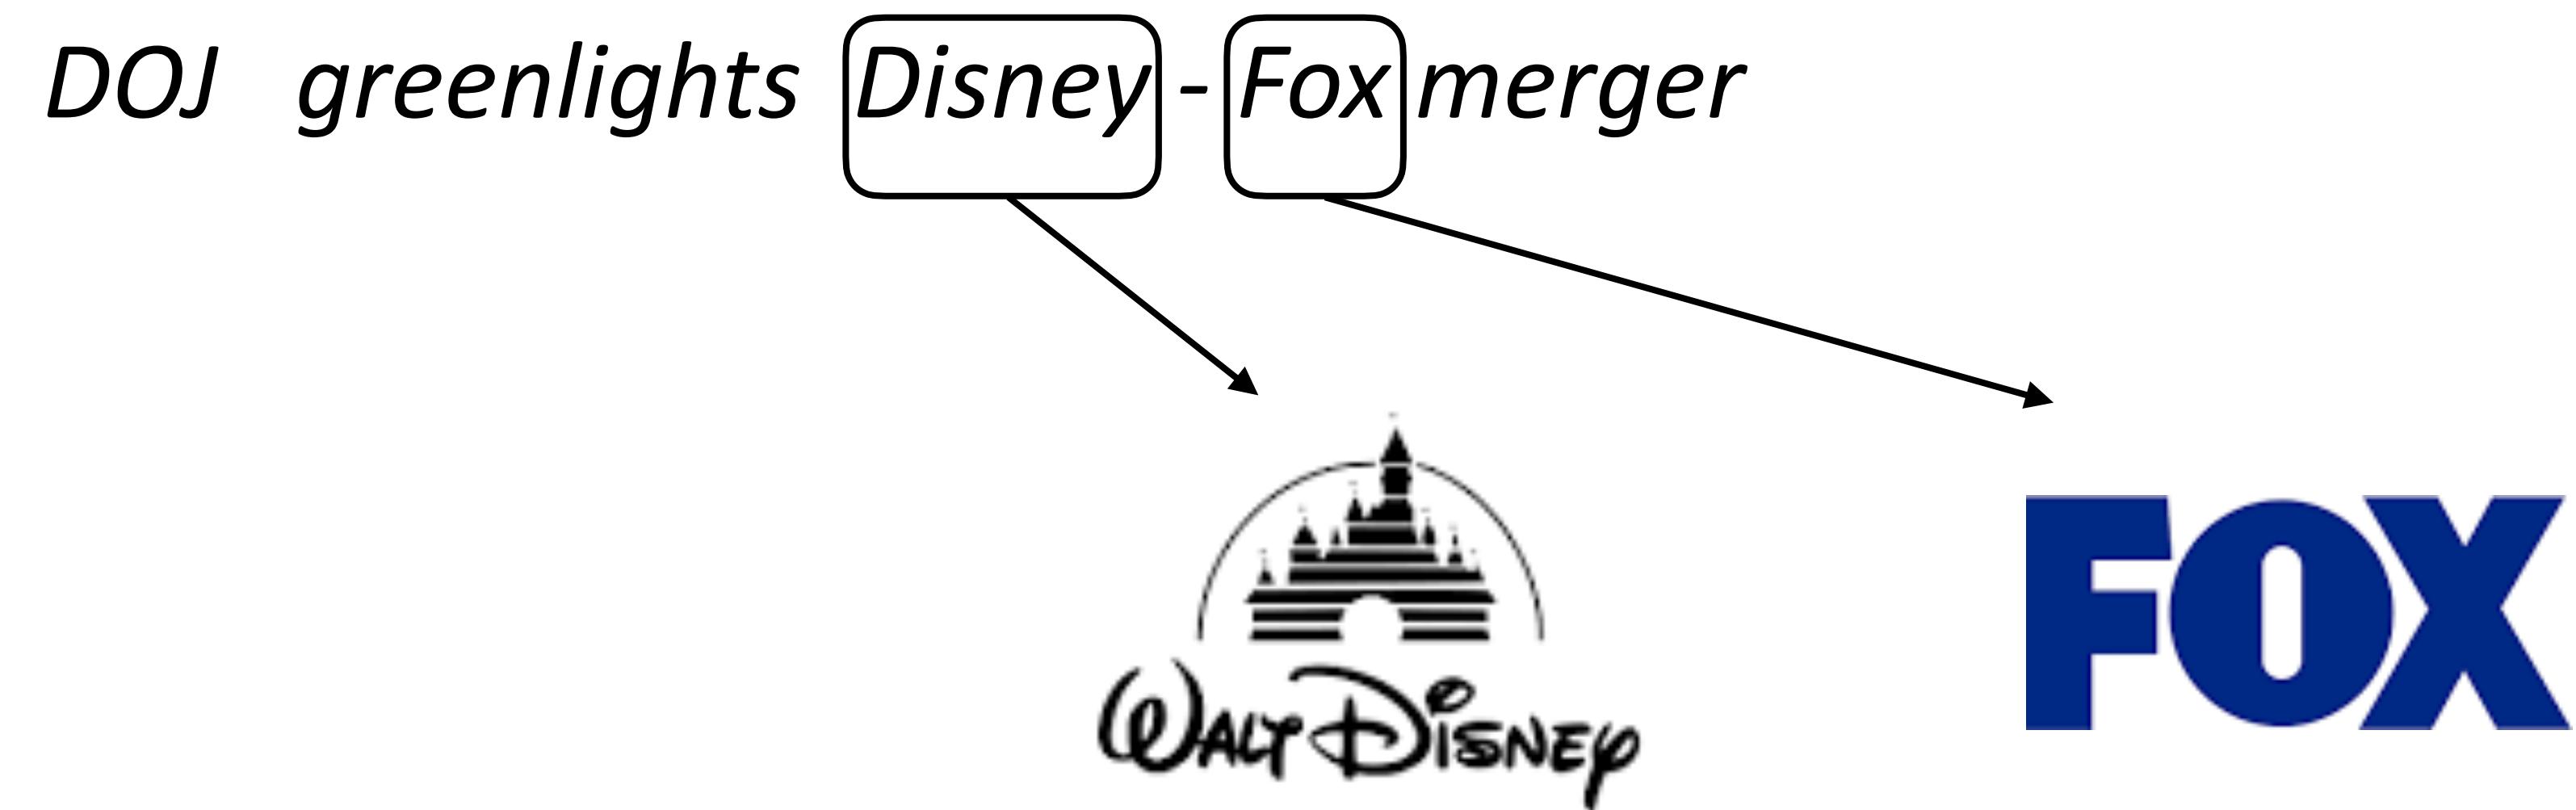
What do we need to understand language?

- ▶ World knowledge: have access to information beyond the training data

DOJ greenlights Disney - Fox merger

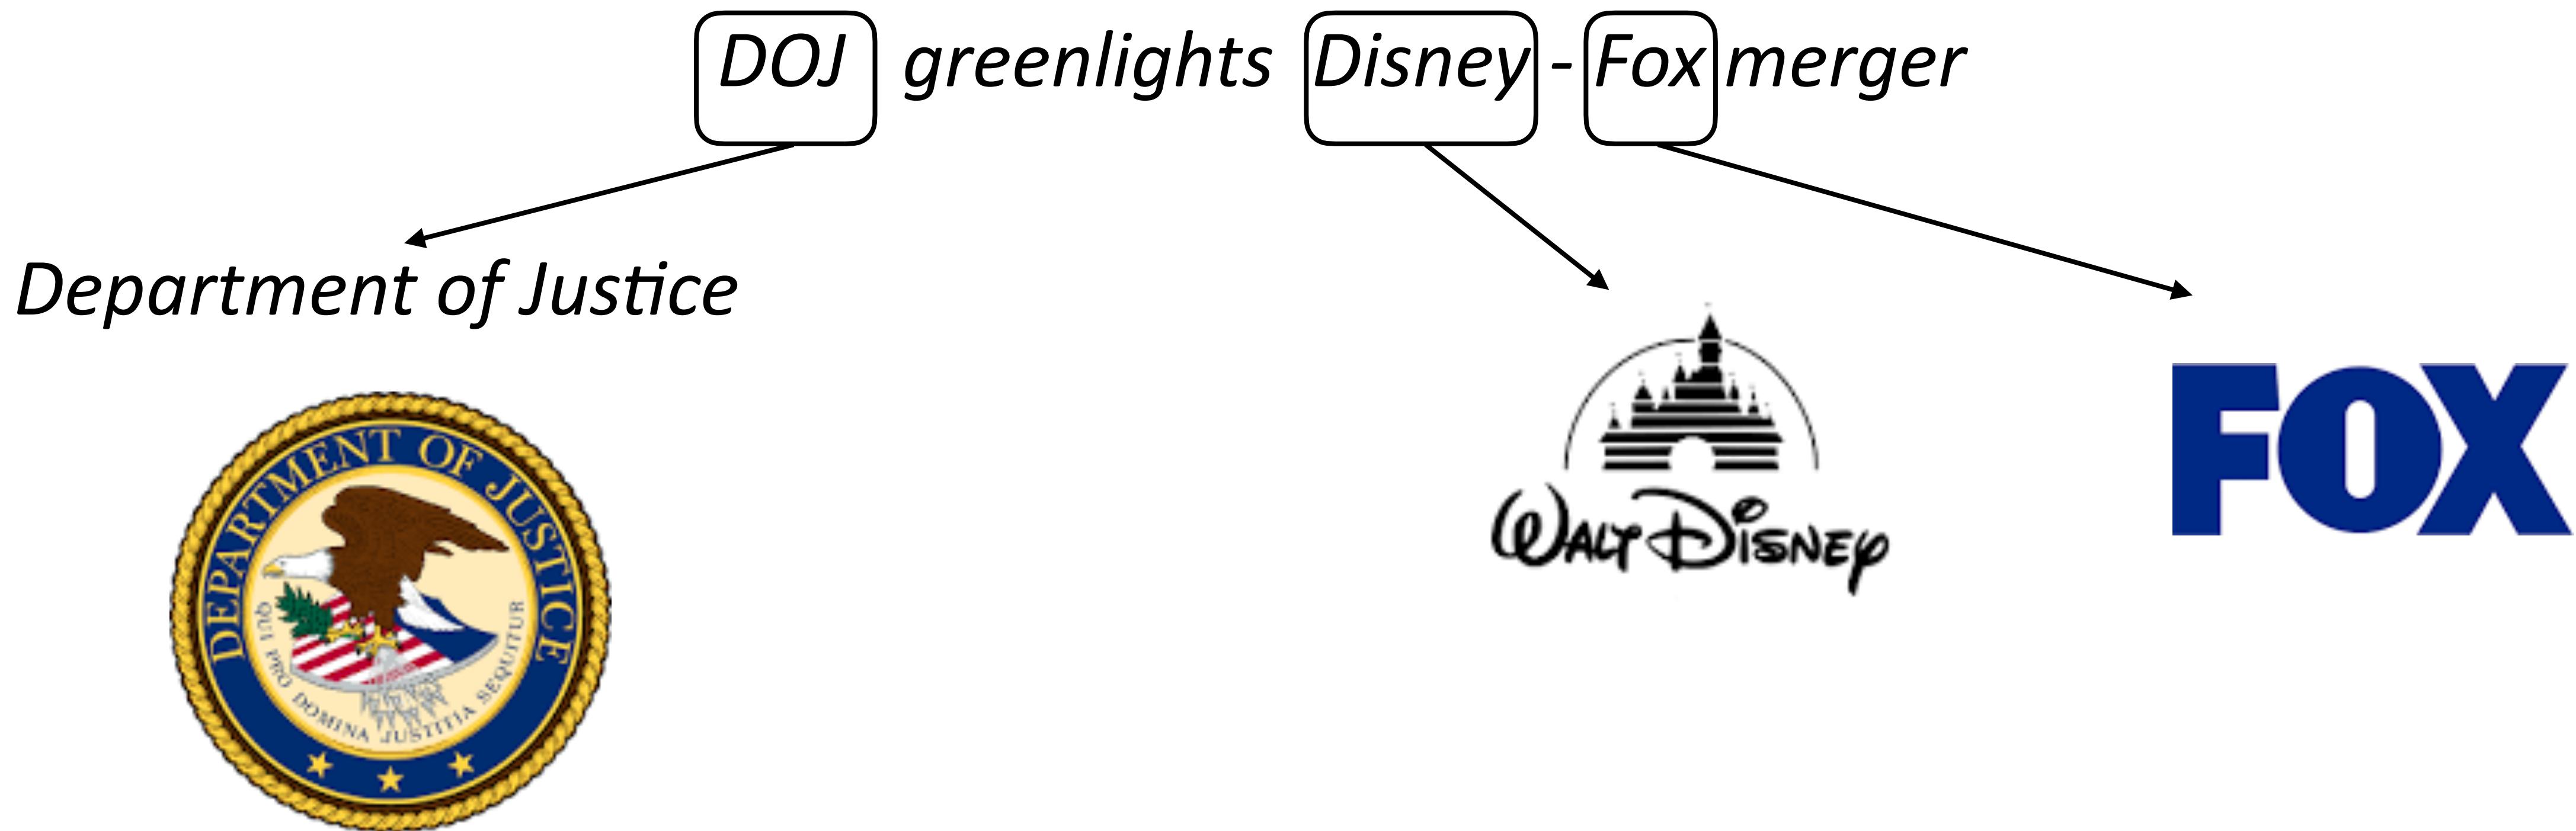
What do we need to understand language?

- ▶ World knowledge: have access to information beyond the training data



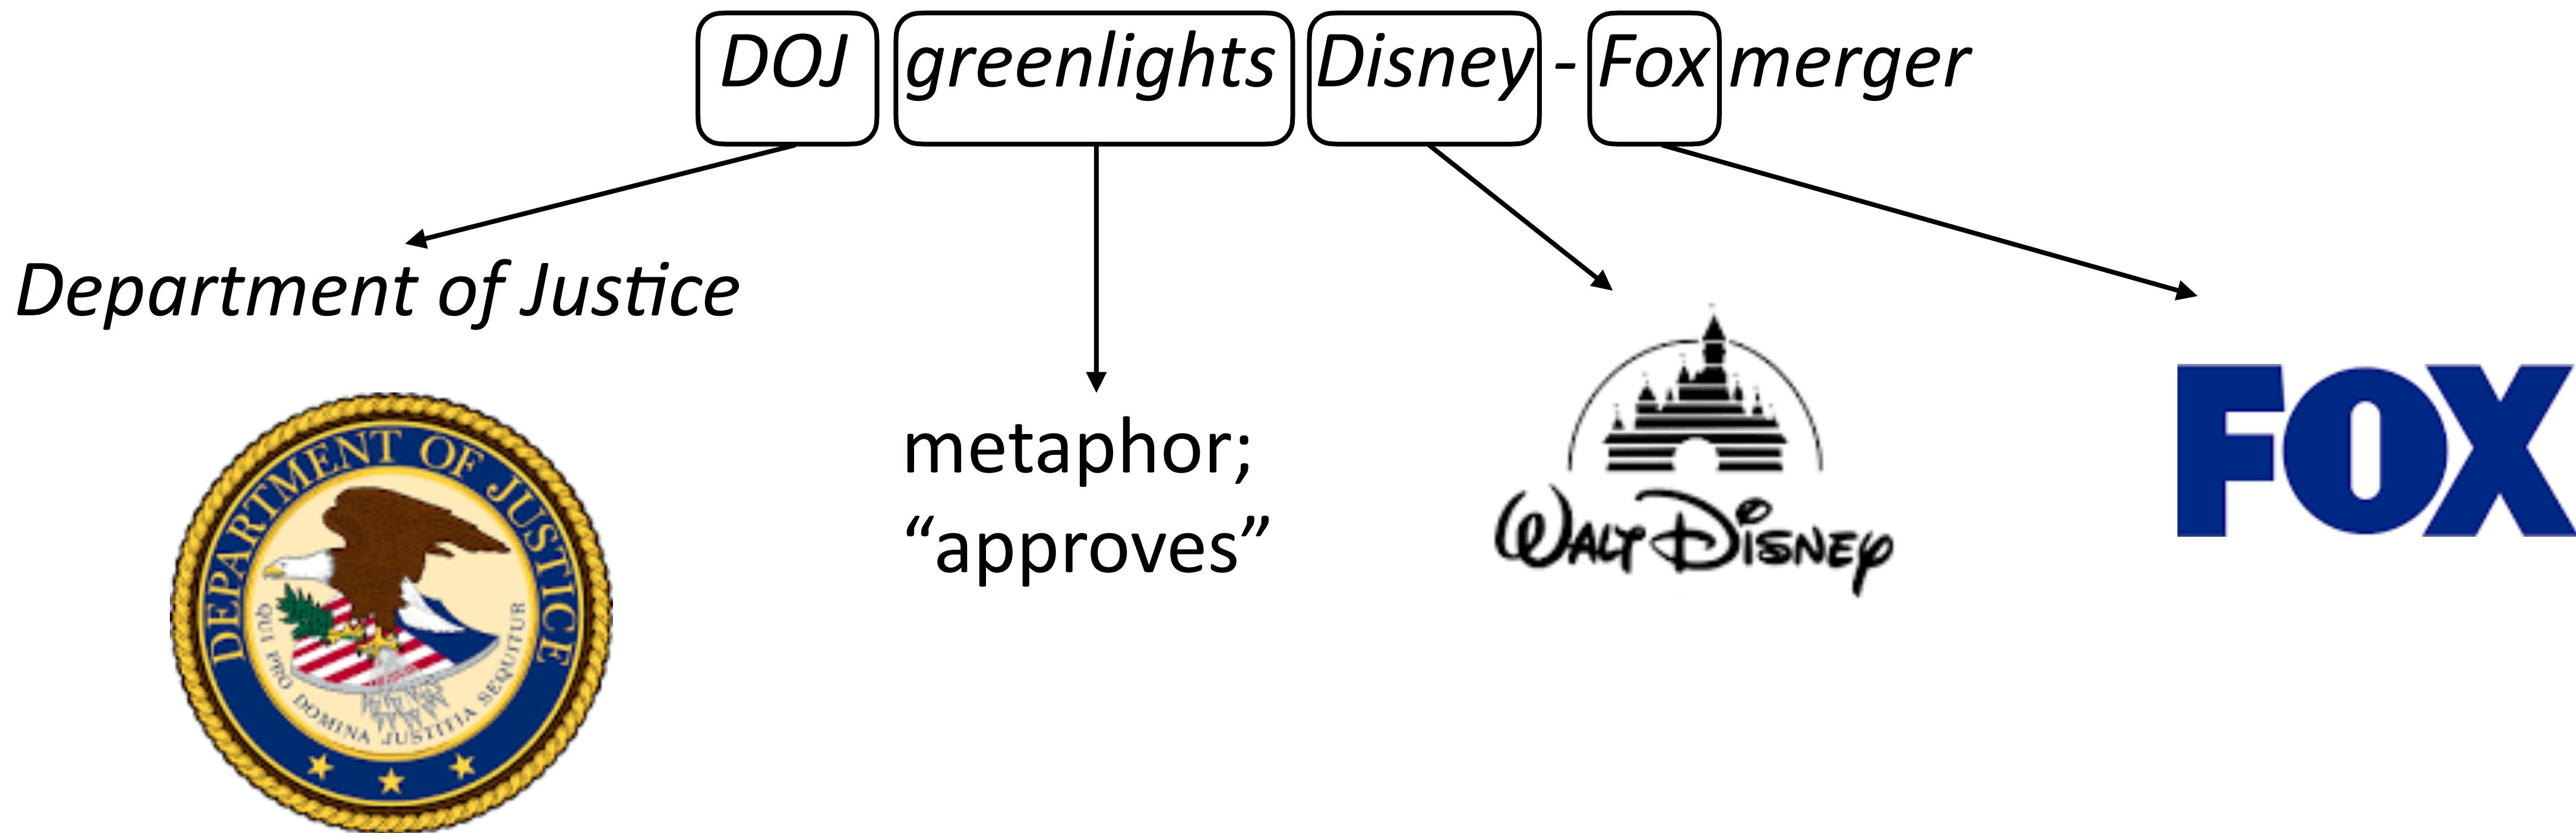
What do we need to understand language?

- ▶ World knowledge: have access to information beyond the training data



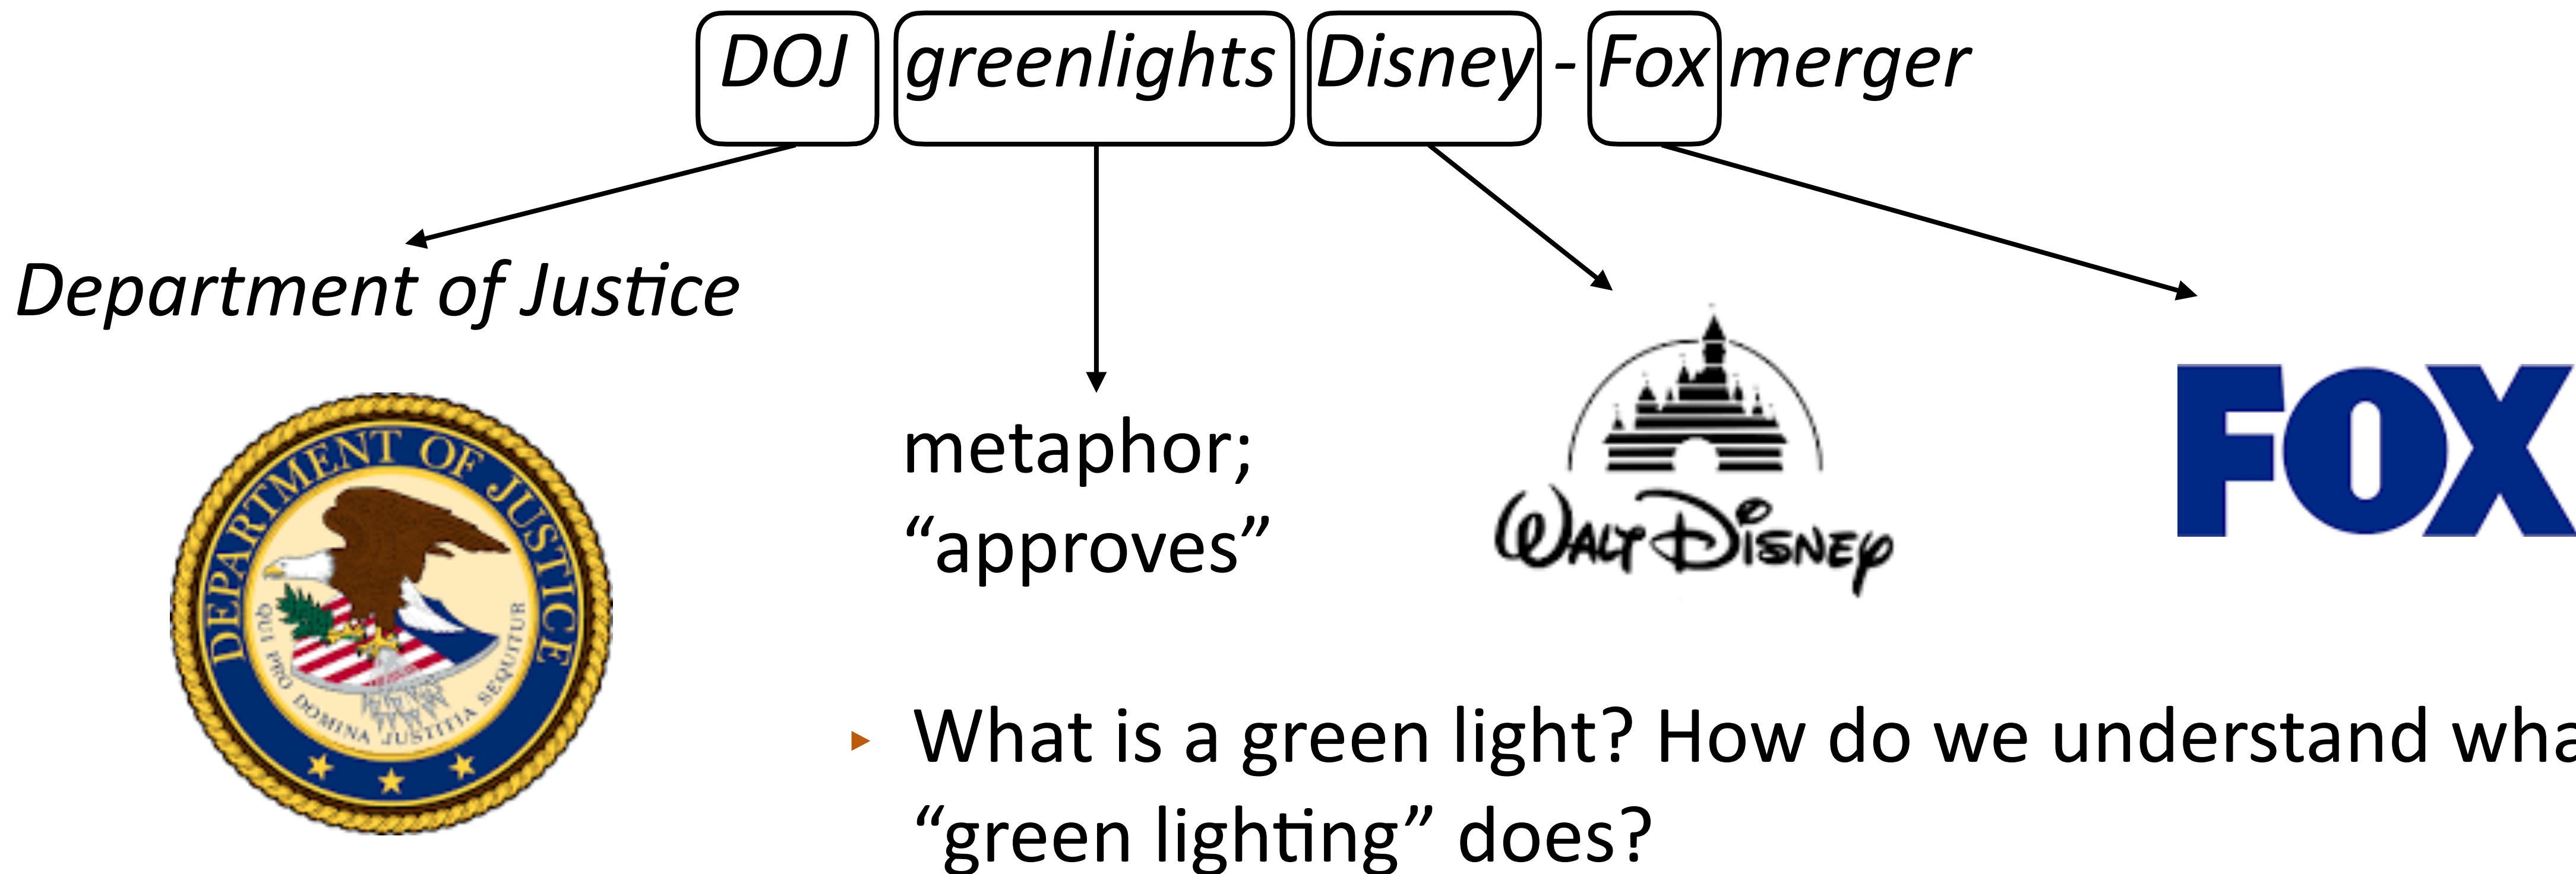
What do we need to understand language?

- ▶ World knowledge: have access to information beyond the training data



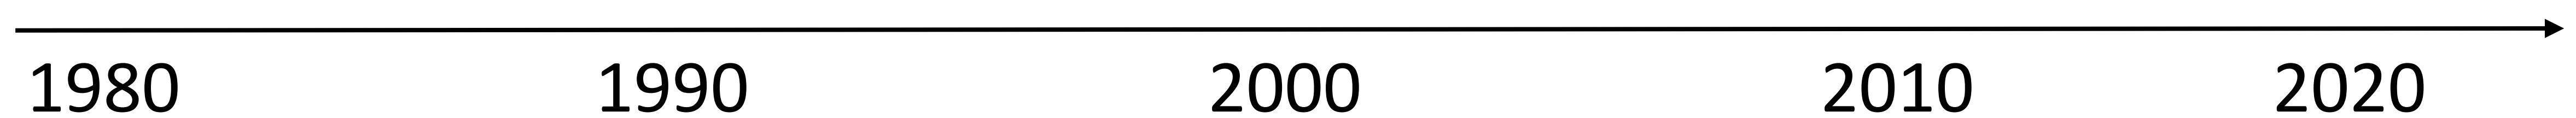
What do we need to understand language?

- ▶ World knowledge: have access to information beyond the training data

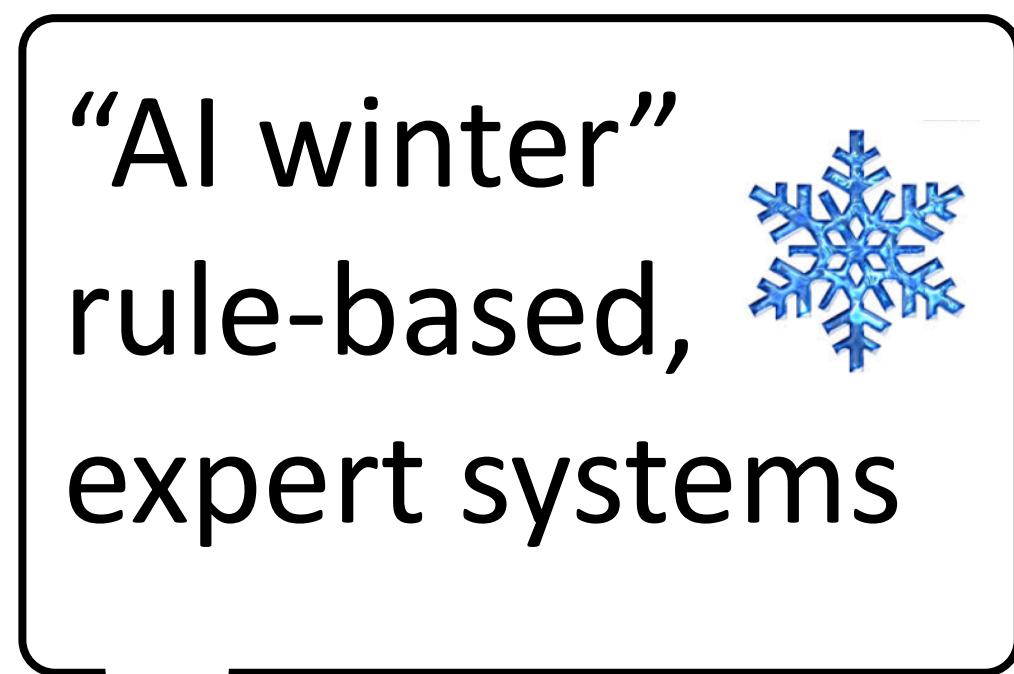


What techniques do we use?
(to combine data, knowledge, linguistics, etc.)

A brief history of (modern) NLP



A brief history of (modern) NLP



1980

1990

2000

2010

2020

A brief history of (modern) NLP

“AI winter”
rule-based,
expert systems



earliest stat MT
work at IBM



1980

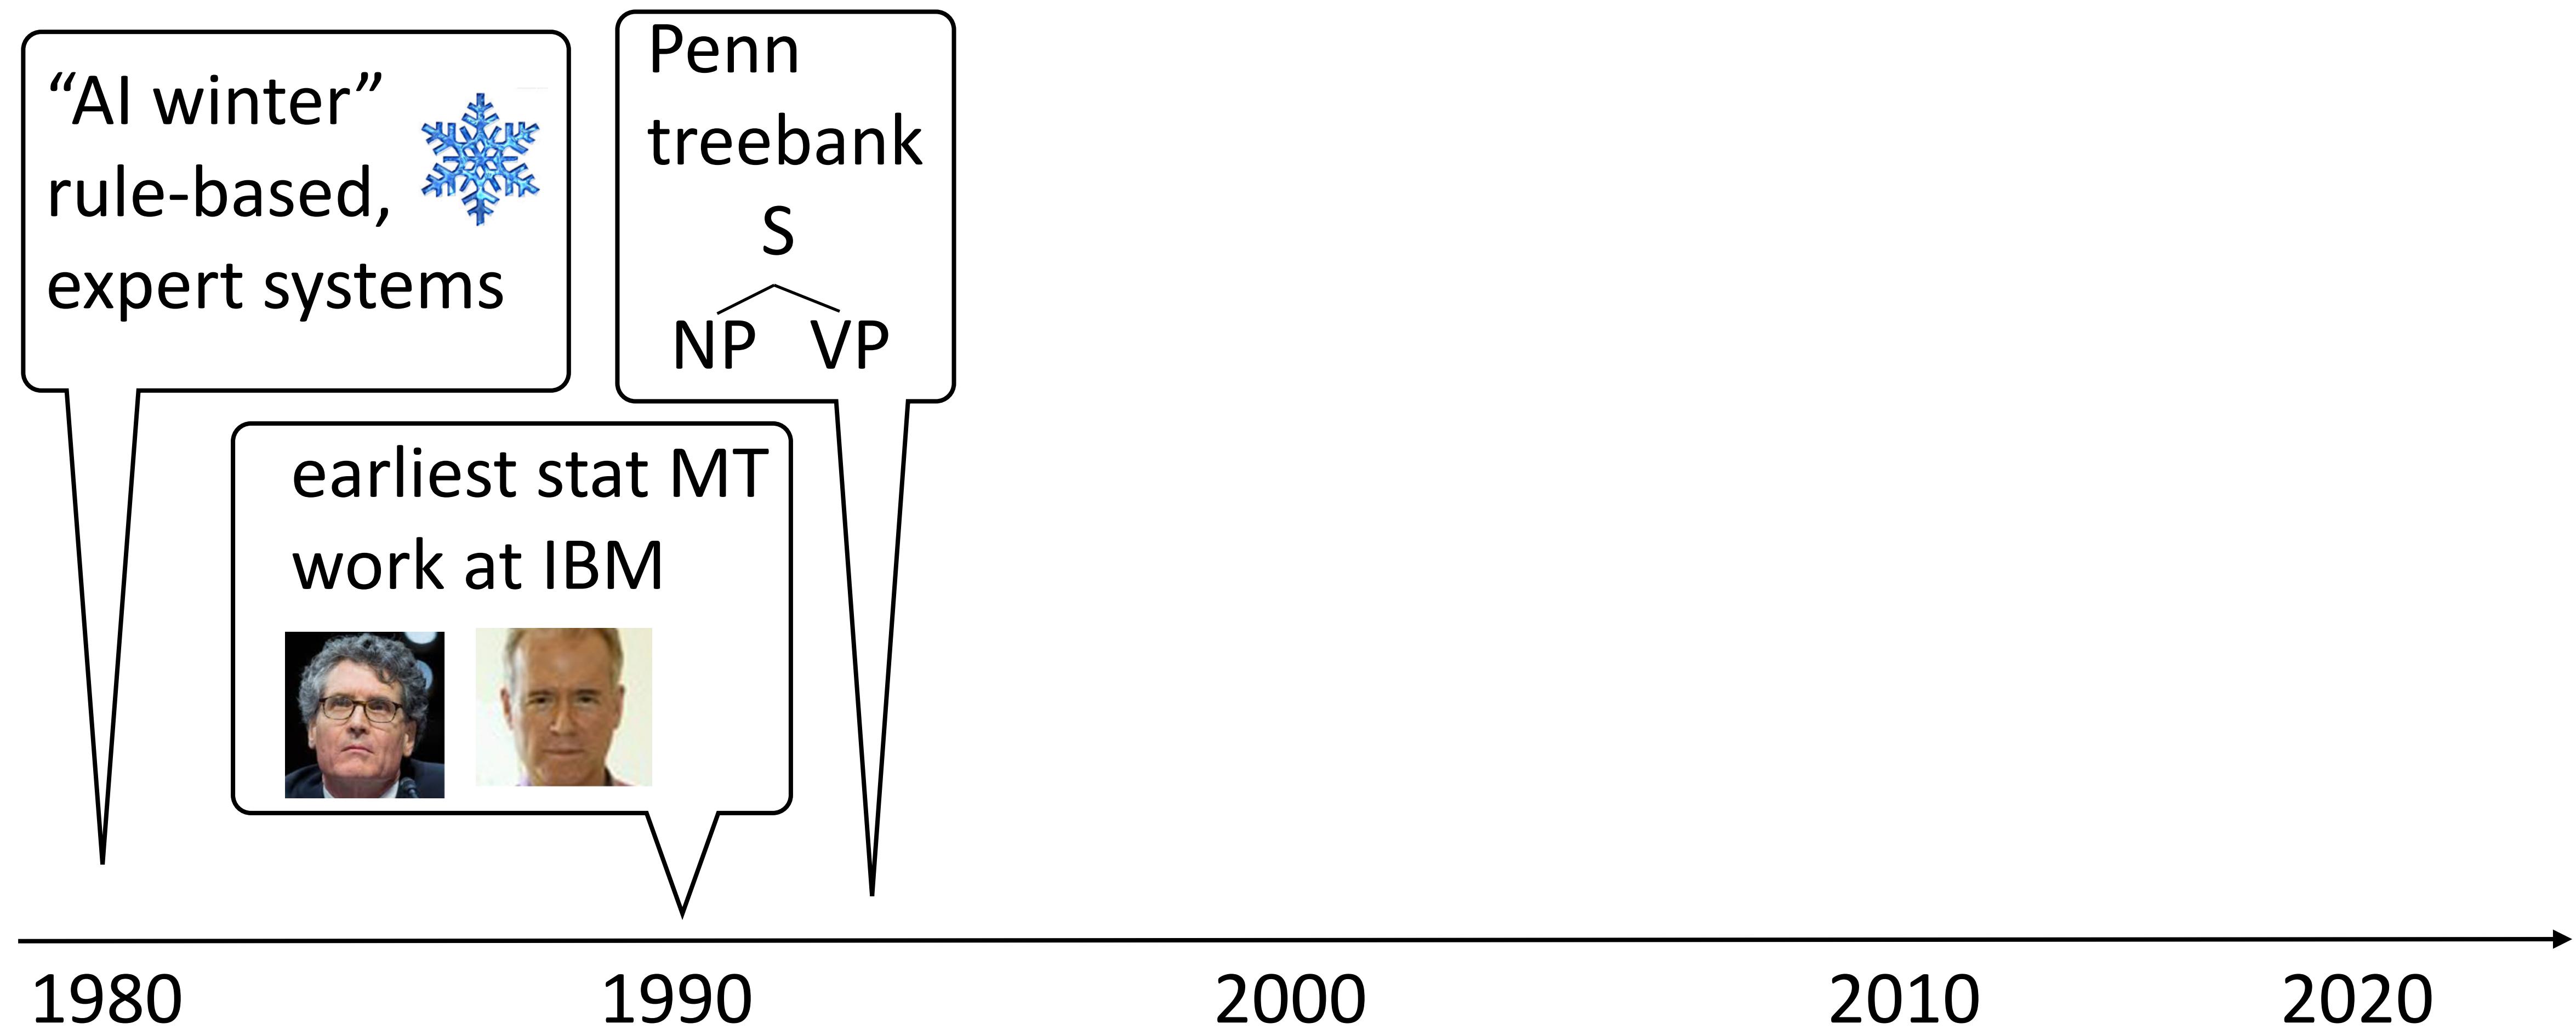
1990

2000

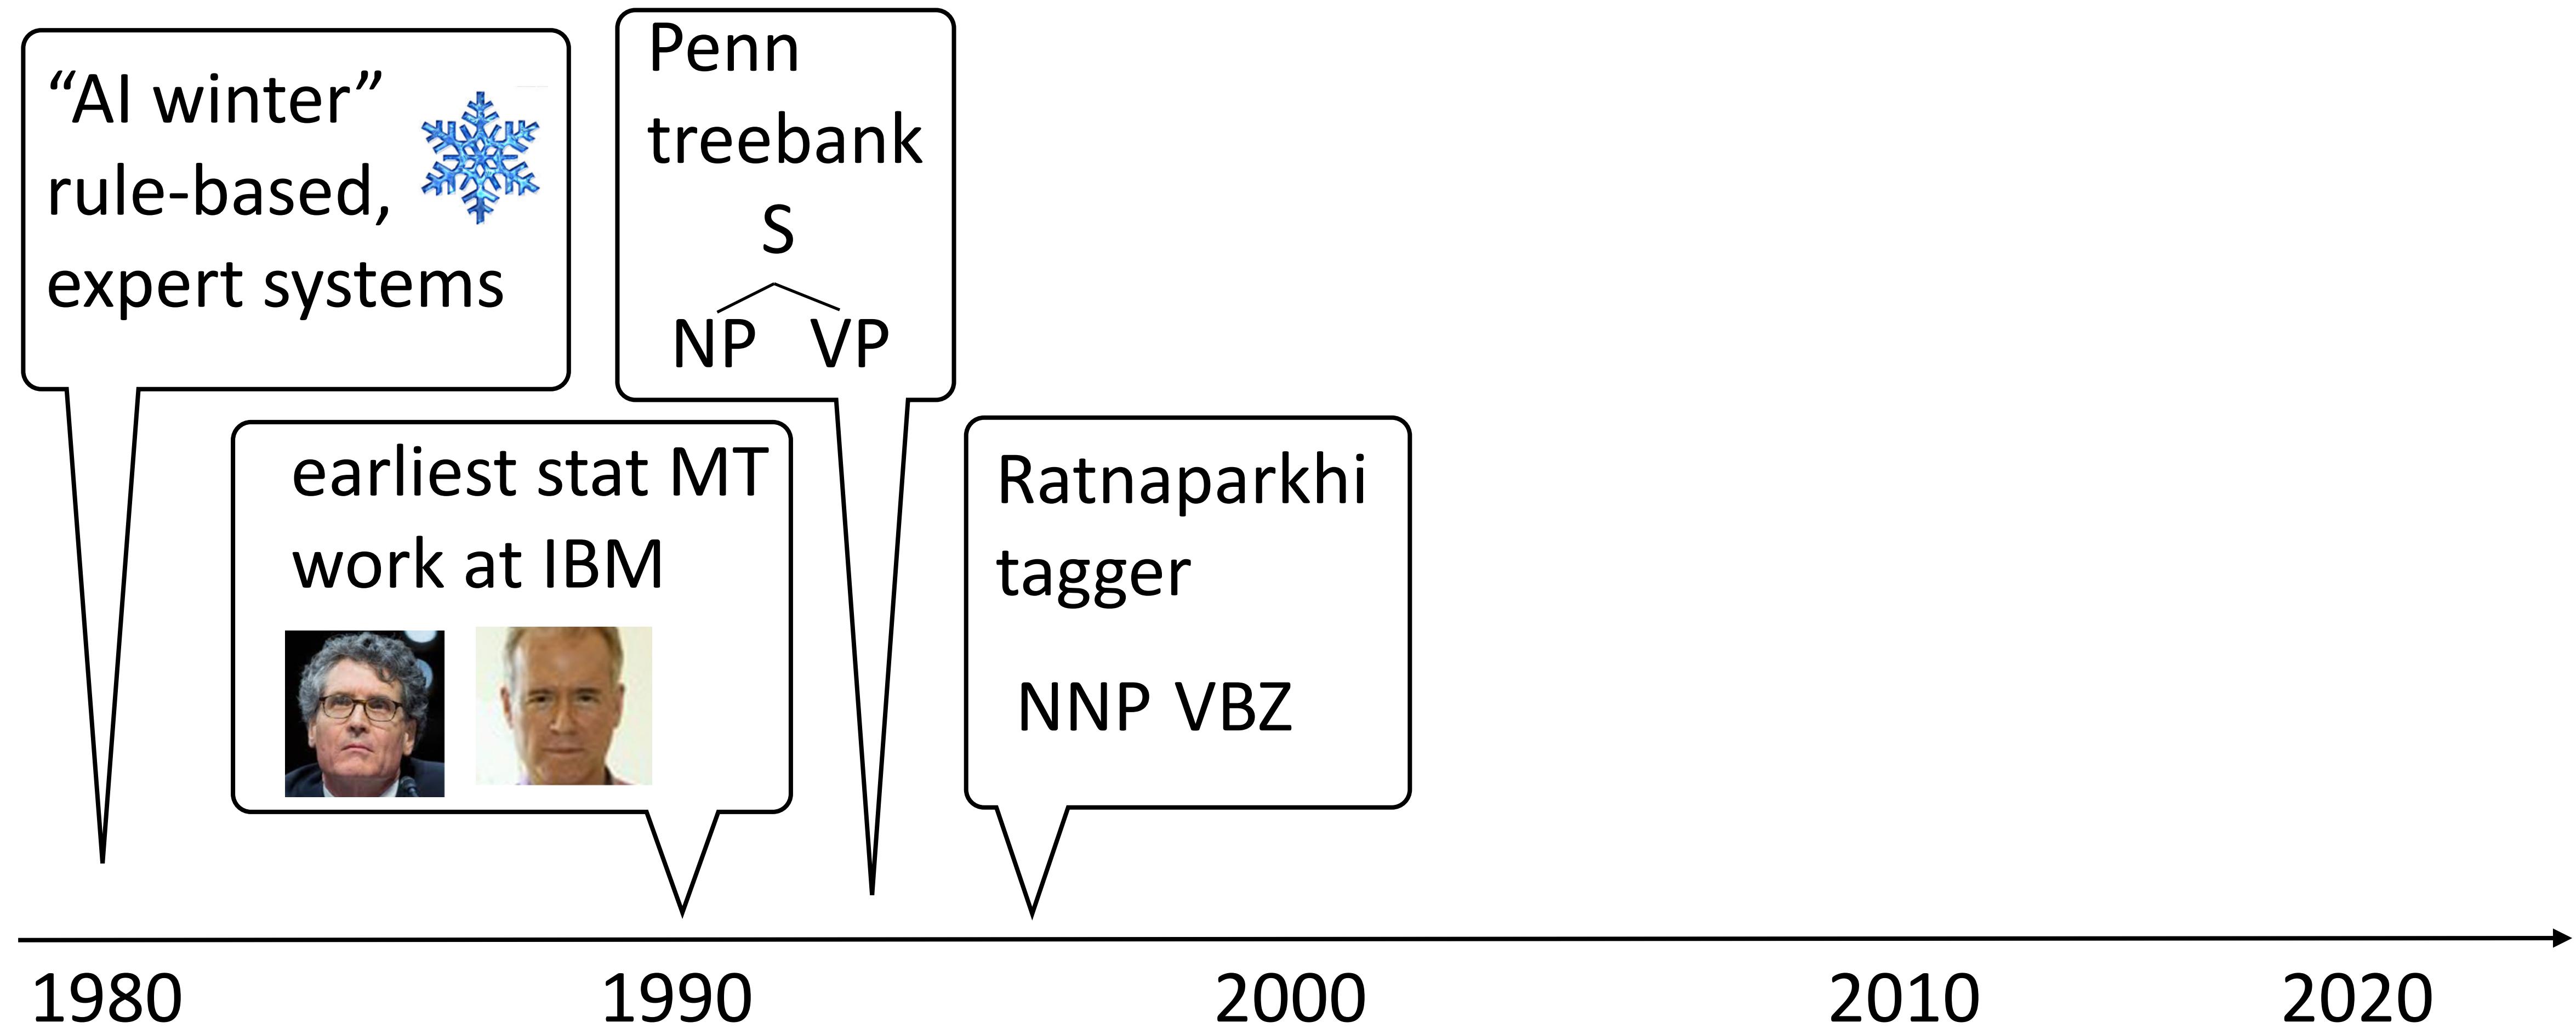
2010

2020

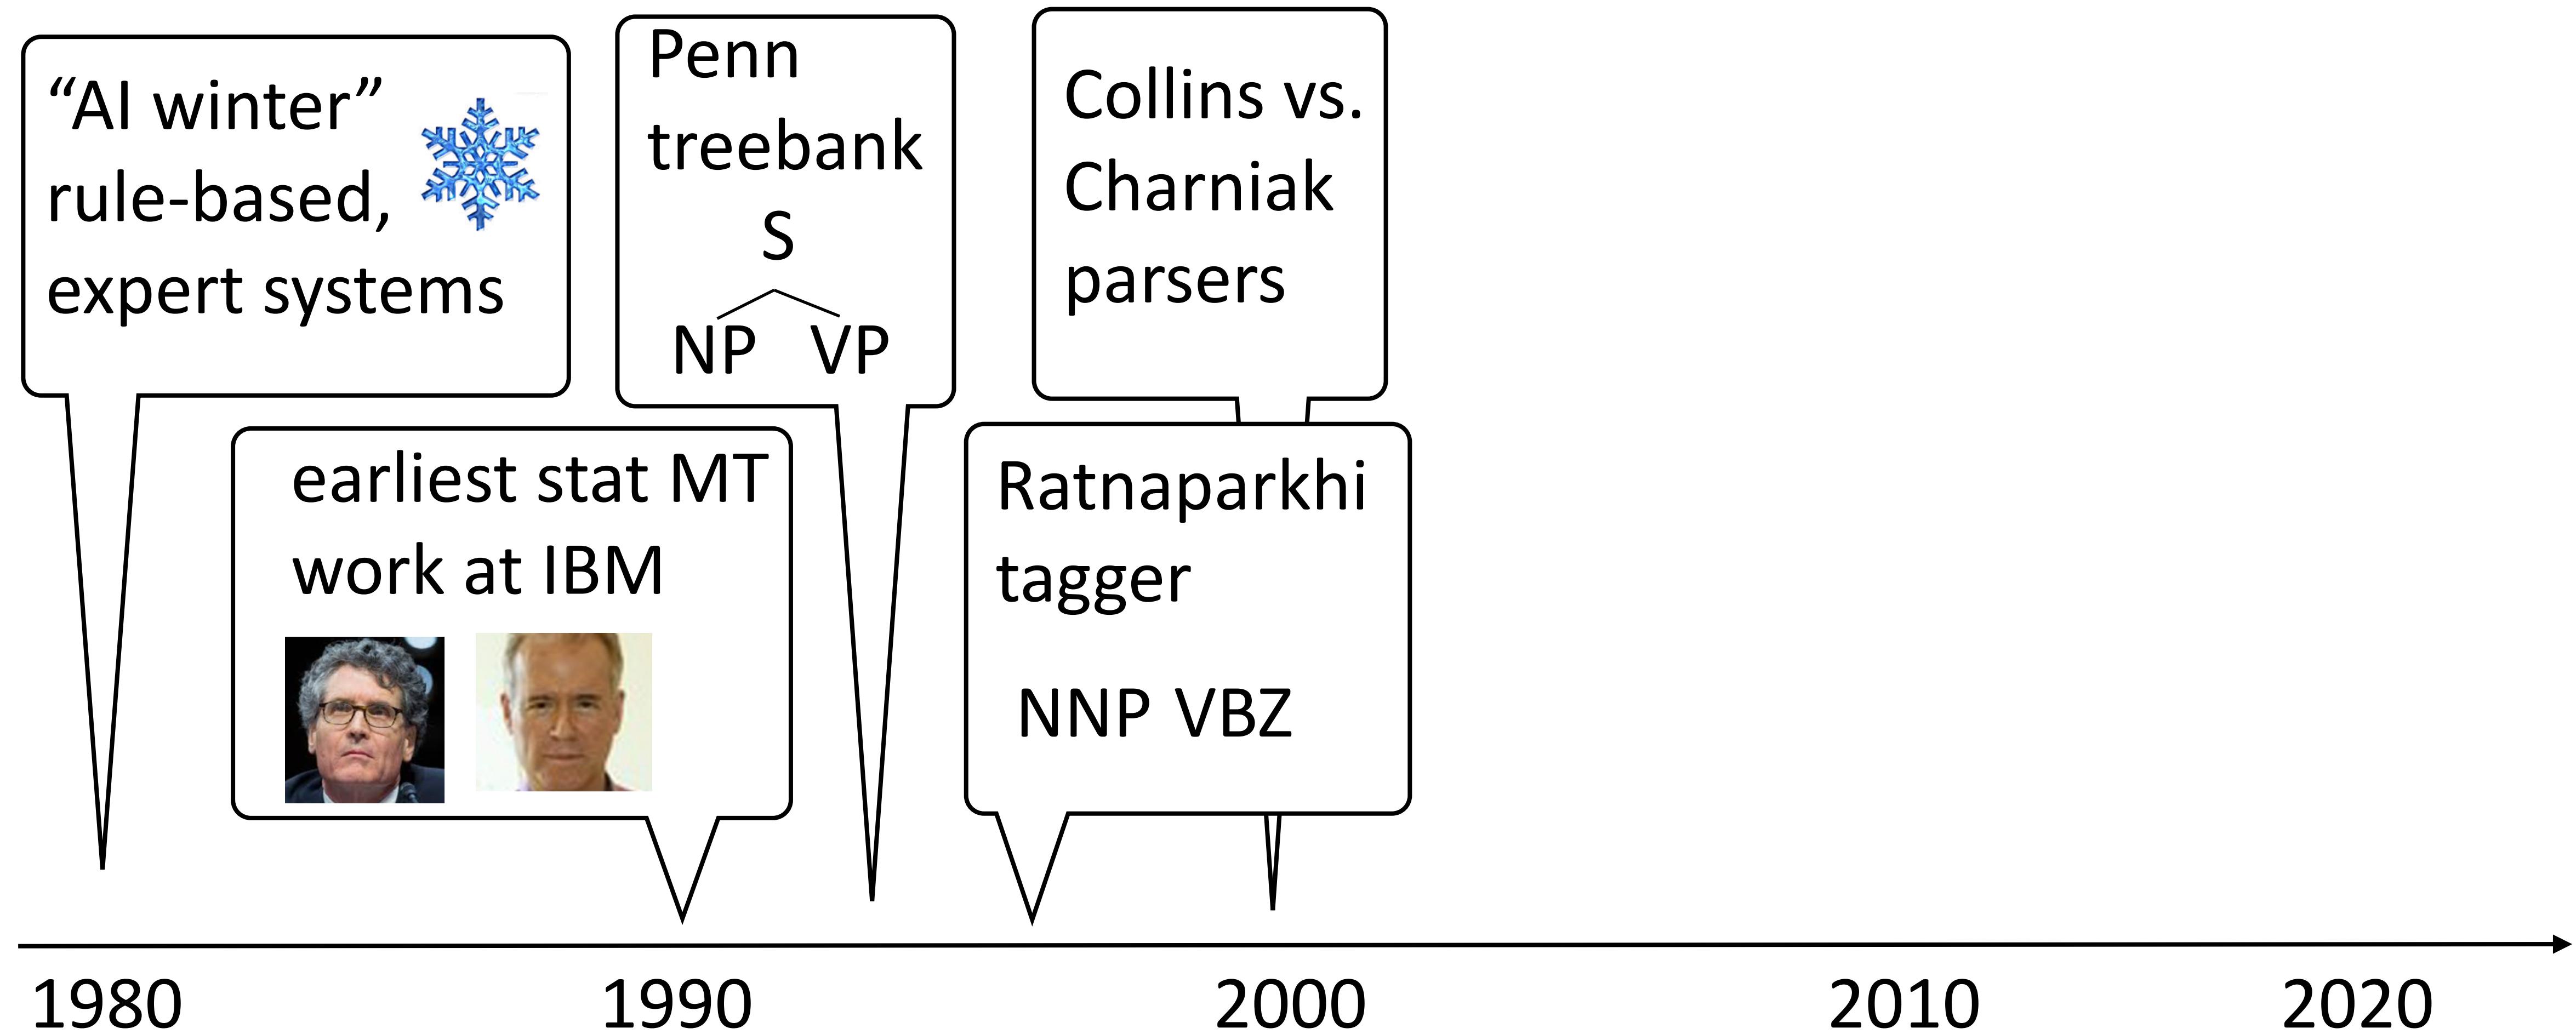
A brief history of (modern) NLP



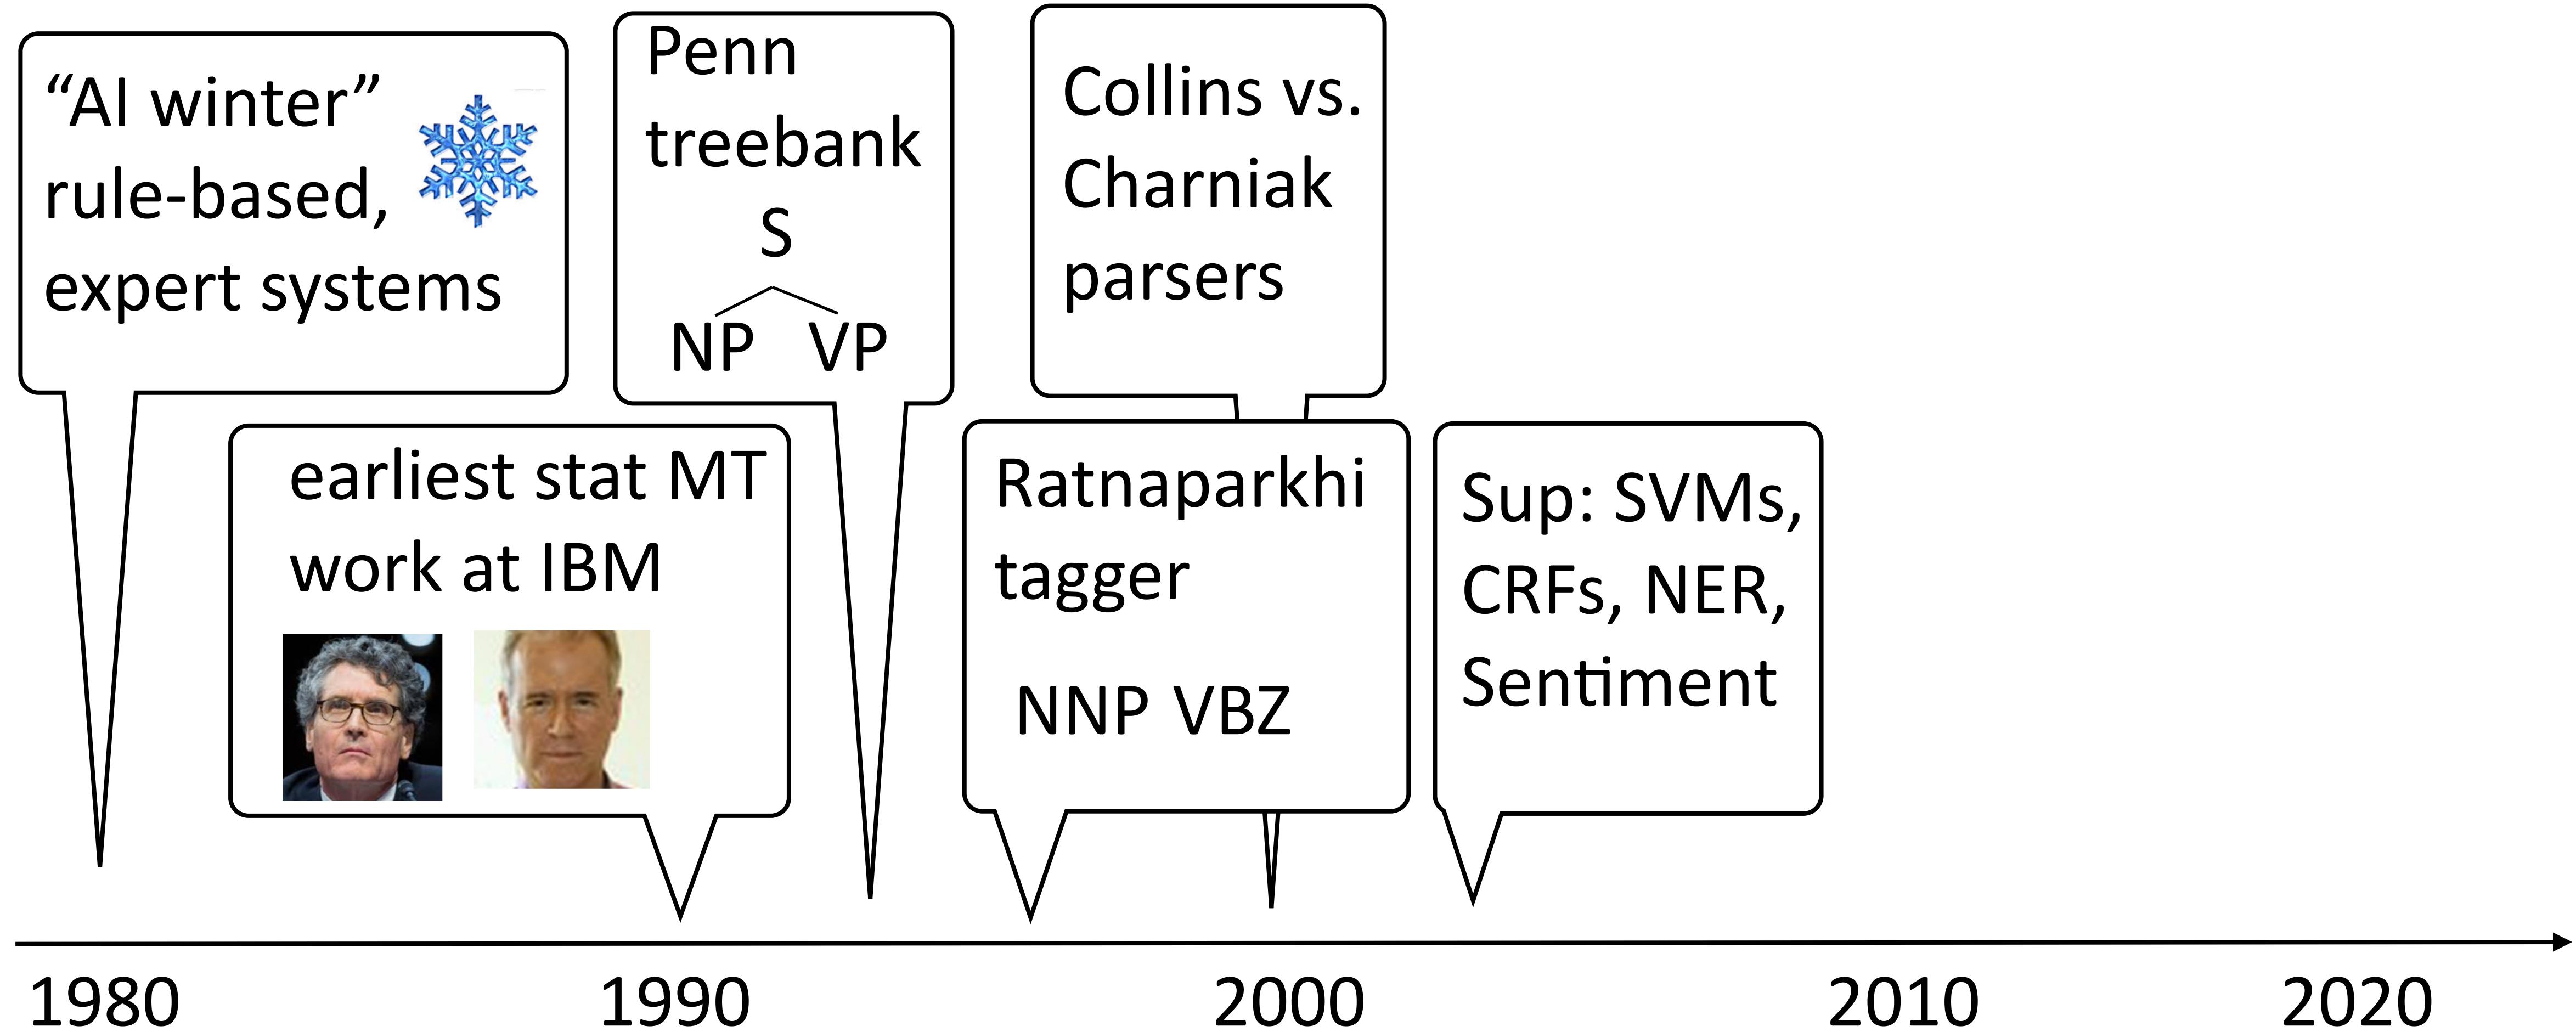
A brief history of (modern) NLP



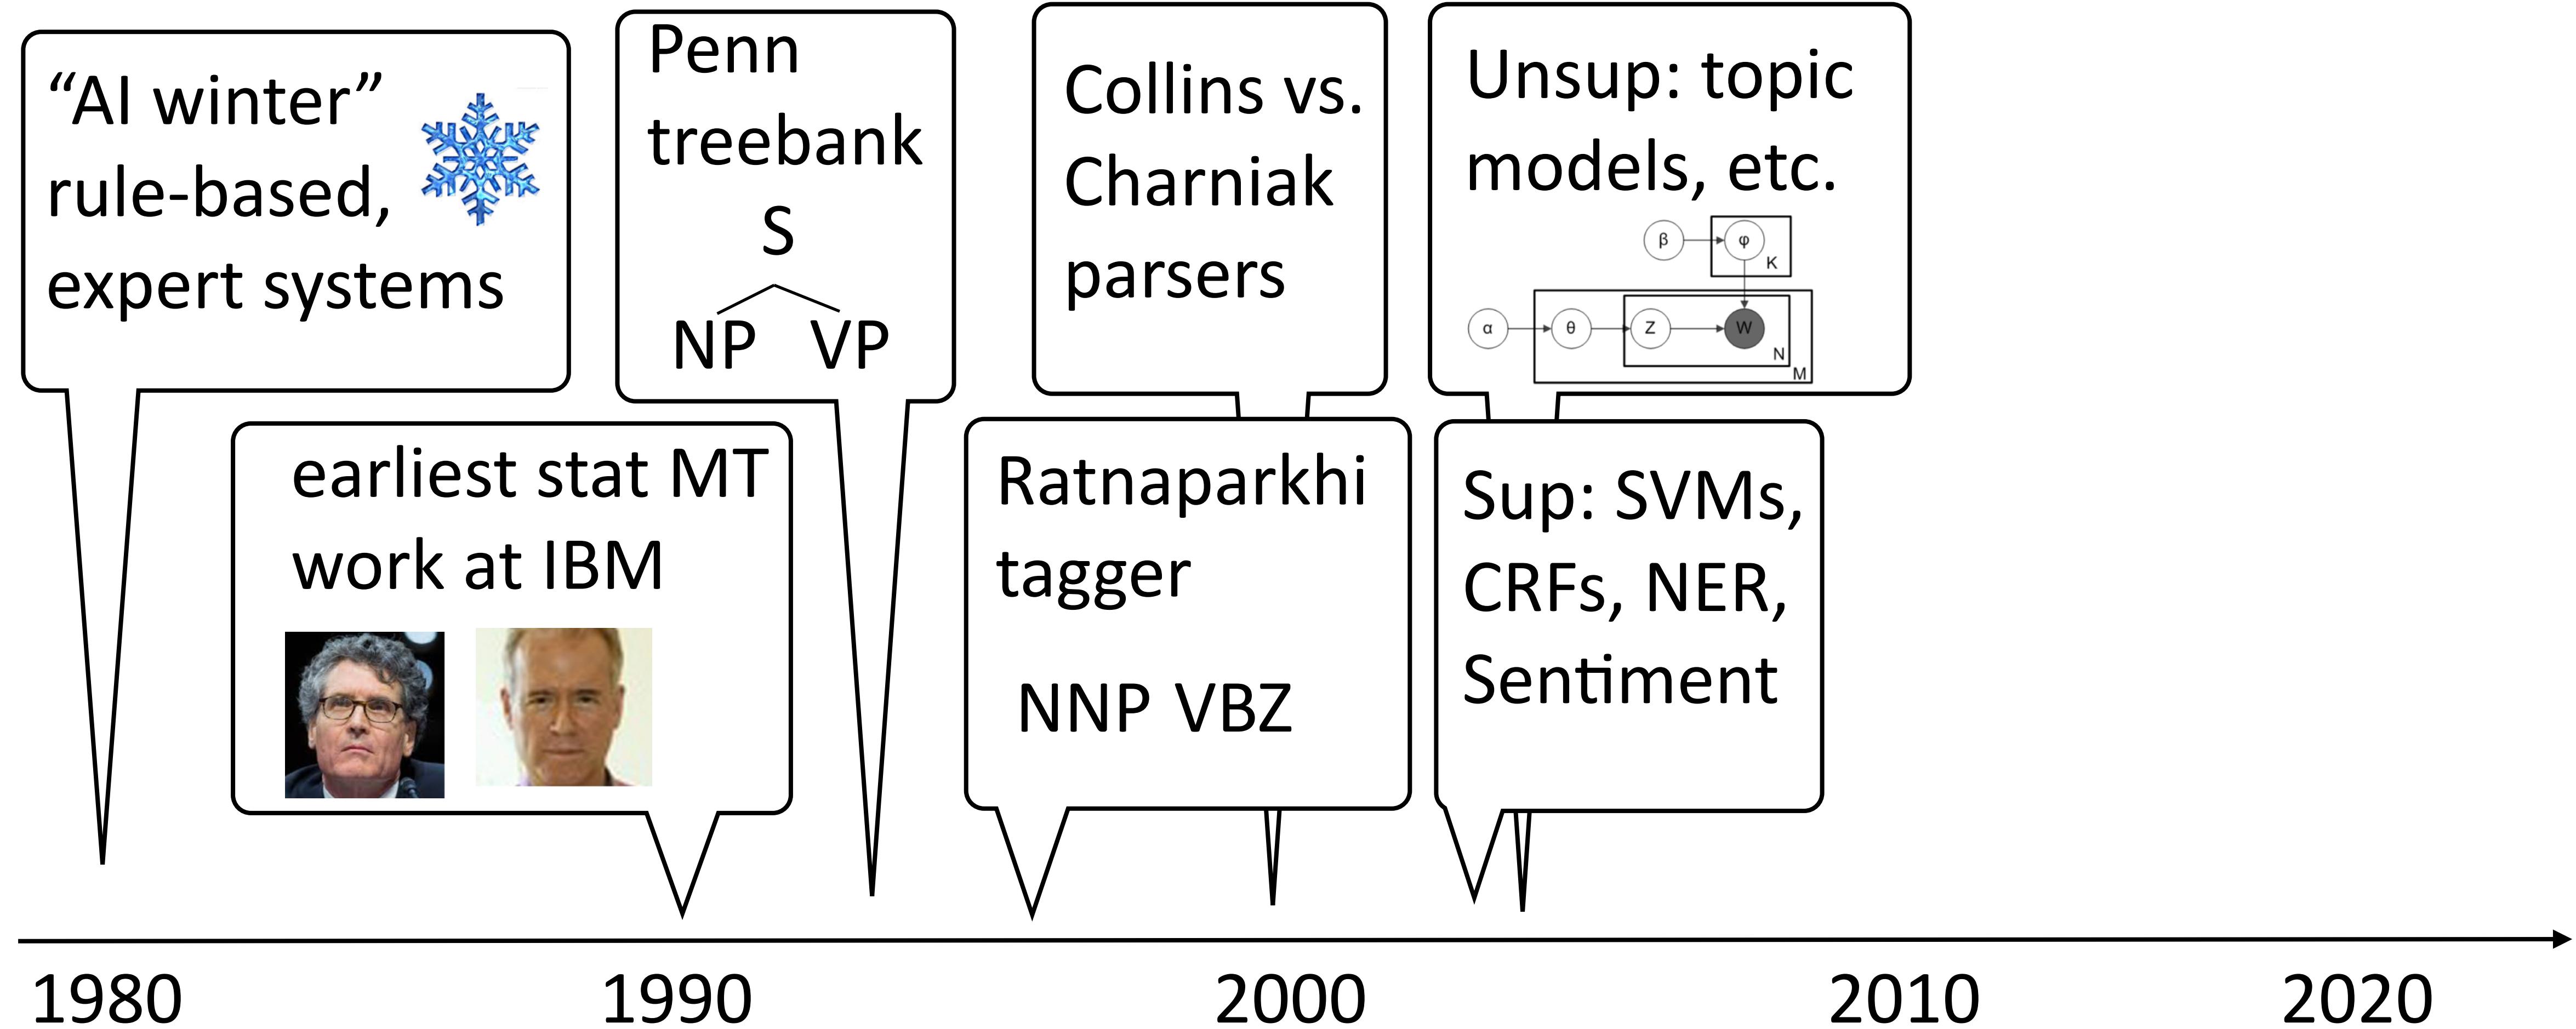
A brief history of (modern) NLP



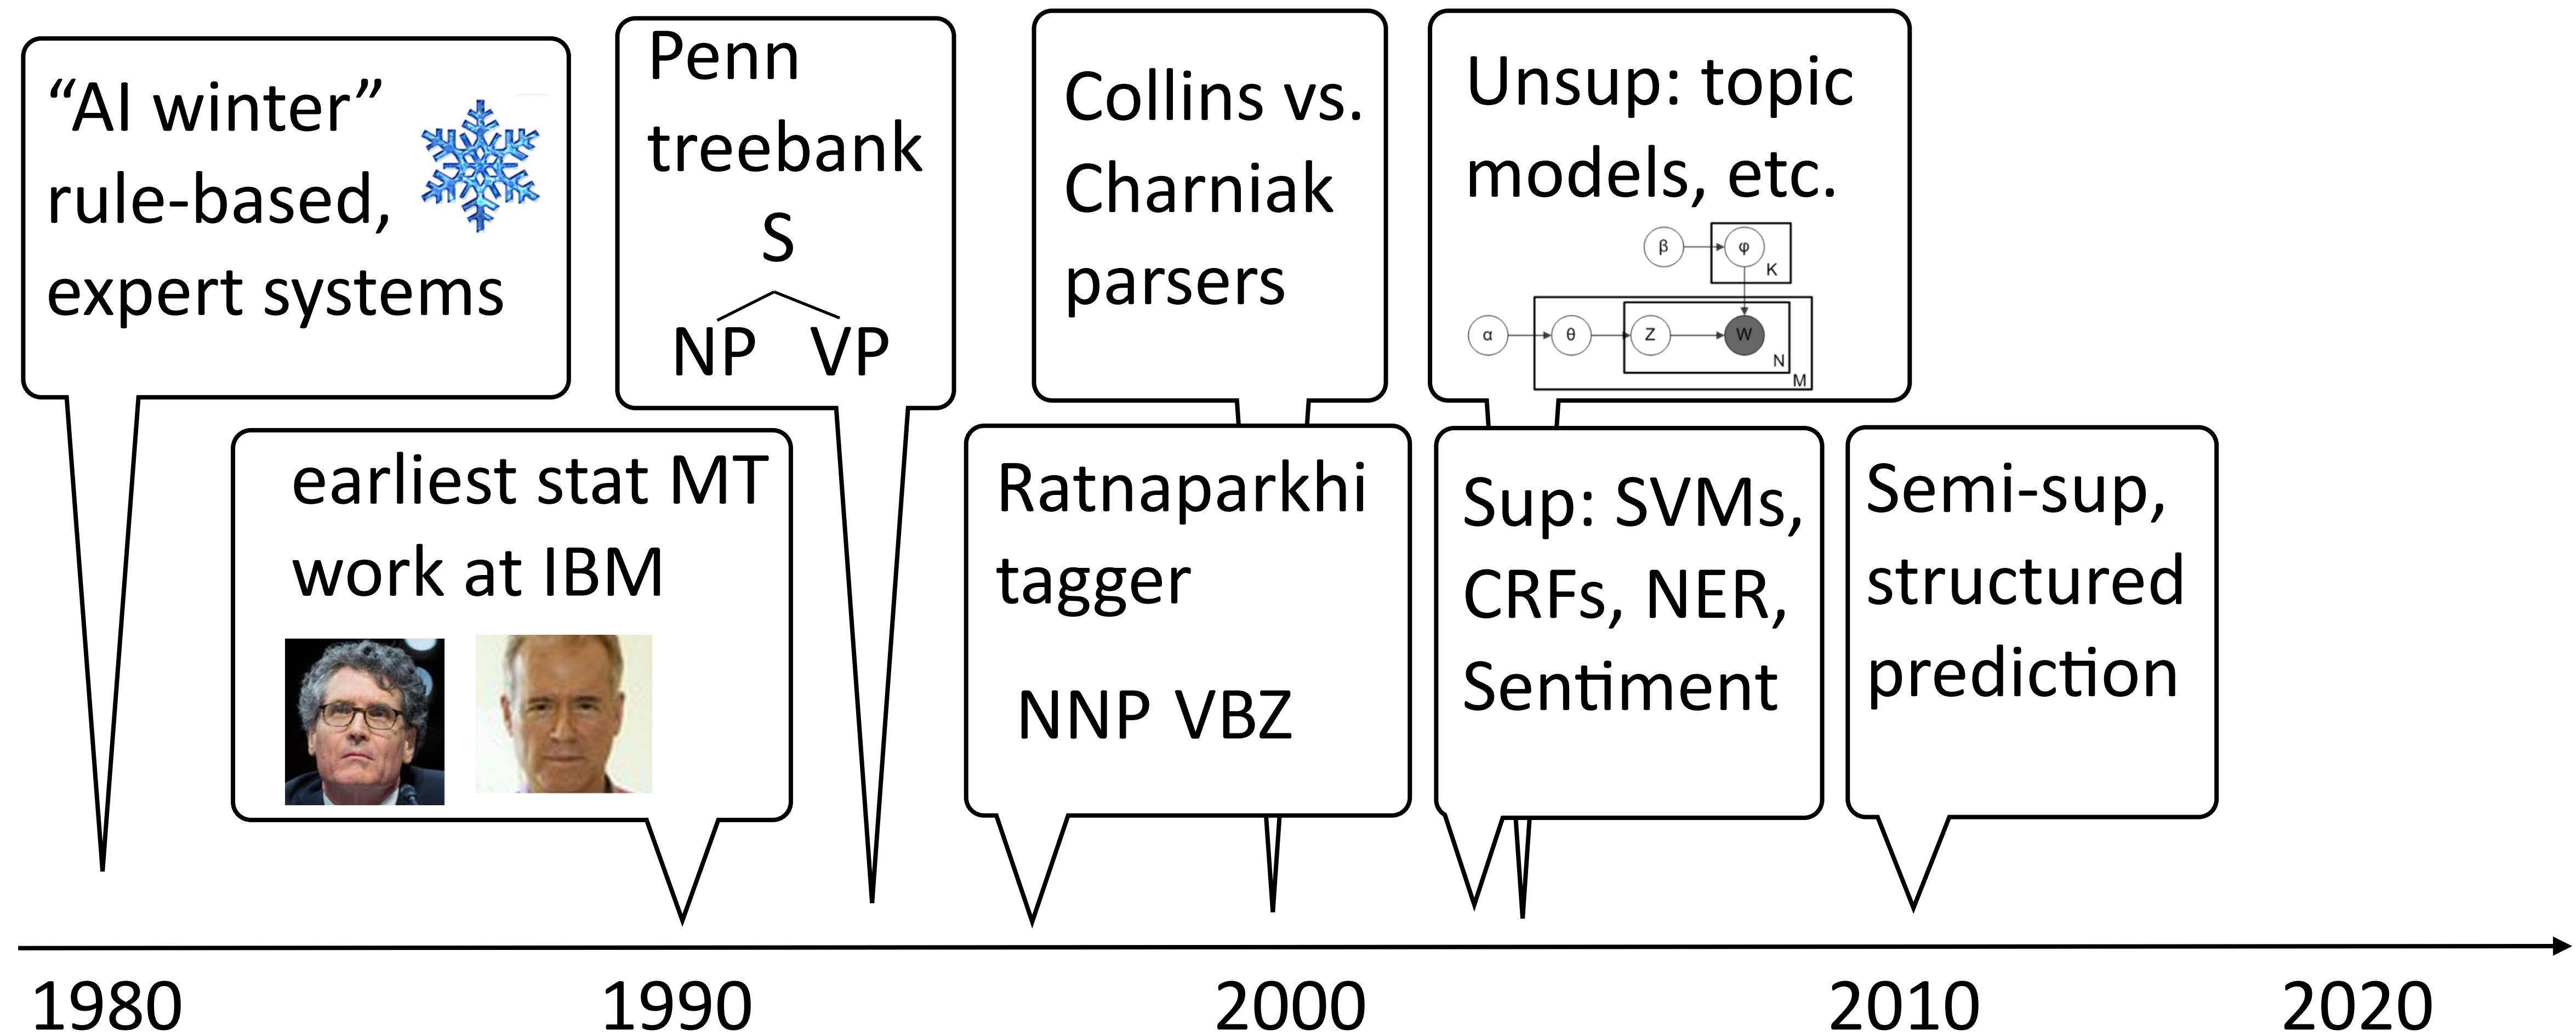
A brief history of (modern) NLP



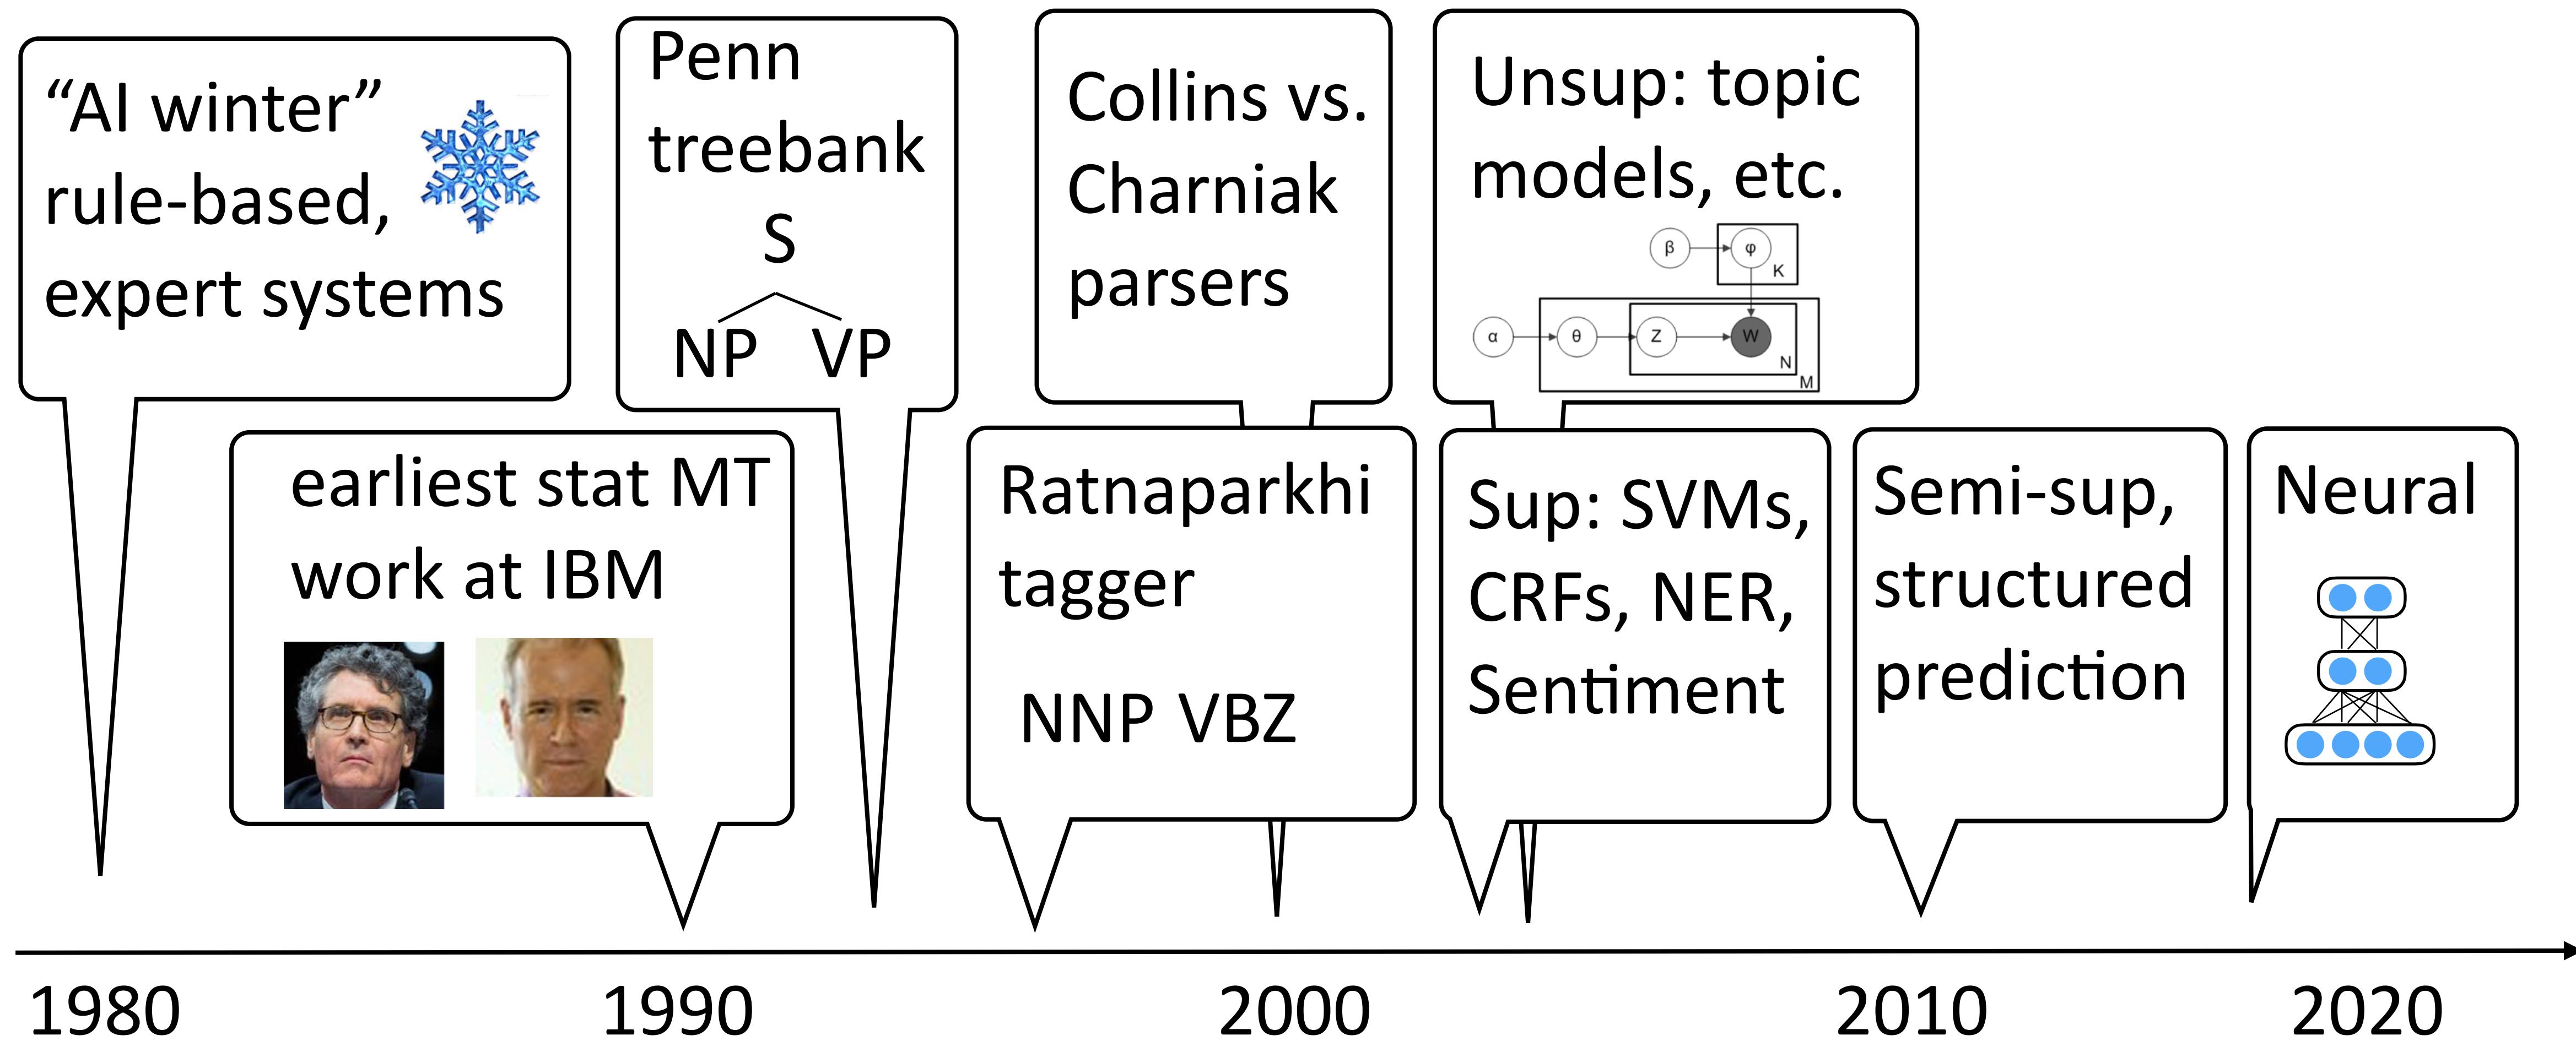
A brief history of (modern) NLP



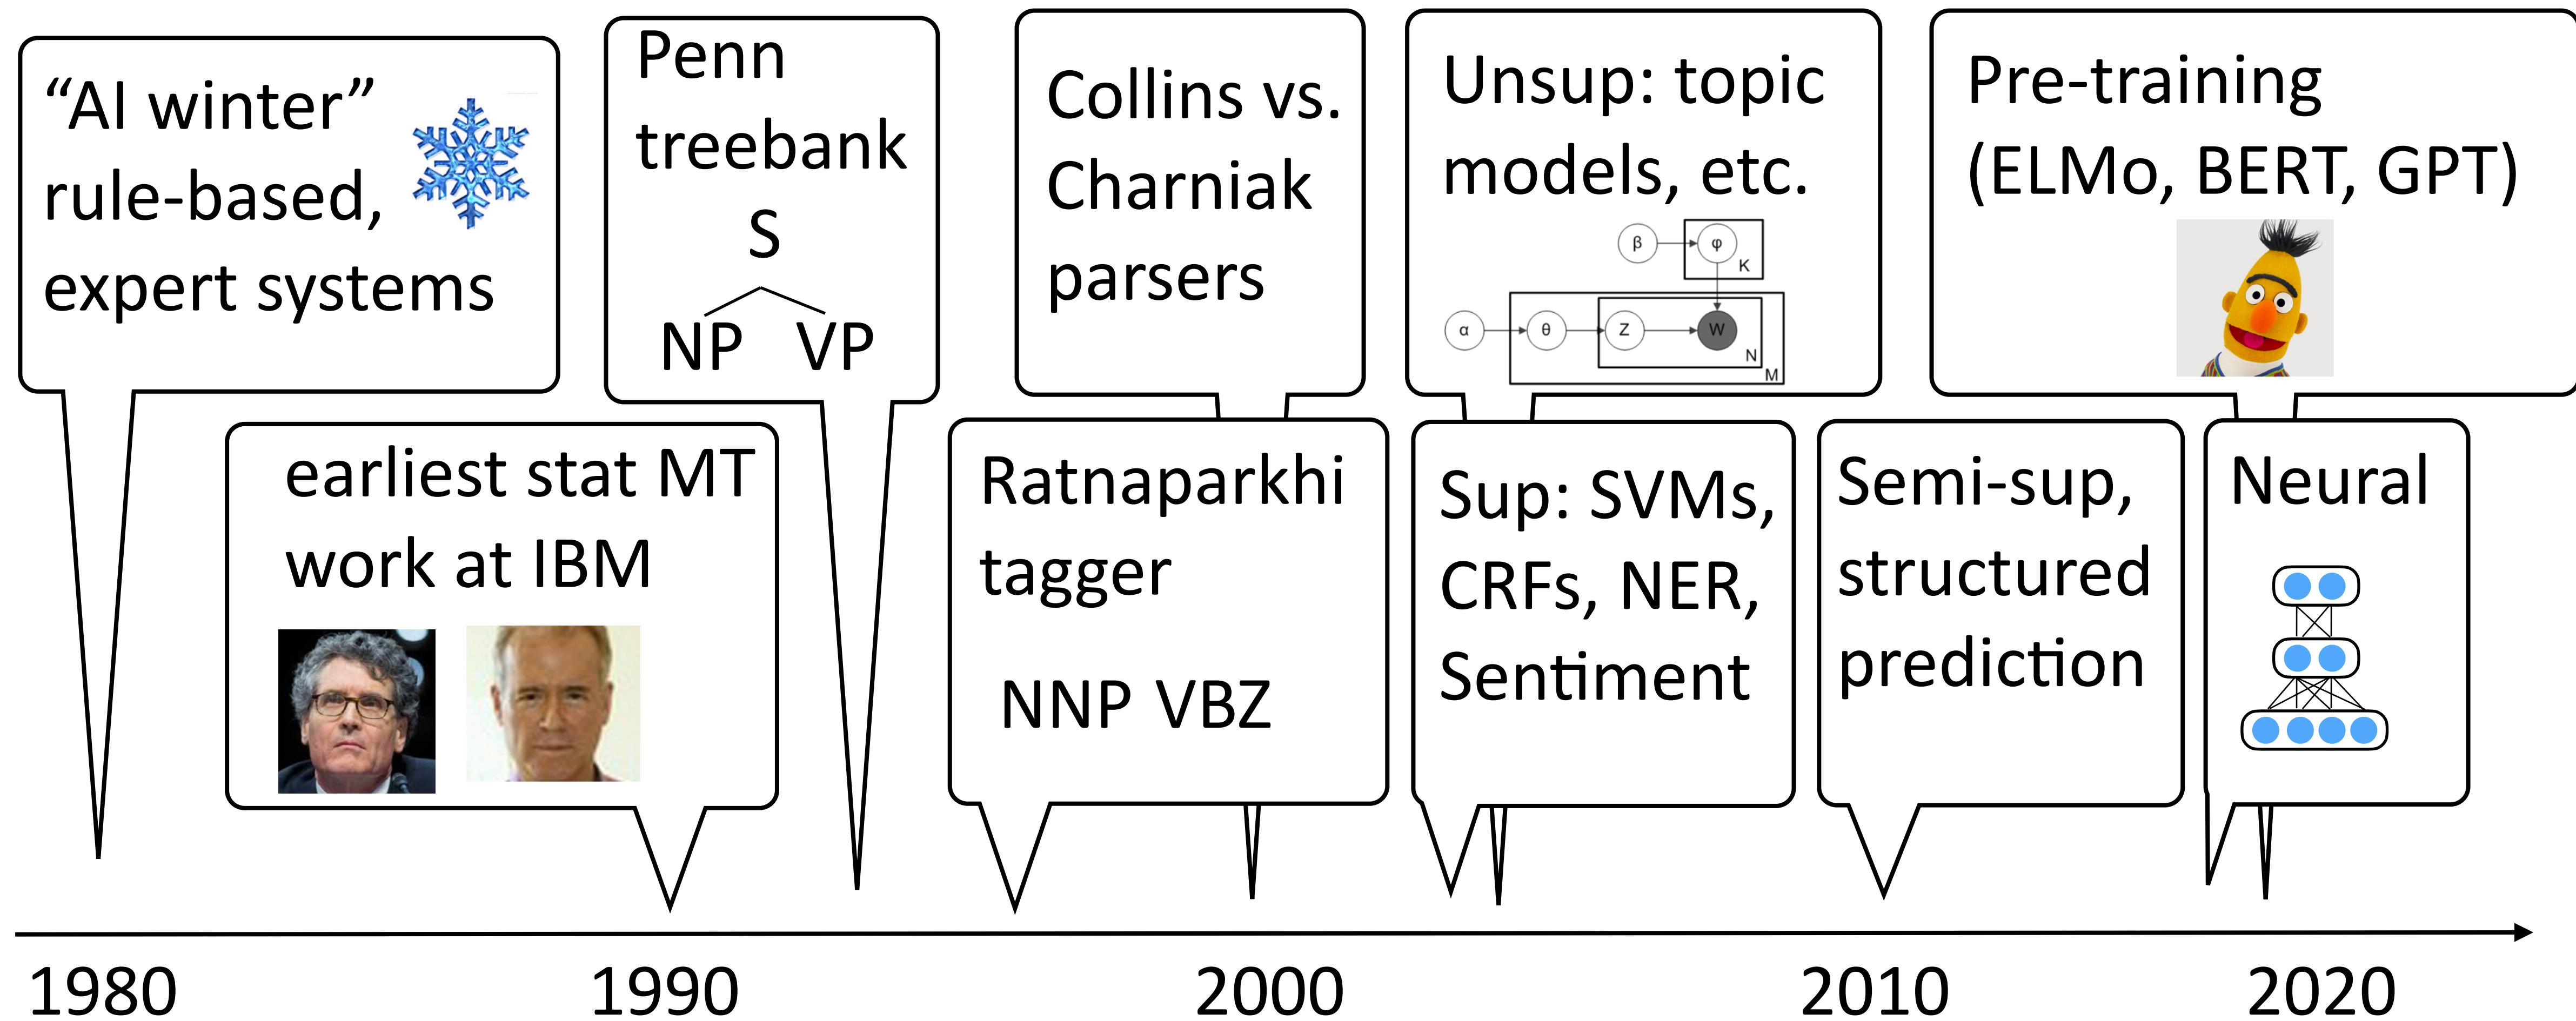
A brief history of (modern) NLP



A brief history of (modern) NLP



A brief history of (modern) NLP



How Much Training Data do we Need?

“Learning a Part-of-Speech Tagger from Two Hours of Annotation”
Garrette and Baldridge (2013)

How Much Training Data do we Need?

- ▶ All of these techniques are data-driven! Some data is naturally occurring, but may need to label

“Learning a Part-of-Speech Tagger from Two Hours of Annotation”
Garrette and Baldridge (2013)

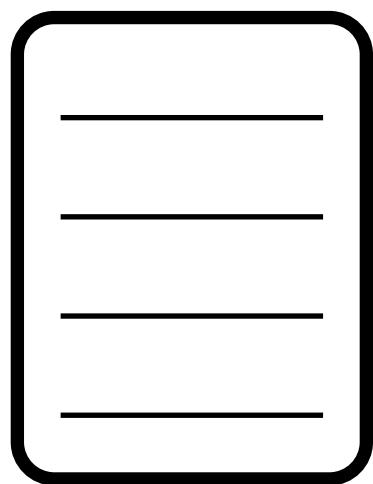
How Much Training Data do we Need?

- ▶ All of these techniques are data-driven! Some data is naturally occurring, but may need to label
- ▶ Supervised techniques work well on very little data

“Learning a Part-of-Speech Tagger from Two Hours of Annotation”
Garrette and Baldridge (2013)

How Much Training Data do we Need?

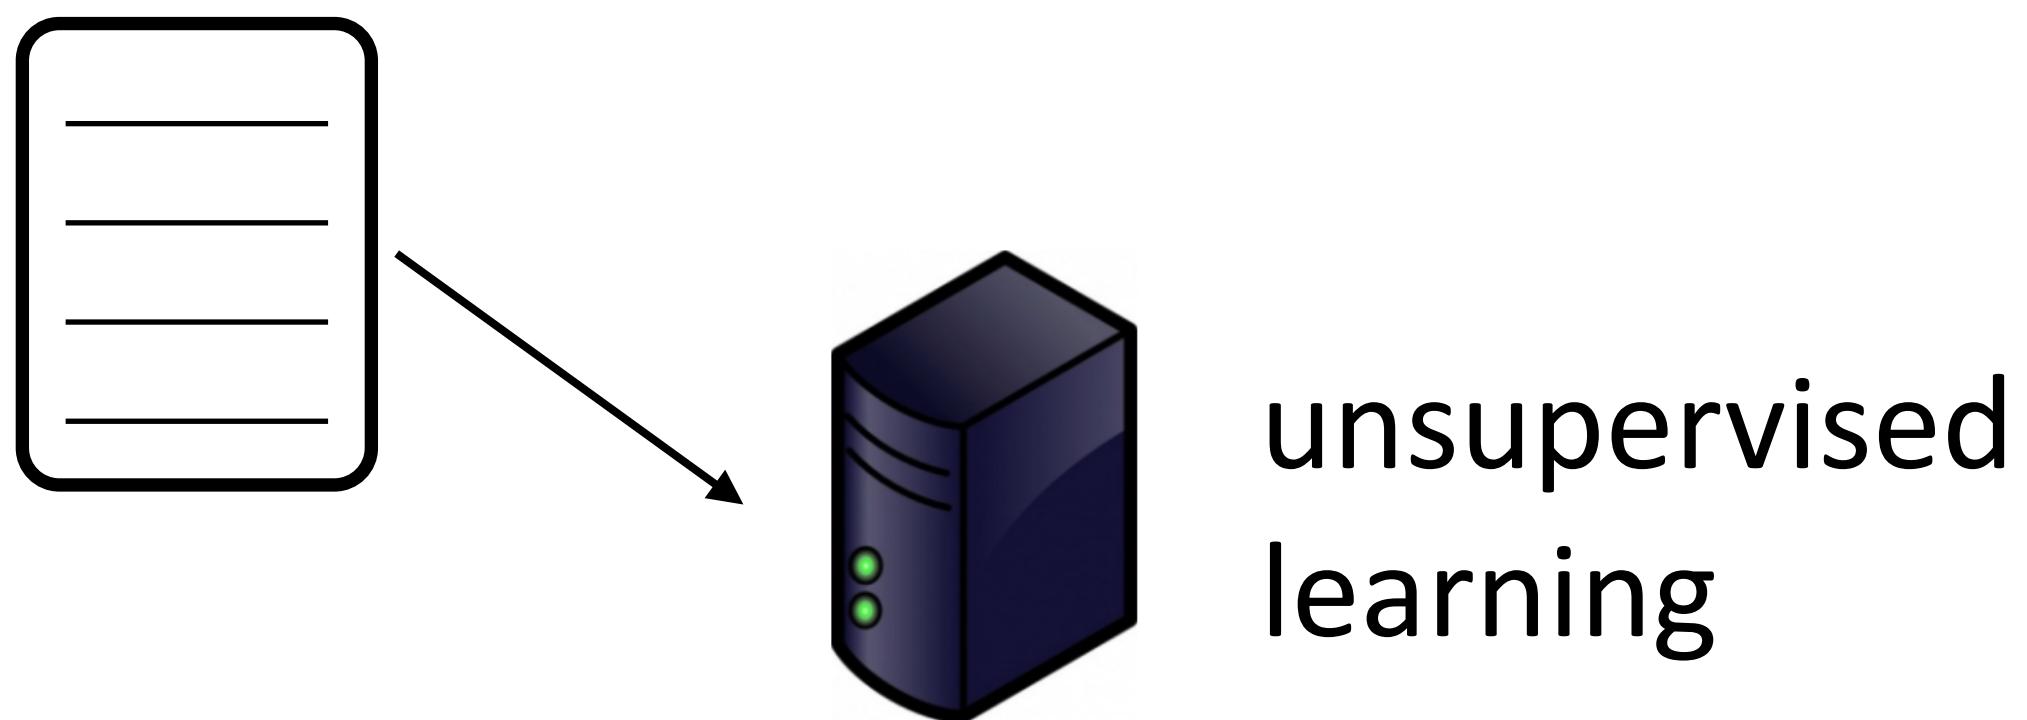
- ▶ All of these techniques are data-driven! Some data is naturally occurring, but may need to label
- ▶ Supervised techniques work well on very little data



“Learning a Part-of-Speech Tagger from Two Hours of Annotation”
Garrette and Baldridge (2013)

How Much Training Data do we Need?

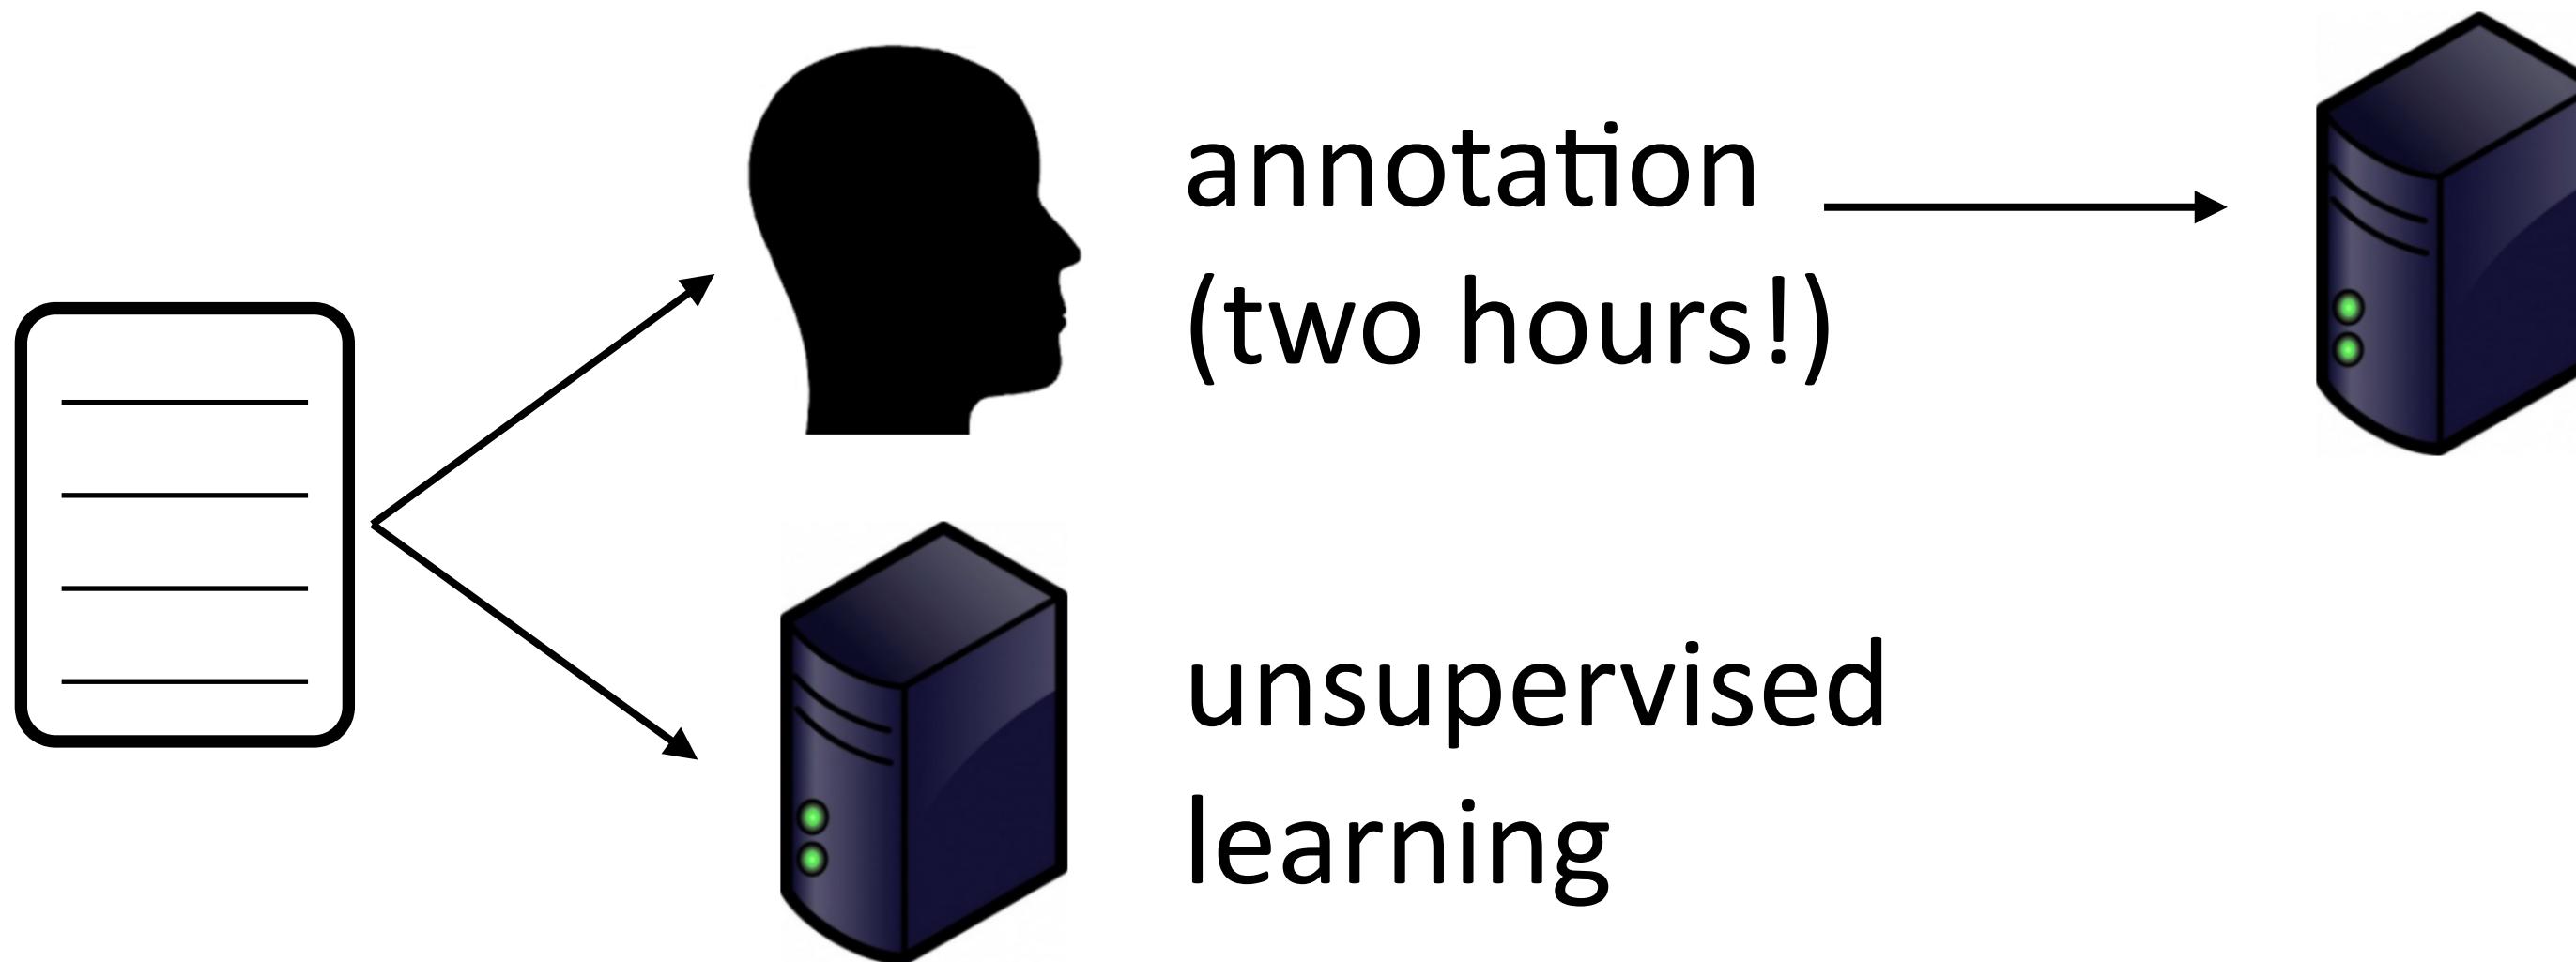
- ▶ All of these techniques are data-driven! Some data is naturally occurring, but may need to label
- ▶ Supervised techniques work well on very little data



“Learning a Part-of-Speech Tagger from Two Hours of Annotation”
Garrette and Baldridge (2013)

How Much Training Data do we Need?

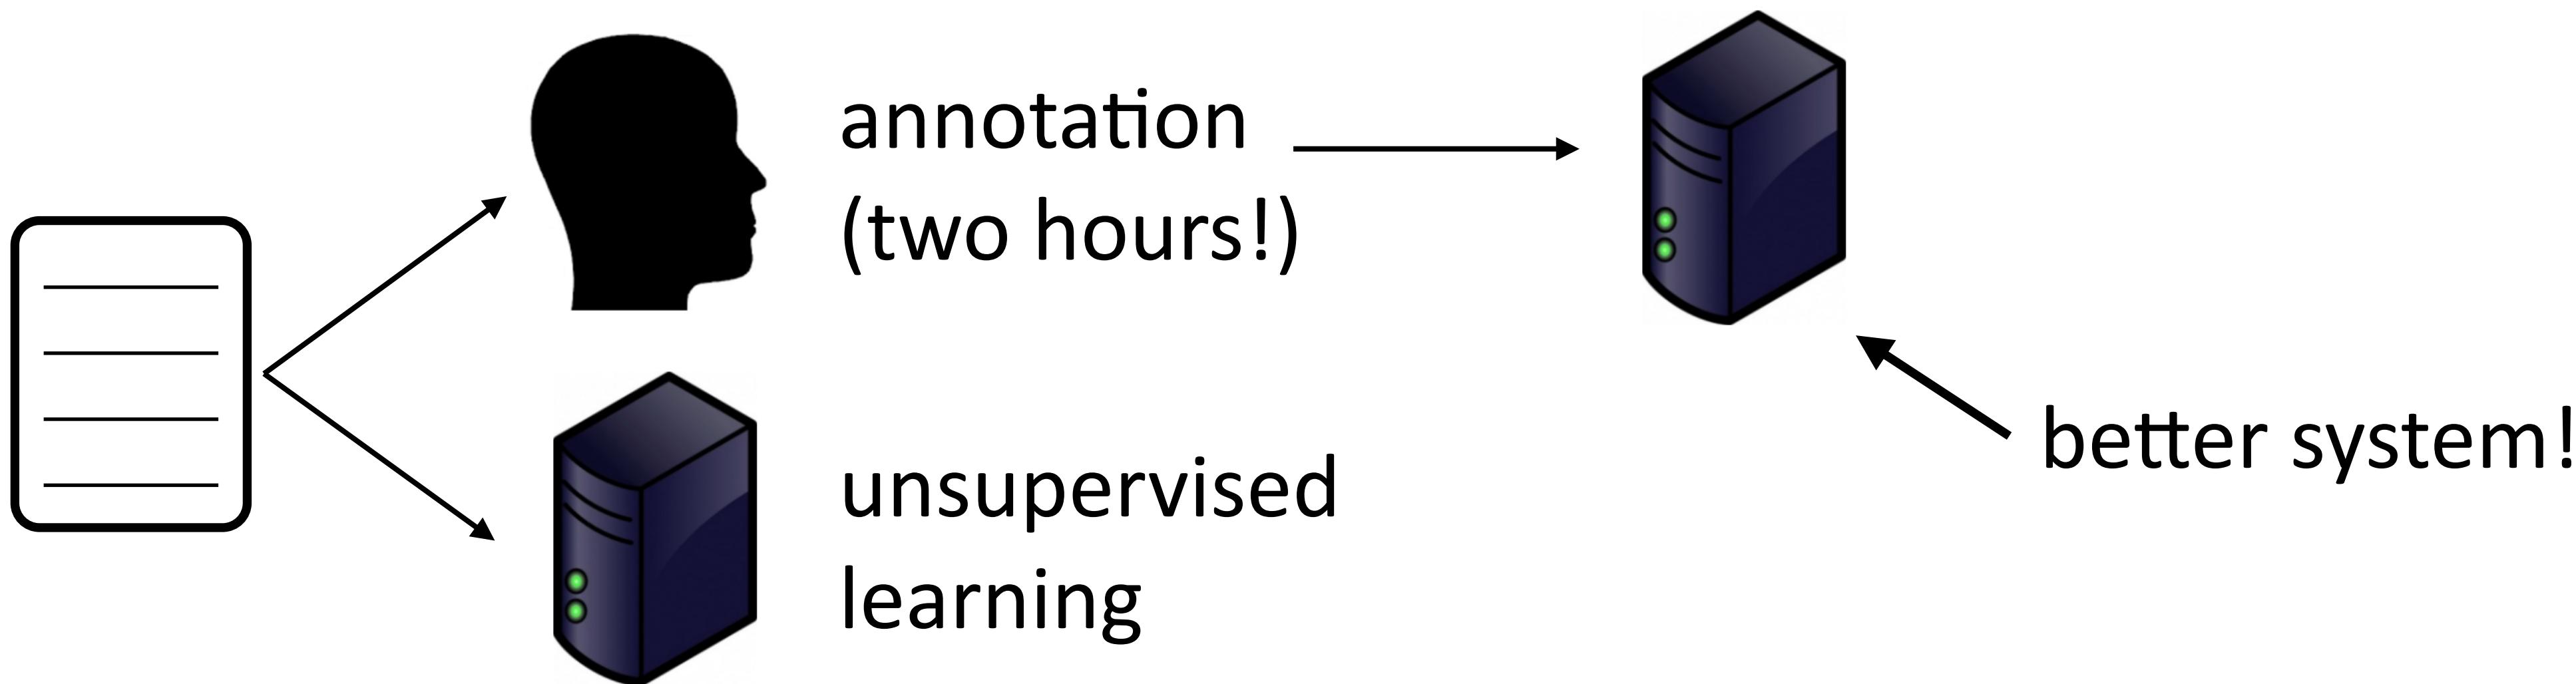
- ▶ All of these techniques are data-driven! Some data is naturally occurring, but may need to label
- ▶ Supervised techniques work well on very little data



“Learning a Part-of-Speech Tagger from Two Hours of Annotation”
Garrette and Baldridge (2013)

How Much Training Data do we Need?

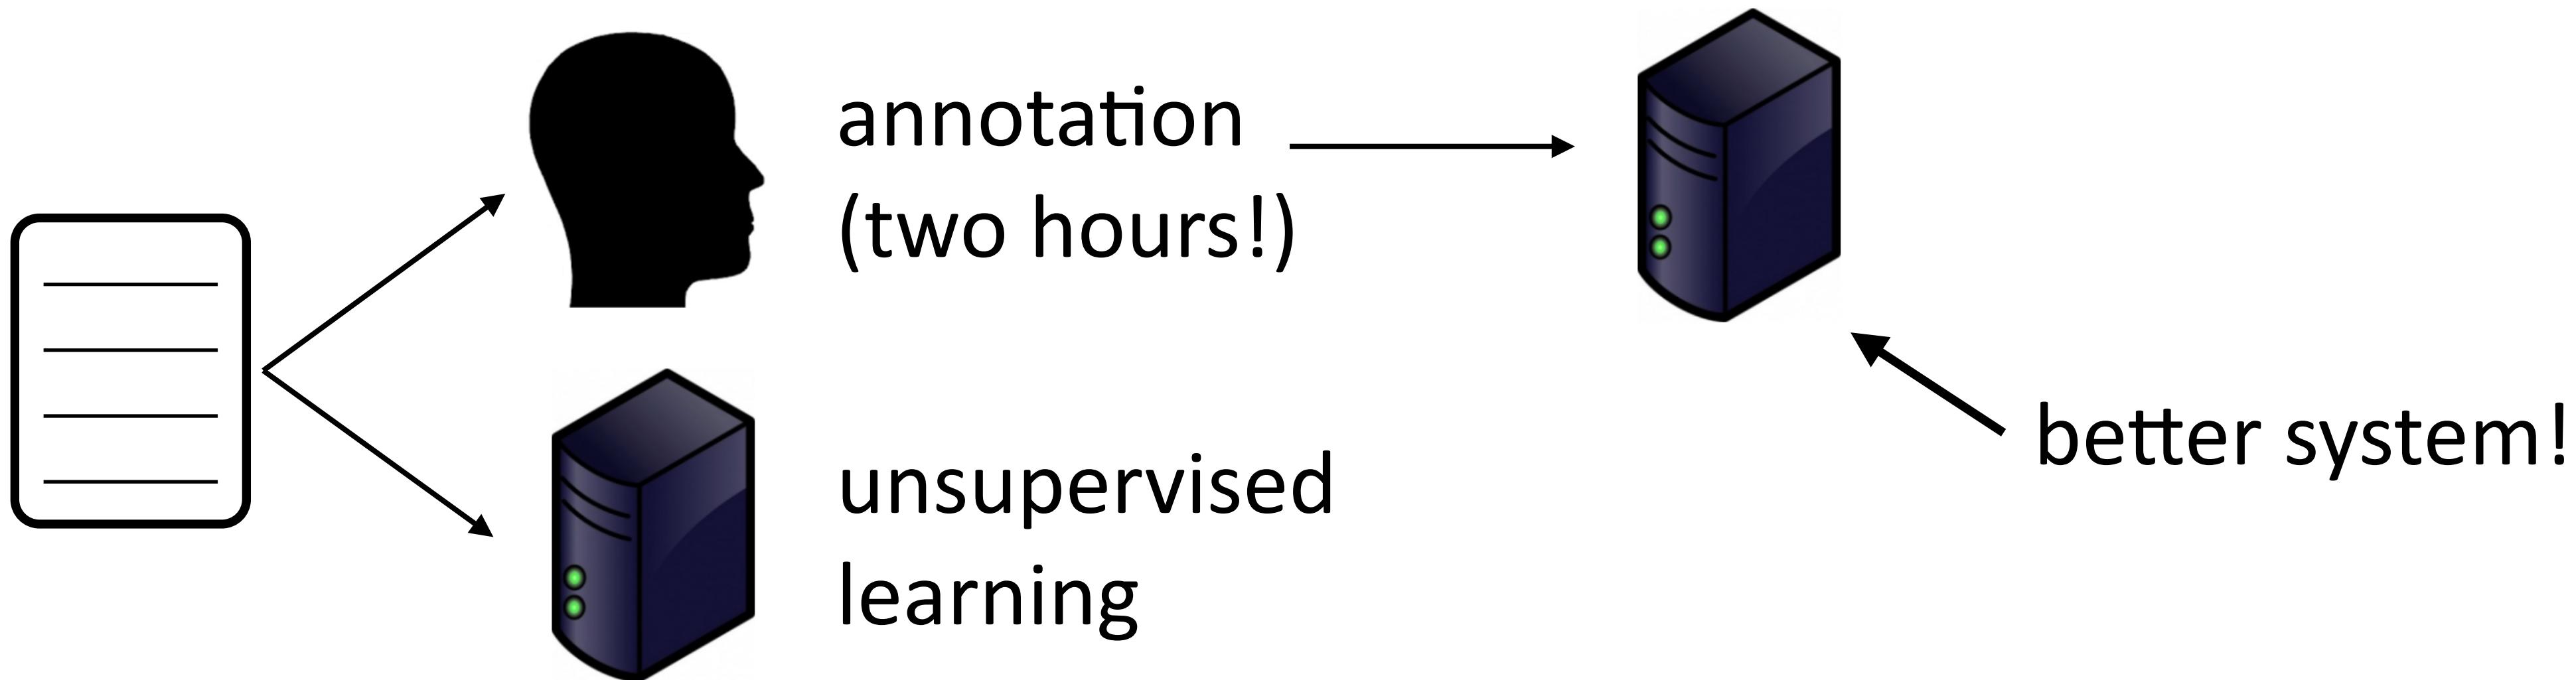
- ▶ All of these techniques are data-driven! Some data is naturally occurring, but may need to label
- ▶ Supervised techniques work well on very little data



“Learning a Part-of-Speech Tagger from Two Hours of Annotation”
Garrette and Baldridge (2013)

How Much Training Data do we Need?

- ▶ All of these techniques are data-driven! Some data is naturally occurring, but may need to label
- ▶ Supervised techniques work well on very little data



- ▶ Even neural nets can do pretty well!

“Learning a Part-of-Speech Tagger from Two Hours of Annotation”
Garrette and Baldridge (2013)

Pretraining

- ▶ Language modeling: predict the next word in a text $P(w_i | w_1, \dots, w_{i-1})$

$P(w | \text{I want to go to}) = 0.01 \text{ Hawai'i}$

0.005 LA

0.0001 class



: use this model for other purposes

$P(w | \text{the acting was horrible, I think the movie was}) = 0.1 \text{ bad}$

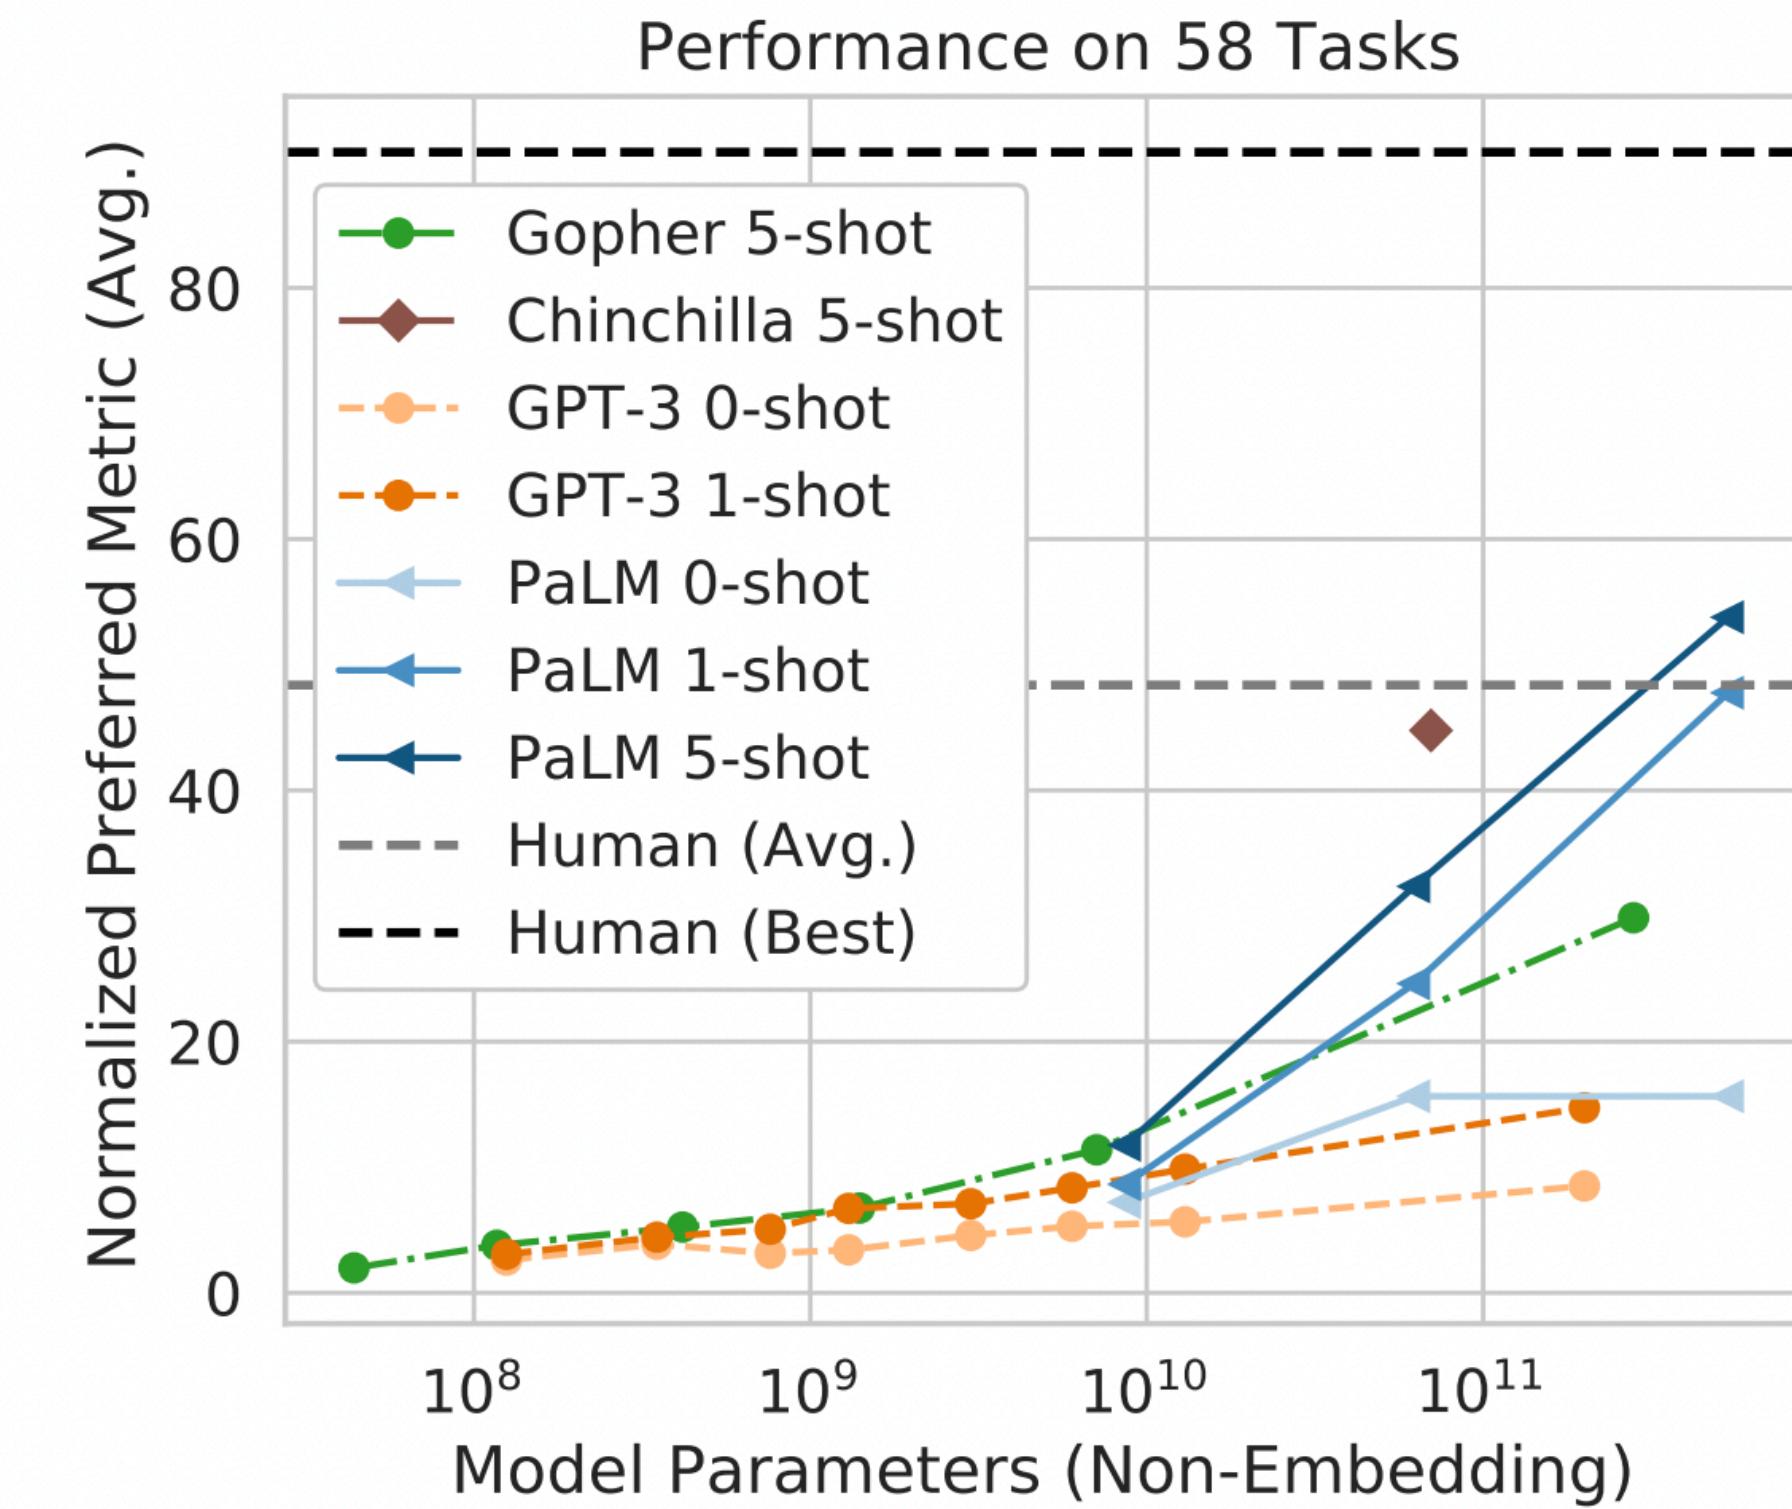
0.001 good

- ▶ Model understands some sentiment?

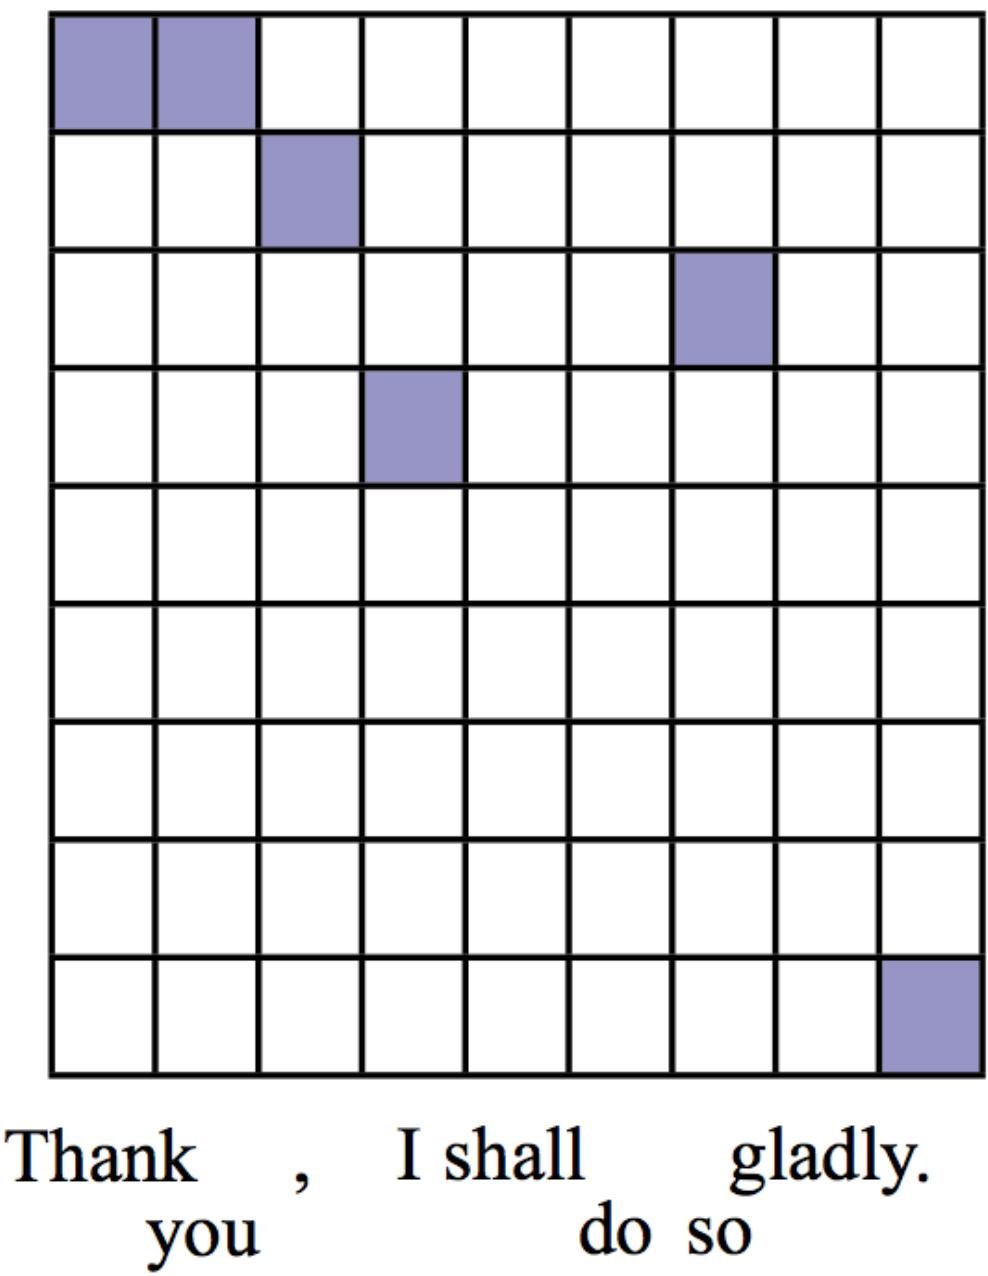
- ▶ Train a neural network to do language modeling on massive unlabeled text, fine-tune it to do {tagging, sentiment, question answering, ...}

Large Language Models

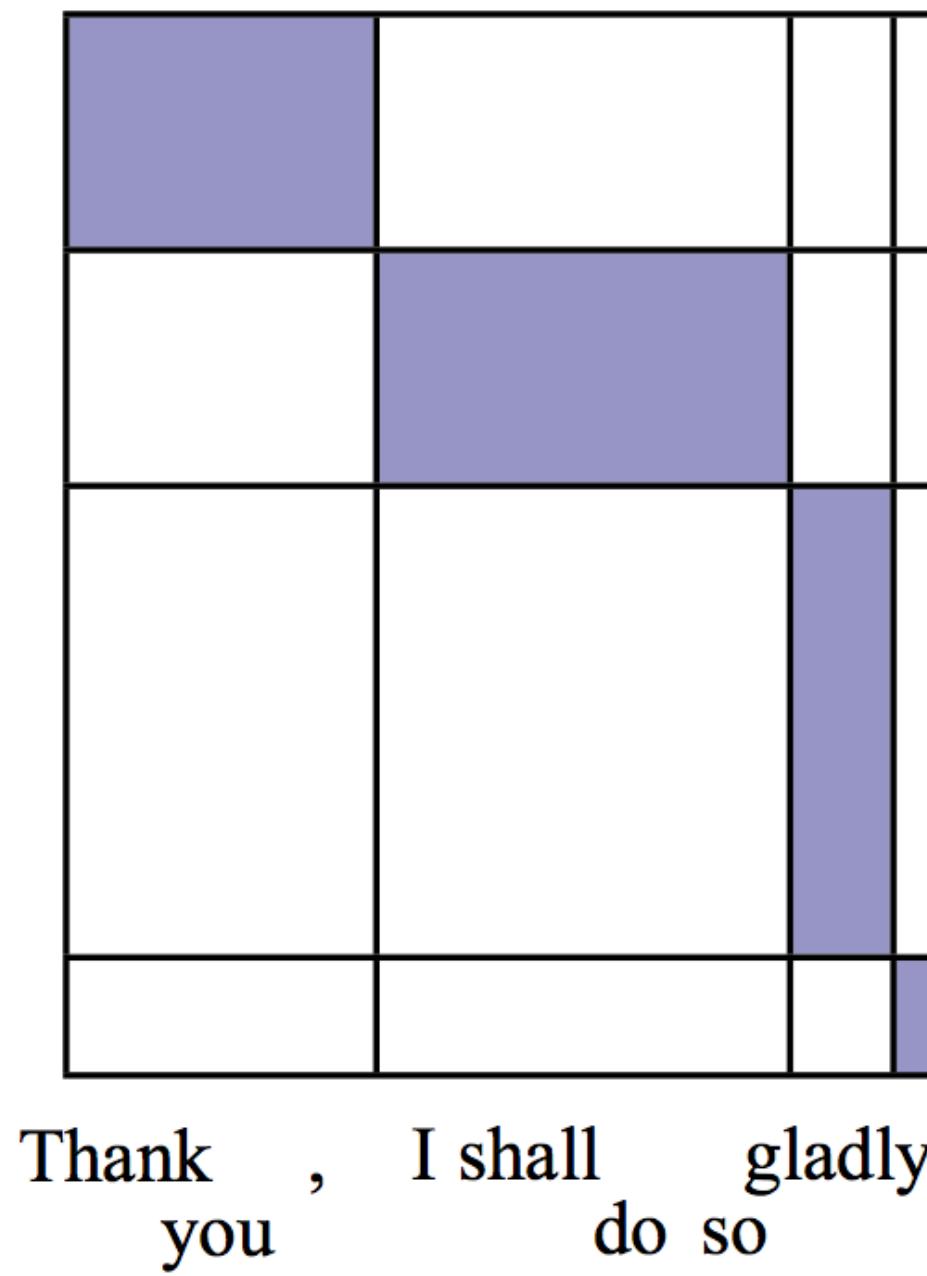
- ▶ Google's PaLM: 540 billion parameter model
 - ▶ Trained on 780 billion tokens (~780 GB)
- ▶ GPT-3 / ChatGPT
 - ▶ <https://chat.openai.com>



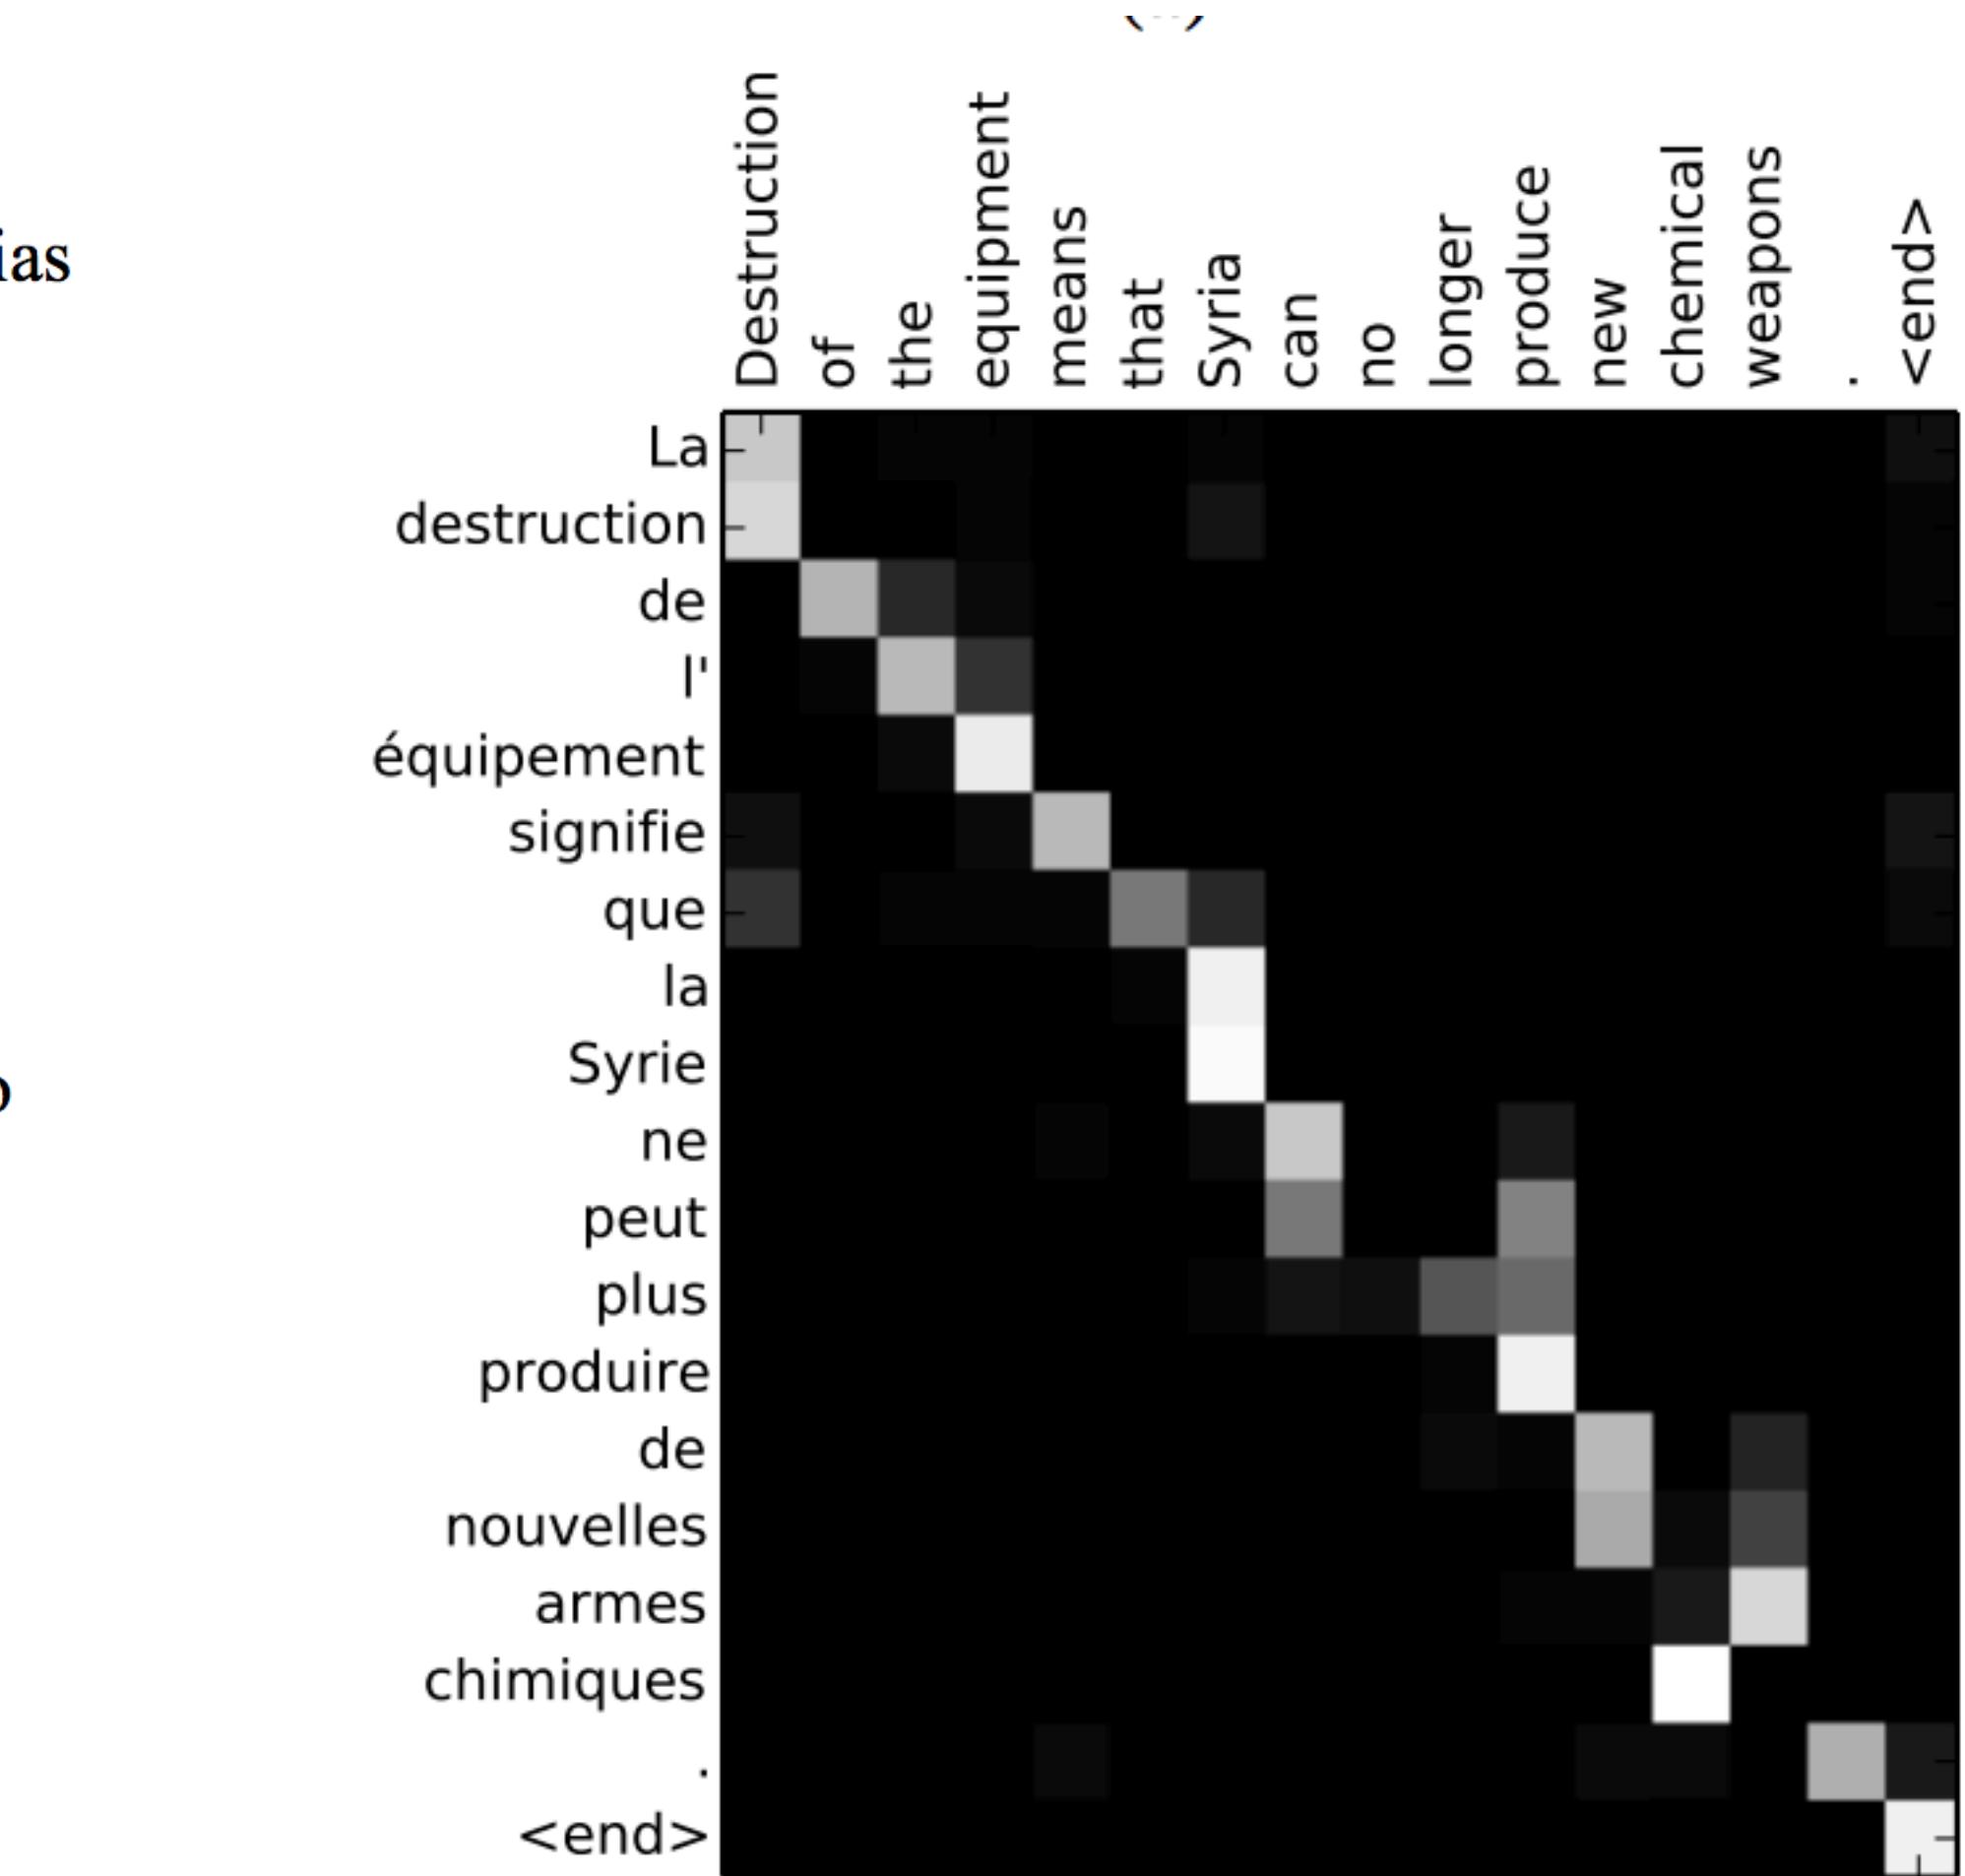
Less Manual Structure?



(a) example word alignment



(b) example phrase alignment



DeNero et al. (2008)

Bahdanau et al. (2014)

Does manual structure have a place?

Does manual structure have a place?

- ▶ Neural nets don't always work out of domain!

Does manual structure have a place?

- ▶ Neural nets don't always work out of domain!
- ▶ Coreference: rule-based systems are still about as good as deep learning out-of-domain

Does manual structure have a place?

- ▶ Neural nets don't always work out of domain!
- ▶ Coreference: rule-based systems are still about as good as deep learning out-of-domain

	CoNLL Avg. F ₁
Newswire	
rule-based	55.60
berkeley	61.24
cort	63.37
deep-coref [conll]	65.39
deep-coref [lea]	65.60
Wikipedia	
rule-based	51.77
berkeley	51.01
cort	49.94
deep-coref [conll]	52.65
deep-coref [lea]	53.14
deep-coref ⁻	51.01

Moosavi and Strube (2017)

Does manual structure have a place?

- ▶ Neural nets don't always work out of domain!
- ▶ Coreference: rule-based systems are still about as good as deep learning out-of-domain
- ▶ LORELEI: transition point below which phrase-based systems are better

	CoNLL Avg. F ₁
Newswire	
rule-based	55.60
berkeley	61.24
cort	63.37
deep-coref [conll]	65.39
deep-coref [lea]	65.60
Wikipedia	
rule-based	51.77
berkeley	51.01
cort	49.94
deep-coref [conll]	52.65
deep-coref [lea]	53.14
deep-coref ⁻	51.01

Moosavi and Strube (2017)

Does manual structure have a place?

- ▶ Neural nets don't always work out of domain!
- ▶ Coreference: rule-based systems are still about as good as deep learning out-of-domain
- ▶ LORELEI: transition point below which phrase-based systems are better
- ▶ Why is this? Inductive bias!

	CoNLL Avg. F ₁
Newswire	
rule-based	55.60
berkeley	61.24
cort	63.37
deep-coref [conll]	65.39
deep-coref [lea]	65.60
Wikipedia	
rule-based	51.77
berkeley	51.01
cort	49.94
deep-coref [conll]	52.65
deep-coref [lea]	53.14
deep-coref ⁻	51.01

Moosavi and Strube (2017)

Does manual structure have a place?

- ▶ Neural nets don't always work out of domain!
- ▶ Coreference: rule-based systems are still about as good as deep learning out-of-domain
- ▶ LORELEI: transition point below which phrase-based systems are better
- ▶ Why is this? Inductive bias!
- ▶ Can multi-task learning help?

	CoNLL Avg. F ₁
Newswire	
rule-based	55.60
berkeley	61.24
cort	63.37
deep-coref [conll]	65.39
deep-coref [lea]	65.60
Wikipedia	
rule-based	51.77
berkeley	51.01
cort	49.94
deep-coref [conll]	52.65
deep-coref [lea]	53.14
deep-coref ⁻	51.01

Moosavi and Strube (2017)

Where are we?

Where are we?

- ▶ NLP consists of: analyzing and building representations for text, solving problems involving text

Where are we?

- ▶ NLP consists of: analyzing and building representations for text, solving problems involving text
- ▶ These problems are hard because language is ambiguous, requires drawing on data, knowledge, and linguistics to solve

Where are we?

- ▶ NLP consists of: analyzing and building representations for text, solving problems involving text
- ▶ These problems are hard because language is ambiguous, requires drawing on data, knowledge, and linguistics to solve
- ▶ Knowing which techniques to use requires understanding dataset size, problem complexity, and a lot of tricks!

Where are we?

- ▶ NLP consists of: analyzing and building representations for text, solving problems involving text
- ▶ These problems are hard because language is ambiguous, requires drawing on data, knowledge, and linguistics to solve
- ▶ Knowing which techniques to use requires understanding dataset size, problem complexity, and a lot of tricks!
- ▶ NLP encompasses all of these things

NLP vs. Computational Linguistics

- ▶ NLP: build systems that deal with language data
- ▶ CL: use computational tools to study language

NLP vs. Computational Linguistics

- ▶ NLP: build systems that deal with language data
- ▶ CL: use computational tools to study language

