

CS7650 Problem Set 0 (Spring 2023)

January 8, 2023

IMPORTANT please read the following paragraphs carefully: Problem Set 0 is a brief review of mathematical concepts necessary to succeed in this class. CS 7650 will cover deep learning and other machine learning methods for natural language processing. These will be discussed in a level of mathematical detail that is commonly understood by modern NLP engineers and researchers. We will do our best to make this material easy to follow, but there is a certain level of mathematical background required, that is not possible to fully cover in this course. To succeed in the course, it is important for students to be familiar with the concepts and notation from probability, linear algebra and calculus. If you see any symbols or concepts in this assignment you don't recognize, this is a sign that you need another math course before taking CS 7650 - please feel free to reach out to the course staff to discuss.

Similarly, the course assumes a strong programming background. students who do not have sufficient programming experience are likely to struggle with the homework assignments. The course will involve a significant programming component to implement the models discussed in lecture. We expect students are comfortable programming in Python and Numpy, or are able to learn a new programming language and environment very quickly. In addition to this assignment (PS0), we also strongly recommend completing the first part of Project 1¹ on the basic Perceptron, to determine if you have the pre-requisite background to succeed in the course.

Please be aware that debugging machine learning algorithms can be time-consuming and requires careful thought. This is because they often involve a certain amount of computation and it is not always immediately clear where implementation issues may be occurring. To complete the assignments, you will need access to GPUs. When you begin working on the PyTorch portion of the assignment, you will need to subscribe to a Colab Pro account, which costs \$10 per month. This subscription is not required for the first part of Project 1 (the perceptron), which uses Numpy, but it will be needed for the second part, and subsequent assignments.

Collaboration is **NOT** allowed. All questions represent material that students are expected to be familiar with before taking this class. Please show your work and write clearly. We will not be able to give credit for answers that are not legible.

Please submit your solutions on Gradescope.²

¹<https://colab.research.google.com/drive/1trrreCuhH0HG45gWKqUZJdtl30ya0SzT?usp=sharing>

²<https://www.gradescope.com/courses/482255>

1 Joint and Marginal Probabilities

Assume the following joint distribution for $P(A, B)$:

$$\begin{aligned}P(A = 0, B = 0) &= 0.2 \\P(A = 0, B = 1) &= 0.5 \\P(A = 1, B = 0) &= 0.1 \\P(A = 1, B = 1) &= 0.2\end{aligned}$$

(a) **(1 point)** What is the marginal probability of $P(A = 0)$?

(b) **(1 point)** What is $P(B = 0|A = 1)$?

(c) **(1 point)** What is $P(A = B)$?

2 Independence

(2 points) Assume X is conditionally independent of Y given Z . Which of the following statements are always true? \mathcal{X}_Z represents the set of all possible values of random variable Z . Note that there may be more than one correct answer.

- (a) $P(X, Y) = \sum_{c \in \mathcal{X}_Z} P(X, Y, Z = c)$
- (b) $P(X, Y, Z) = P(X) + P(Y) + P(Z = c), c \in \mathcal{X}_Z$
- (c) $P(X, Y|Z) = P(X|Z)P(Y|Z)$
- (d) $P(X, Y, Z) = P(X) + P(Y) - P(Z)$
- (e) $P(X, Y) = P(X)P(Y)$

3 Bayes Rule

(2 points) There is a 20% chance that a thunderstorm is approaching at any given moment. You own a dog that has a 60% chance of barking when a thunderstorm is approaching and only a 40% chance of barking when there is no thunderstorm approaching. If your dog is currently barking, how likely is it that a thunderstorm is approaching? *Hint:* write down the relevant probabilities and apply Bayes' Rule to answer this question.

		y	
		0	1
x	0	0	1/5
	1	2/5	2/5

4 Entropy

The entropy of a random variable x with a probability distribution $p(x)$ is given by:

$$H[x] = - \sum_x p(x) \log_2 p(x)$$

Consider two binary random variables x and y having the joint distribution:

- (a) **(1 point)** Evaluate $H[x]$
- (b) **(1 point)** Evaluate $H[y]$
- (c) **(1 point)** Evaluate $H[x, y]$

5 Probability

- (a) A probability density function is defined by

$$f(x) = \begin{cases} Ce^{-x} & \text{if } x > 0, \\ 0 & \text{otherwise.} \end{cases}$$

- (i) **(1 point)** Find the value of C that makes $f(x)$ a valid probability density function.
 - (ii) **(1 point)** Compute the expected value of x , i.e., $E(x)$.
- (b) **(2 points)** Three locks are randomly matched with three corresponding keys. What is the probability that at least one lock is matched with the right key?

6 Calculus Review

Consider the following function:

$$f(x) = x^2 \log_e(x)$$

- (1 point)** For what range of x is $f(x)$ differentiable and continuous?
- (2 point)** What is the minimum value of $f(x)$ (approximated to 3 decimal points)?
- (2 point)** Evaluate the following expression:

$$\lim_{x \rightarrow 0^+} f(x)$$

Hint: Use L'Hopital's Rule.

7 Multivariate Calculus

- (a) (**2 points**) The number of members of a gym in Midtown Atlanta grows approximately as a function of the number of weeks, t , in the first year it is opened: $f(t) = 200(80 + 9t)^{1/2}$. How fast was the membership increasing initially (i.e., what is the gradient of $f(t)$ when $t = 0$)?
- (b) Let \mathbf{c} be a column vector. Let \mathbf{x} be another column vector of the same dimension.
- (i) (**1 point**) Consider a linear function $f(\mathbf{x}) = \mathbf{c}^\top \mathbf{x}$. Compute the gradient $\frac{\partial}{\partial \mathbf{x}} f(\mathbf{x})$. For this question, consider the gradient as a row vector.
- (ii) (**1 point**) Consider a quadratic function $g(\mathbf{x}) = \frac{1}{2} \mathbf{x}^\top \mathbf{H} \mathbf{x}$. Compute the gradient $\frac{\partial}{\partial \mathbf{x}} g(\mathbf{x})$.
- (iii) (**2 points**) Let $h(\mathbf{x}) = \frac{1}{2} \mathbf{x}^\top \mathbf{H} \mathbf{x} + \mathbf{c}^\top \mathbf{x}$, where $\mathbf{H} = \begin{bmatrix} 4 & 0 \\ 0 & 3 \end{bmatrix}$ and $\mathbf{c} = \begin{bmatrix} 2 \\ 6 \end{bmatrix}$. When the gradient $\frac{\partial}{\partial \mathbf{x}} h(\mathbf{x}) = 0$, what is \mathbf{x} ? Is it a local minimum, maximum or saddle point?